

Nuts and Bolts of New Ventures

Joseph Hadzima

MIT 15.393 OpenCourseWare entrepreneurship course
Lecture notes organized by [LazyingArt LLC](#) with [Video2Book](#)

Original lectures by Joseph Hadzima for MIT OpenCourseWare. Lecture notes organized by [LazyingArt LLC](#).

Contents

1	Introduction: Most Startups Fail; How to Improve Your Odds	1
1.1	Why We Are Here	1
1.2	The Entrepreneur's First Questions	2
1.3	Why Do This If Most Startups Fail?	2
1.4	Why This Course, and Who Is In The Room?	3
1.4.1	Question & Answer	4
1.5	Frameworks for Judging a Venture	5
1.6	Ideas, Execution, Timing, and People	6
1.6.1	Ideas	8
1.6.2	Execution	8
1.6.3	Timing	8
1.6.4	People	8
1.6.5	Question & Answer	9
1.7	From Vision to Pitch, Without Losing the Substance	11
1.8	Summary	12
2	Market Identification and Sales: Finding Your Customer	14
2.1	Logistics, Teams, and the Hand-Off to Customer-Finding	14
2.2	Start at the End: Money Needed, Customer Value, Customer Count	15
2.2.1	Question & Answer	15
2.3	Users, Payers, Champions, and What We Are Really Selling	16
2.3.1	Question & Answer	17
2.4	The Non-Negotiable Requirement: Customers	18
2.5	Customer Discovery by Failure: The Regain Case	18
2.5.1	Question & Answer	19
2.6	Rebuilding the Business Around the Real Customer: Nighttime Diabetes	20

2.6.1	Question & Answer	21
2.7	Channels, Free Sales Forces, and Selling Customers to Retailers	21
2.8	Segment by Motivation, Not Demographics	23
2.8.1	Question & Answer	23
2.9	Summary	24
3	How Are You Going to Make Money? The Business/Venture Model	26
3.1	From Customer Discovery to Business Model	26
3.2	MolecularWare and the First Operating Model	27
3.2.1	Question & Answer	28
3.3	Definition, Flexibility, and the Nine-Part Canvas	28
3.4	A Menu of Canonical Models	29
3.5	Netflix as a Moving Target	30
3.5.1	Question & Answer	30
3.6	Google, Data, and Predictive Value	31
3.6.1	Question & Answer	32
3.7	Industry-Specific Models: MedTech, Diagnostics, and Pharma	33
3.7.1	Question & Answer	33
3.7.2	Switching Costs, Segmentation, and the Late Examples	35
3.8	Summary	36
4	Legal Issues	37
4.1	Framing Legal Issues Along the Venture Life Cycle	37
4.2	Why Intellectual Property Matters, and What Kind of Protection We Actually Have	38
4.2.1	Question & Answer	39
4.3	Patents as Exclusion Rights, Timing Constraints, and Strategic Assets	40
4.3.1	Question & Answer	41
4.4	Patent Process, Prior Art Mapping, and Cost Discipline for Startups	41
4.4.1	Question & Answer	43
4.5	Working With Lawyers and Moving University IP Into a Startup	44
4.5.1	Question & Answer	45
4.6	Entity Choice, Tax Elections, and the Section 83 Trap	45
4.6.1	Question & Answer	48
4.7	Founder Equity as Negotiated Structure, Not Myth	48
4.8	Summary	49

Chapter 1

Introduction: Most Startups Fail; How to Improve Your Odds

Overview. Joseph Hadzima begins this MIT OpenCourseWare course the right way: not with a doctrine, not with a slogan, but with a room full of motives and questions. He first asks why we are here, then what we do not know, then why we would attempt something so painful given such poor odds. Only after that does he permit himself equations, frameworks, and diagrams. These notes follow that sequence closely. They are curated by LazyingArt LLC from the source lecture, and they are meant to read as companion notes for a practical course in venture formation rather than as startup mythology.

1.1 Why We Are Here

The lecture opens with a guide's promise. We are beginning a six-session journey through the nuts and bolts of new ventures, and the first task is to identify the motives that have brought the room together. Hadzima does not sort those motives into pure and impure kinds. He simply lays them out and lets us hear ourselves in them.

Some are here because they have never taken a formal class in entrepreneurship and want to see what the subject even looks like. Some are here because they have encountered some miserable product or process and want to make it less bad. Some have an invention and want to bring it to life. Some want to change the world. Some are drawn by the glamour and money that entrepreneurship still advertises.

The lecture's early tone matters. We are not yet being asked to judge a venture. We are being asked to notice why the possibility of a venture exerts a pull at all.

The Dropbox anecdote gives this opening its first concrete shape. A student, stuck without access to his files while traveling, does not merely complain. He converts irritation into a design imperative: make a storage system that does not "suck." The lecture uses this story carefully. It is not yet a theory of ventures. It is an example of how dissatisfaction, if pursued, can become the beginning of disciplined entrepreneurial thought.

1.2 The Entrepreneur's First Questions

From motive, the lecture widens into uncertainty. Hadzima says that he will give us guideposts and asks us to carry his questions from session to session. The effect is cumulative. We do not begin with a neat framework because the subject does not first arrive to us neatly.

1. How do we start?
2. What, exactly, do we do first?
3. What problem are we solving?
4. Does the proposed solution really answer that problem?
5. Who cares about the idea, specifically, and who is the customer?
6. How will the venture make money?
7. How long will it take to get to market, what will it cost, and what resources will be required?
8. What legal entity, if any, should be formed?
9. How can the idea be protected, if protection is the right objective?
10. Will cofounders be needed, and how will the relationship be structured?
11. How will negotiation enter, with employees, consultants, customers, advisors, and investors?
12. How do we discover what we do not know?

The lecture then adds two important qualifications. First, entrepreneurs are always negotiating because they are assembling people and resources they do not yet fully control in order to pursue a vision they do not yet fully own. Second, mistakes are not the scandal. Repeating the same mistake is the scandal. The people teaching the course, he says, have scars, and part of their use is that they may help us recognize recurring potholes before we drive directly into them.

By this point the chapter has already adopted the lecture's rhythm. We begin from personal motives, but we are not left in autobiography. We are pushed outward into operational questions. The subject becomes practical before it becomes formal.

1.3 Why Do This If Most Startups Fail?

At exactly the point where the lecture might have become merely energizing, Hadzima turns and creates resistance. Why do this at all? Why enter a line of work defined by difficulty, embarrassment, and attrition?

The success image is easy to summon: fame, fortune, public recognition. The lecture then opposes to that image the much harsher testimony of startup life. The Nvidia anecdote serves this purpose. The founder, looking back across decades, says in effect that if he had fully known what the road required, he might not have taken it. The point is not anti-entrepreneurial. It is corrective. We are meant to feel the cost before the course offers any technique for improving our odds.

The lecture gives the base rates in blunt numerical form:

$$\Pr(\text{startup success}) \approx 0.10 \approx 10\%, \quad (1.1)$$

$$\Pr(\text{startup failure}) \approx 0.90 \approx 90\%, \quad (1.2)$$

$$\Pr(\text{top-tier VC success per investment}) \approx \frac{3}{10} = 0.300. \quad (1.3)$$

The baseball analogy is part of the argument, not a decorative aside. A venture capitalist succeeding three times out of ten is likened to a .300 hitter: someone who fails often and is still a star. The analogy has a pedagogical use. It takes a discouraging ratio and translates it into a more familiar domain where high failure and excellence already coexist.

One slide, cited in the lecture, sharpens the point further by attributing most failure to poor planning and poor experience. We should not over-read the exact number attached to that slide, but the argumentative purpose is clear. Hadzima is preparing the answer to the question he has just raised. If failure were purely random, there would be little reason for a course like this. If planning and preparation matter, then a course may genuinely improve the odds.

That is why he then introduces the local MIT counterweight. The Kauffman-style numbers are not given as a proof that ventures succeed here routinely; they are given as evidence that when ventures launched in this environment do succeed, they may succeed at very large scale:

$$\frac{26,000}{120,000} \approx 0.217 \approx \frac{1}{5}. \quad (1.4)$$

On the lecture's telling, that means roughly one active MIT-founded company for every five living alumni in the period covered by the report. The same framing attributes annual revenues on the order of \$2 trillion to those companies and says that, if aggregated, they would amount to roughly the eleventh largest economy in the world.

So the lecture does not deny the ugliness of the odds. It accepts them. Then it narrows the ambition. The course will not make venture formation safe. It will try to improve the probability of landing on the successful side of an intrinsically asymmetric game.

1.4 Why This Course, and Who Is In The Room?

After the odds have been stated baldly, the lecture repairs confidence by explaining what kind of course this is. It is not about theory. It is about doing. The speakers are not academics who study entrepreneurship from a distance; they are people who either are doing the work or have already done it.

The course's origin story reinforces the same point. Students effectively demanded a January course on how to start a company. Hadzima, short on time, assembled practitioners and asked them for the most useful possible format: tell the room the three or four or five things you wish somebody had told you when you started. That origin story is itself a small entrepreneurial parable. Demand appeared first; the product was assembled in response.

Then the lecture narrows from the course as a whole to the room immediately in front of it. This matters because the first equation in the lecture does not appear on a pure mathematics slide. It appears on a slide about audience heterogeneity: students and non-student participants, different backgrounds, different interest groups, different expectations.




Figure 1.1: Audience mix and the course happiness equation.

The slide is really a room map. The audience breakdown at left and the small interest-group table at right explain why the equation belongs here. The lecture needs a compact rule for thinking about how one course can satisfy people who arrive with different prior knowledge and different hopes.

The slide writes the quotient in a stylized stacked form. We standardize it in ordinary mathematical notation:

$$H = \frac{R}{E}. \quad (1.5)$$

Definition 1.1 (The Happiness Equation). In the lecture's notation,

$$H = \frac{R}{E},$$

where H denotes happiness, R denotes reality, and E denotes expectations.

Hadzima then states the course's objective in deliberately mathematical shorthand. The goal, he says, is to maximize H . We should hear that as lecture rhetoric rather than as a literal optimization problem, but the rhetoric is useful because it forces us to ask what in the ratio can be responsibly changed.

1.4.1 Question & Answer

Question. What are R and E in $H = \frac{R}{E}$?

Answer. R is reality and E is expectations. The lecture briefly records a revealing false start: some Sloan students suggest that R must be revenue and E must be expenses. Hadzima accepts

that reading as a subset and then enlarges it. The equation is not an accounting identity. It is a more general rule for disappointment and satisfaction.

Once we write it cleanly, the lecture's verbal logic becomes immediate:

$$H = \frac{R}{E}, \quad (1.6)$$

$$H > 1 \iff R > E, \quad (1.7)$$

$$H = 1 \iff R = E, \quad (1.8)$$

$$H < 1 \iff R < E. \quad (1.9)$$

So if we underpromise and overdeliver, then $R > E$, and the ratio exceeds one. If we overpromise and underdeliver, then $R < E$, and the ratio falls below one. The lecture's little classroom equation is therefore already doing more than managing expectations for a single course. It is establishing a general rule about promises, delivery, and emotional aftermath.

Remark 1.2. The lecture jokes that if we could drive expectations to zero, we might produce infinite happiness. We should not literalize the joke into a usable model. The mathematical content we actually need is comparative: if expectations rise while reality does not, the quotient worsens.

This is also why the equation returns at the end of the lecture, after the discussion of pitch decks and Theranos. It is not a throwaway line. It is one of the lecture's organizing devices.

1.5 Frameworks for Judging a Venture

Now that the room has been calibrated, Hadzima begins compressing entrepreneurship into reusable filters. The first filter is almost embarrassingly simple, and that is part of its strength.

A venture must create value. If it creates no value, it has no subject matter. But creation alone is not enough. It must also capture enough of that value to continue. In the for-profit setting, that capture eventually takes the form of profit and cash flow. In the not-for-profit setting, it may take the form of enough attention, support, and incoming resources to keep operating. The structural distinction is important. Hadzima is careful not to reduce entrepreneurship to one legal or financial form.

The next compression is the "three whys," followed by a fourth question that appears only after the first three have begun to work:

1. **Why this?**
2. **Why now?**
3. **Why this team?**
4. **Why won't this work?**

Why this? The lecture's strongest illustration is Amy Smith's water-incubator pitch. Its power begins with scale: 1.9 billion people lacking access to clean water. Only after that scale has been felt does the mechanism appear. Biological contamination testing requires incubation. Existing incubators require electricity. The affected population frequently lacks electricity. Therefore a non-electric incubator attacks a real bottleneck. The structure is exact: large problem, narrow obstacle, specific bridge.

Why now? The FAA example gives this question its form. A technical solution may be genuinely good and still not be timely. But if institutional demand appears at the same moment that the solution becomes available, then timing and opportunity converge. The lecture's point is severe: five years early and five years late may both be failures even when the idea itself is sound.

Why this team? Prior experience can help, but it is not magic. Hadzima's Encore Computer example makes that clear. A team may have astonishing prior credentials and still fail. Prior experience can attract funding; it does not guarantee strategic coherence or executional success. For early founders without that prior experience, the question becomes sharper: what model, what judgment, what surrounding talent gives us reason to believe that this team is the right team?

Why won't this work? This question does not replace the first three. It arrives after them. Hadzima even frames it in sales language: by the time a prospect begins seriously probing limitations, interest already exists. That is why this fourth question belongs to de-risking. We are now asking not whether the venture sounds attractive, but whether its failure modes are understood well enough that we can plan around them.

By this point the lecture has moved from personal motive to evaluative machinery. But it has not yet become abstract. Each filter is anchored in a story, and each story exists to sharpen the filter rather than to replace it.

1.6 Ideas, Execution, Timing, and People

Hadzima says he dislikes frameworks, and then, sensibly, gives us one. The next compression is the lecture's central four-part model of venture success: ideas, execution, timing, and people. The lecture does something subtle here. It first presents the four-part structure in a general, almost static form, and only afterward begins to make one part of it dynamic.




Figure 1.2: Four startup factors above an empty importance-time graph.




Figure 1.3: Pocket redraw of the four-factor framework before the lower graph is interpreted.

The upper panel says simultaneous alignment. The lower panel says, in effect, not yet. Although the lecture verbally introduces it as a “curve graph,” the visible slide at this stage is only a scaffold: axes, labels, and empty space. That is exactly right narratively. The lower panel is being reserved for a specific problem that has not yet been solved.

1.6.1 Ideas

The lecture is impatient with the mere existence of ideas. Around MIT, ideas are everywhere. The real question is whether an idea is valuable, to whom it is valuable, how much value it creates, and whether that value can be defended long enough to support a venture. An idea that can be copied instantly may still be admirable; it may even be socially useful. But the lecture insists that admiration and venture viability are not the same thing.

1.6.2 Execution

Execution is where the lecture becomes particularly concrete. Hadzima uses examples of deals he passed on not to boast about discernment, but to show how execution uncertainty feels before an outcome is known.

Zipcar looked conceptually important. Renting a car by the hour is a genuine break in the consumption model of transportation. But the operational burden, especially around parking and city-by-city implementation, was for him obscure enough to justify passing. The precursor-to-eBay story works the same way. The market concept was visible early, but the trust and fraud machinery required for large-scale execution was not yet persuasive.

This is why the lecture pauses over *mens et manus*. Ideas and execution belong together. Yet even that pair is not sufficient.

1.6.3 Timing

Timing is treated with unusual seriousness. Hadzima says plainly that he has lost more money and time by being ahead of the curve than in many other ways. Timing, in this lecture, is not a matter of trendiness. It is a matter of whether the surrounding world has become compatible with the venture.

3D printing is the clean example. The technology was not “new” in the simple sense. What had changed was the supporting environment: materials, sensors, microelectronics, and the expiration of earlier constraints. Prodigy, by contrast, spent heavily to do online shopping before the user environment and browser infrastructure were ready. Fusion makes the point in an even harder form. It does not work in pieces. All of it must work together, and that makes timing inseparable from the completeness of the system.

So when the lecture asks “Why now?”, it is not asking for a slogan about urgency. It is asking whether the world and the venture have arrived at the same time.

1.6.4 People

People, Hadzima says, are the single biggest source of failure in most ventures. He first frames this with the Cheshire Cat story: if we do not know where we are going, route choice is meaningless. He then makes the point vivid with the three-founder example. One founder wants the technology to become a standard, perhaps even open source. Another wants a fast-growing company that can go public. A third wants something important but more collegial and less intense. These are not three versions of the same ambition. They are three different journeys mistakenly housed in one venture.

The lecture then adds two more people problems. Sometimes teams are funded before they understand what they are doing, and then the pivot becomes internal rather than customer-facing. And sometimes teams are too certain too early. The E-Ink example belongs here. Not knowing what we do not know is itself a people problem because it determines whether a team is teachable.

1.6.5 Question & Answer

Question. How do technical and business contributions change over time inside a founding team?

Answer. The lecture answers this with a qualitative time-dependent graph. Early on, the technical side has very high relative importance because until something works technically, there is nothing to finance, market, or scale. The business founder, by contrast, may initially be working nights and weekends, perhaps still holding another job, trying to understand the market and the eventual commercial structure. Over time the balance changes. The business side rises in importance because distribution, team building, financing, negotiation, and organization become increasingly central.

The lecture's clean symbolic summary is qualitative rather than quantitative:

$$I_{\text{technical}}(t) \downarrow, \quad I_{\text{business}}(t) \uparrow. \quad (1.10)$$




Figure 1.4: Technical importance falls as business importance rises.




Figure 1.5: Pocket redraw of the lecture’s qualitative team-balance graph.

Remark 1.3. The graph is schematic. The arrows are straight, unscaled, and not tied to a marked crossing time. The lecture explicitly says that whether the lines cross is not the point. The point is that the relative balance changes, and with it the conversation about legitimacy, structure, and equity.

Once the graph has opened the team problem, the lecture does not leave the point in diagrammatic form. It immediately turns to actual ventures and asks which of the four components aligned and which failed.

SpeechCo. This is the lecture’s success case of alignment. The idea was strong. The timing was good because compute power had finally arrived. Execution improved by restricting the system to a domain of language use rather than trying to solve everything at once. And the people had spent enough time together before launch to form a functioning team.

VideoCo. This case initially looks like an execution failure in the making. The founders seem incapable even of choosing a company name. The payoff is that they were not merely naming the company; they were choosing a strategy. The DIVA/Avid connection reveals that what looked like hesitation had a directional logic.

HIV Co. Here the idea, timing, and scientific credibility were all strong. The weakness was people. Cofounder conflict drained the venture even in the presence of first-rate backers and advisers. This is Hadzima’s cleanest case of a venture weakened primarily by internal human structure.

NanoCo. This is the lecture’s timing case. NanoFuel had technical promise and environmental appeal, but the market arrived in the wrong price regime. The lecture even gives a back-of-the-

envelope threshold:

$$p_{\text{diesel}} \approx \$0.50/\text{gal}, \quad (1.11)$$

$$p_{\text{neutral}} \approx \$2.00/\text{gal}. \quad (1.12)$$

At launch,

$$p_{\text{diesel}} < p_{\text{neutral}}, \quad (1.13)$$

so the product fell below the lecturer's rough economic-neutral point. The lecture makes the reasoning even more vivid by observing that one could not buy a gallon of distilled water for \$0.50. The argument is intentionally simple, but it is enough. A technically interesting product may fail because the world supplies the wrong price at the wrong time.

1.7 From Vision to Pitch, Without Losing the Substance

Only after all this does the lecture turn to communication. The order is essential. We are not being taught to substitute rhetoric for venture substance. We are being taught how substance is compressed when it must be explained to other people.

Hadzima begins with the mission statement or vision, sometimes also treated as a value proposition. He then immediately insists that the vision must be supported. The lecture's pyramid is a hierarchy of compression.




Figure 1.6: A supported vision from mission statement to pitch deck.




Figure 1.7: Pocket redraw of the supported-vision pyramid.

The original slide is worth keeping because the arrow targets fix the order of the tiers. The apex is the most compressed claim, and the lower tiers are progressively larger supports. One can read the entire device as repeated compression of the same venture at different bandwidths:

$$\text{mission statement} \sim 1 \text{ sentence or paragraph}, \quad (1.14)$$

$$\text{elevator pitch} \sim 30 \text{ seconds}, \quad (1.15)$$

$$\text{executive summary} \sim 1\text{--}4 \text{ pages}, \quad (1.16)$$

$$(\text{slides, minutes, font size}) = (10, 20, 30). \quad (1.17)$$

That last line is Guy Kawasaki's 10/20/30 rule: ten slides, twenty minutes, thirty-point font. Hadzima presents it as discipline. Each level asks the same question at a different scale: can the venture still make sense when compressed?

This is why the lecture then moves directly to the Theranos warning. Communication matters, but unsupported communication is not persuasion. It is deferred collapse. The lecture's earlier happiness equation becomes useful again. If reality is fixed while expectations are inflated, the ratio worsens:

$$R \text{ fixed}, \quad E_2 > E_1 \quad \implies \quad \frac{R}{E_2} < \frac{R}{E_1}. \quad (1.18)$$

So the lecture's final position is not anti-pitch. It is anti-separation. Sizzle without substance is empty, but substance that cannot be articulated will often fail to recruit the people, capital, and trust required for the venture to exist.

1.8 Summary

We can now see the lecture unfolding in the order Hadzima intended. It begins with motives, expands into operational ignorance, then creates the central tension by insisting on bad odds and painful realities. Only after that tension is fully present does it justify the course as a practical instrument for preparation. The room itself then becomes the site of the first equation, $H = \frac{R}{E}$, and that local rule about reality and expectations becomes a general rule about promises, delivery, and hype.

From there the lecture moves by successive compressions. First: create value and capture value. Then: why this, why now, why this team, and why won't this work? Then: ideas, execution, timing, and people. Finally: mission statement, elevator pitch, executive summary, and pitch deck. Each

compression is followed by examples, because the lecture never allows a framework to float free of lived ventures.

The ending is modest, and that modesty is part of the seriousness of the lecture. Luck remains. Preparation does not abolish it. It changes the odds of being ready when luck matters. That is why the lecture ends not with startup triumphalism but with a practical handoff: break the room into interest groups, meet collaborators, return for logistics, and then go find customers.

Chapter 2

Market Identification and Sales: Finding Your Customer

These notes follow Lecture 2 of MIT 15.393, in which Joseph Hadzima handles the course machinery and then turns the room over to Bob Jones for a sustained discussion of finding customers. The curation is by LazyingArt LLC. No validated mathematical frames survive from this lecture, so every displayed relation and every figure below is a cautious reconstruction from the transcript. The lecture's method is itself part of the lesson: begin with a case so simple that we are tempted to dismiss it, and then keep asking the next honest question until the real business either comes into focus or falls apart.

2.1 Logistics, Teams, and the Hand-Off to Customer-Finding

The lecture opens with logistics, but the logistics are already saying something entrepreneurial. Hadzima speaks about attendance, the written requirement, and the preference that the work be done in teams. One reason is comic and administrative: fewer separate submissions are easier to review. The serious reason is the one the lecture wants us to keep. Teams, he says, have a higher probability of success than sole entrepreneurs, which we may summarize as

$$\Pr(\text{success} \mid \text{team}) > \Pr(\text{success} \mid \text{solo}). \quad (2.1)$$

That inequality is a compact formalization of a spoken claim, not a theorem proved in the room; but it is the first practical relation of the evening, and it explains why the lecture begins with team formation rather than with product genius.

Only after the course structure is in place does Hadzima hand the room to Bob Jones. Jones immediately makes a contract with the audience. If he is going to ask for ninety minutes of attention, he says, he ought to return something worthwhile for that investment. He therefore asks two questions before he teaches anything substantive: who are you, and what do you want? The room turns out to be mixed. Some people are already on a second venture, some are in their first, some are seriously considering one after graduation, and some are simply exploring. The jokes matter here: medication is available for first-time founders; exploratory work can be as unacademic and as useful as springboard diving during IAP. The room is heterogeneous, and Jones wants the argument to be usable across that range.

This is why he makes his anti-academic turn. He mocks the sort of lecture that begins with a sentence about risk aversion, mean-variance utility maximization, and quadratic utility functions, gets dutifully written down, and leaves nobody any wiser. That joke is not the content of the lecture. It is the permission for the method. Jones tells us he will go in the opposite direction: start with examples so simple they are almost trivial, and let the complexity arrive only when the example itself demands it. That is the pedagogical pivot on which the whole chapter turns.

2.2 Start at the End: Money Needed, Customer Value, Customer Count

Jones now gives us the toy model. Suppose, he says, that he flees the overcharged intellectual world, moves to the Berkshires, becomes a grumpy old man in a tiny house, and decides to give guitar lessons. He can imagine the whole thing very easily: beginner chords, tuning, groove, rhythm, more advanced work, even how to play a mixolydian solo that still sounds like the blues. He would need a small studio, a few amplifiers, perhaps a drum machine. Then he stops himself: let us stop right there.

2.2.1 Question & Answer

What mistake did he make? The answer is the first conceptual obstacle of the lecture. We are often quick to ask whether we *could* do a venture. Jones insists that we must first ask whether we *should* do it. The bridge from the first question to the second is money. Most companies die, he says, because they run out of money or because they are on a path where they will never make any. So he reverses direction and begins at the end of the story. We want to make enough money. That means we need revenues, and those revenues should exceed our costs:

$$R > C. \quad (2.2)$$

Revenues come from customers. So the real problem is not “I have a product idea” but “how many customers do I need, and is that number sane?”

The notation here is a cleaned reconstruction of Jones’s spoken arithmetic, not notation visibly written in the room. Let M denote the amount of money we need over a chosen period, and let V_c denote the value of one customer over that same period. Then the required customer count is

$$N = \frac{M}{V_c}. \quad (2.3)$$

In this lecture V_c is not a grand lifetime-value formula. It is simply what one active customer is worth on the same monthly or yearly clock used for M .

Jones gives the numerical skeleton in one line:

$$\frac{1,000,000}{1,000} = 1,000. \quad (2.4)$$

If we need a million dollars and each customer is worth a thousand, then we need a thousand customers. Only after that division do we ask the practical question: how are we going to do that?




Figure 2.1: Transcript-derived reconstruction of Jones’s backward logic. We do not start from enthusiasm for the offering; we start from the money requirement and work backward until the venture either looks feasible or does not.

Now the toy business is run through the arithmetic. Jones wants only a modest living. On a monthly time base,

$$M = \$5,000/\text{month}, \quad (2.5)$$

$$V_c = 25 \text{ \$/hour} \times 2 \text{ sessions/month} = 50 \text{ \$/customer/month}, \quad (2.6)$$

$$N = \frac{M}{V_c} = \frac{5,000}{50} = 100. \quad (2.7)$$

A hundred customers is not frightening for a large organization. It is a very different number for one grumpy guitarist starting alone. Jones does not hide that discouraging answer. The arithmetic is supposed to hurt a little. That pain is the point. It is cheaper to discover the problem in five lines of arithmetic than in two years of effort and most of a life savings.

2.3 Users, Payers, Champions, and What We Are Really Selling

Having forced the first go/no-go calculation, Jones now asks the question he had postponed: what exactly is broken here, and who pays us to fix it? The default answer is schoolchildren. They are the obvious users. But they are not the obvious buyers. Past a certain point, Jones says quite plainly, the users are not the customers. The parents are the ones who pay. Once we notice that, the economics remain the same but the meaning changes. The child may want fun, status, escape from boredom, or maybe music itself. The parent may be buying activity, discipline, social development, or merely an experiment.

The lecture then performs its first major reset. Jones says, in effect, erase that and start over. What if the relevant customer is not the child but the adult who used to play, used to sing, used to feel cool, and misses all of that? He gives a much richer version of the service: teach a recognizable tune on Tuesday, assemble an ensemble, and let the person play it with others on Thursday. The

Customer definition	Primary payer	Approx. customer value per month	Customers needed for \$5,000/month
Schoolchildren taking lessons	Parents	\$50	100
Lapsed adults wanting to play again	Adult learner	$\approx \$400$	≈ 12.5
Older adults seeking music and company	Adult learner	$\approx \$400$	about a dozen

Table 2.1: Transcript-derived comparison of three customer definitions for roughly the same underlying teaching activity. The lesson is that the economics change when the payer and the perceived benefit change.

ticket is no longer mere instruction. It is competence regained, identity recovered, and performance made possible. The arithmetic shifts with it:

$$V_c \approx 50 \text{ \$/hour} \times 2 \text{ sessions/week} \times 4 \text{ weeks/month} \quad (2.8)$$

$$\approx 400 \text{ \$/customer/month}, \quad (2.9)$$

$$N \approx \frac{5,000}{400} = 12.5. \quad (2.10)$$

Jones rounds the answer in words: about a dozen or so.

He pushes the point once more with the older adults he has seen through volunteer music work. Here the same visible format meets yet another need: social life, purpose, companionship, getting out of the house, perhaps even the possibility of romance. The lesson is not that one should randomly mutate the market. It is that changing the customer definition changes the economics because it changes the benefit being purchased.

2.3.1 Question & Answer

Who wins if I win? Jones slows the room down over this question because it widens the frame. The first answer is customers, and that answer is correct but incomplete. Customers may become a sales force by talking to one another. End users may become champions. Economic buyers may not be the same people as the users. Adjacent businesses may benefit from our success.

His music-store example is the cleanest illustration. A lapsed guitarist goes into a shop to buy strings, stares at the instruments on the wall, and explains that he has lost his chops and has no real setting in which to play. The helpful person in the store says, in effect, I know a guy who can fix that. If the customer starts taking lessons, the store may later sell strings, better instruments, and more equipment. We win; the store wins; the player wins. This is not generic referral marketing. It is a structural point about aligned incentives.

At this stage Jones asks the other crucial question: what are we really selling? One student says romance. Jones half laughs, half accepts, and then clarifies. We are selling comfort, socialization, purpose, collaboration, and the excitement that comes when ensemble performance suddenly clicks. The anecdote about Jones himself walking into a blues jam, playing with strangers, and feeling “magic” happen is not decoration. It is a phenomenology of the benefit. The guitar lesson is the visible product. The invisible product is the life the customer gets to step back into.

2.4 The Non-Negotiable Requirement: Customers

Having extracted a surprising amount from a tiny guitar business, Jones tightens the screws. He now lists the venture questions in a blunter form. We should ask:

1. What is broken?
2. Who specifically wants a solution?
3. Are there enough of those people?
4. Will they spend money to solve the problem?
5. How are they solving it now?
6. Why is our answer compellingly better than the alternatives?
7. Who actually pays?
8. How do we find them and let them know about us?

The questions are easy, he says. The answers are not. But if we do not at least face these questions, we will drive straight into the hole and only discover it when it is too late.

This is also the moment where the lecture pivots toward investor logic. Jones repeats Hadzima's hard framing: most investors lose most of their money. He jokes that perhaps only eight out of ten fail at MIT instead of nine out of ten elsewhere, but the odds are still dreadful. So the investor wants the answer to two questions: how will you make money, and how will I make money? The order matters. If the venture cannot answer the first question, there is no point even asking the second.

From that perspective, there is only one non-negotiable requirement. We must have customers. Everything else can be renegotiated. Jones says this with the force of a room trying to get away with something and being told that it cannot. We do not have to be rich. We do not have to be beautiful. We do not even have to look especially brilliant. If we have customers, the rest of the world begins to reinterpret us. That is the bridge into the first serious case study. Enough of the abstract, he says. Here is what it looks like when smart people get this wrong.

2.5 Customer Discovery by Failure: The Regain Case

The first medical case arrives only after the scaffold is in place, and the sequencing matters. Jones wants us to hear how reasonable a bad plan can sound. He was running the clinical nutrition division of a medical company in California, a division whose economics were excellent:

$$\text{cost per liter} \approx \$3, \quad \text{sale price per liter} \approx \$70. \quad (2.11)$$

The broader company had been doing roughly two hundred million dollars a year for a long time, but managed care threatened the old profit pool. So Jones went looking for the next revenue source and found, as he says, an opportunity nobody else was pursuing. He immediately adds the correct diagnosis in hindsight: that should have been a red flag.

The field was kidney dialysis. The background numbers looked attractive:

$$400,000 \text{ patients}, \quad 2,300 \text{ dialysis centers}, \quad 8\text{--}12 \text{ lb fluid gain between treatments}. \quad (2.12)$$

The contradiction also looked obvious. Dialysis patients often need nutritional supplementation, but the standard supplements were liquids, precisely the thing these patients should not be loading up on. So the company designed a bar-like product, high in what patients needed, low in what they should avoid, and without added fluid. The product was named *Regain*.

The company did the respectable things. It ran a clinical trial even though it was not required. It published the result. The result was good: blood chemistries improved, lives were extended. It then asked opinion leaders, clinicians, and dieticians whether they would recommend the product. The answer, extraordinarily, was yes. For how many patients? All of them. How often? Seven days a week. At what price? Roughly \$3 per bar. The arithmetic looked almost too good. Jones says they began to think they had replaced the business they were afraid of losing.

Then the product was launched, and the floor dropped out:

$$\text{\$400,000 forecast} \quad \text{versus} \quad \text{\$35,000 actual.} \quad (2.13)$$

As a cleaned-up ratio,

$$\frac{35,000}{400,000} = 0.0875 < 0.10. \quad (2.14)$$

Month after month, the product came in at less than ten percent of forecast.

2.5.1 Question & Answer

What did we do wrong? Jones lets the question sit in the room because the answer comes in layers.

1. We interviewed clinicians, but the patients were the buyers.
2. We judged taste with the wrong palate.
3. We mistook clinical need for willingness to pay.
4. We misunderstood the motivational composition of the market.
5. We forgot that retailers were customers too.

The first two mistakes are concrete. Clinicians did not buy the product and did not eat it. Patients did both. When Jones's team finally spoke to patients, they learned that many did not want the product and could not afford it. The taste problem was equally sobering. In kidney failure the palate shifts, and a heavily protein-based product can taste like spoiled meat. The team had asked the wrong people how it tasted.

The third and fourth mistakes cut deeper. Jones says the real question should have been: who gets kidney failure? The transcript gives a rough decomposition,

$$33\% + 44\% = 77\% \approx \frac{3}{4}, \quad (2.15)$$

which he interprets as follows: about three quarters of the customer base had arrived there through long-term mismanagement of diabetes or high blood pressure. His conclusion is brutal but plain. A great many of these patients had not taken care of themselves before. Many were not going to start now. Clinical benefit was real, but motivation to spend was weak.

The final mistake is the one the lecture keeps widening toward. Patients still had to buy the product somewhere. Jones's company understood hospitals; it did not understand retail pharmacy.

That ignorance was expensive. A pharmacy manager has only seconds to hear a pitch. A shelf is a monetized asset. Remove a known product, insert an unknown one, and the retailer wants compensation. That is the logic of slotting fees, and Jones gives the number that made the lesson unforgettable:

$$\text{\$1,000,000 per quarter for CVS,} \quad (2.16)$$

before guaranteed sales, refunds, and chain-wide stocking requirements were even fully counted.

The failure therefore was not one error but a stack of errors, each caused by a different bad definition of customer. Jones later found a way to reach the roughly twenty-five percent of the market that had not arrived through self-neglect, and the broader company recovered. But the lesson of the year is not rescue. It is diagnosis. The consequences of not thinking through the customer system got worse and worse.

2.6 Rebuilding the Business Around the Real Customer: Night-time Diabetes

With characteristic dark humor, Jones says that entrepreneurship may be a form of mental illness, because after all this he went off and did more of it. The next medically grounded venture is the constructive answer to the Regain failure. This time the company enters diabetes, and the starting point is not a product gap but a lived problem.

The market is large:

$$10,000,000 \text{ diagnosed patients,} \quad 4,000,000 \text{ insulin users.} \quad (2.17)$$

The practical aim of diabetes management, as Jones states it, is to keep blood glucose from fluctuating wildly. If we denote the overnight glucose level only schematically by $g(t)$, then we should read the next picture as a qualitative shape, not as a clinical graph with calibrated axes. During the day, disciplined patients can spread intake and insulin across small meals and small snacks. At night that discipline is much harder to sustain. Instead people tend to eat a large bedtime snack, take a large insulin bolus, and hope the two will titrate through the night.

The hope fails for a simple reason. The food turns to glucose quickly. The spike is higher, but the support does not last longer. Jones identifies the dangerous interval as

$$2 \text{ a.m.} \leq t \leq 6 \text{ a.m.}, \quad (2.18)$$

the period in which food support has run out but insulin may still be working.

At this point the lecture changes pace. Jones explains focus groups, then rebels against the standard focus-group script. After too much science and too much intellectual talk, he bursts out and asks what the customers are actually afraid of. The answer is the first true emotional discovery of the case: some people no longer sleep in a bed because they are afraid they will get too comfortable and die in their sleep. That is the sentence around which the business is rebuilt.

So the company creates a product whose components convert to glucose at different rates and last all night. Jones says that, being marketers at heart, they called it time-release glucose. Later in the lecture he refers to the resulting business by product name, *Night Bite*. This time, he adds pointedly, it also tasted good.




Figure 2.2: Transcript-derived schematic of the overnight problem. The ordinary bedtime snack gives a sharper early spike without widening the safe interval, whereas the redesigned product is meant to spread glucose availability further into the night.

2.6.1 Question & Answer

Why would putting “for diabetes” on the label kill the business? Jones raises this as a genuine puzzle because he originally fought to be able to say exactly that on the package. The customers corrected him. It is nobody’s business, they say, that I have diabetes. I am not “a diabetic” in the public sense; I am a banker, a lawyer, an entrepreneur, an active professional. If I pull out something that looks like medicine during a meeting, people will think I am sick, frail, and professionally diminished. That is a career-limiting signal.

So the product has to look like something an elite athlete might use, not like medication. The chemistry can be medical; the public identity cannot. That is why the company gives it a non-medical name and why Jones later notes, with some satisfaction, that marathon runners ended up buying it too.

The redesign also becomes more precise because the company now talks not only to clinicians but to parents and patients. The unit is made deliberately small and divisible:

$$100 \text{ calories per unit,} \quad (2.19)$$

with scored use that allows 50, 100, or 150 calories depending on the user. A better product has now emerged. But Jones is careful not to stop there. A better product is still only half a business. The next problem is distribution.

2.7 Channels, Free Sales Forces, and Selling Customers to Retailers

Jones now slows the story down and turns to what he calls the nuts and bolts. The first discovery is that channel changes perceived value. In a focus group the company puts paper in front of users and asks two questions: where would you expect to buy a product like this, and how much would you expect to pay? The answers sort into two piles:

$$\$0.49 \text{ in a grocery store,} \quad \$1.25 \text{ in a pharmacy.} \quad (2.20)$$




Figure 2.3: Transcript-derived reconstruction of the Night Bite customer system. The patient is the economic buyer, but the educator, the pharmacy, and the wholesaler each carry distinct incentives. The venture only scales when those roles are aligned rather than collapsed into one vague “customer.”

Same product, different channel, different value. Jones says, dryly, call me silly, I like pharmacies.

The next discovery is organizational. A business like this does not sell to “the market” in one undifferentiated move. It sells through a customer system. The crucial intermediary turns out to be the certified diabetes educator. These professionals can literally be found by zip code through their association. More importantly, they have a professional frustration: patient compliance. They tell patients what to do, and the patients do not do it. If Jones can give them something patients will actually use and thank them for recommending, he has found people who win if he wins.

Several thousand educators eventually become a kind of unpaid sales force. Jones is careful about why. They are not unpaid because they are fools; they are unpaid because they are healers, and what they want is not commission but something their patients will comply with. The educators can then recommend the product to patients and even help pressure local pharmacies to stock it.

The company first sells directly, not because that is the intended endpoint but because Jones wants information. He wants reorder frequency so he can infer usage. Advertising generates inbound calls, and the numbers become a kind of live demand function:

$$100, 200, 300, 500 \text{ calls/day.} \quad (2.21)$$

At the same time, long-run usage turns out to be strikingly strong:

$$\text{reuse rate} > 300 \text{ occasions/year.} \quad (2.22)$$

That is not a market-research fantasy. It is behavioral evidence.

Jones also gives a small methodological lesson from inside sales. Closed-ended questions stop the conversation too quickly. “Did you like it?” can end with “no.” An open-ended question such as

“What did you think?” keeps the door open long enough to discover that taste may have been imperfect, efficacy was good, and the customer wants to reorder anyway.

The pharmacy-entry tactic is equally concrete. A store manager may have a purchase amount below which corporate approval is unnecessary. Jones guesses that threshold is about one hundred dollars and prices the starter kit just below it:

$$\$99.95 < \$100. \quad (2.23)$$

The risk is removed: if it does not sell, throw it away and get a refund. Then the company calls its existing customers and tells them where the product is now stocked. When the first batch sells out, the pharmacy itself becomes interested. The phrase Jones uses is exact and worth preserving: he bought a bunch of customers and then sold those customers to the retailers.

The same logic is pushed upstream. A large pharmacy may carry about eight thousand SKUs and does not want eight thousand invoices. So the company gets into wholesalers by seeding nearby pharmacies, giving the wholesaler its normal margin, and asking only to be entered into the system before reorder demand appears. Once that is done, scale becomes possible.

The planogram meeting at a major chain is the final demonstration. Jones and his colleague go into a fifteen-minute meeting with no money to pay slotting fees. They spend ten minutes saying who they are and why the product is good, and then open a briefcase full of printouts showing ten thousand existing diabetes customers in the states where the chain operates. Those customers all have one question: which pharmacy carries the product? How, they ask the buyer, would you like us to answer them? Jones then says nothing. He invokes the negotiating rule of thumb that the next person who speaks loses. After the silence, the buyer says yes. What won the meeting was not persuasion in the abstract. It was customers.

2.8 Segment by Motivation, Not Demographics

By the end of the Night Bite story the lecture has earned the right to become doctrinal. Jones turns back to the room and asks what he should have learned from all these experiences.

2.8.1 Question & Answer

The audience answers, and Jones sharpens the answers as he goes. The resulting principles are the real retrospective of the lecture:

1. Talk to your customers.
2. Let the customer definition evolve as the business evolves.
3. Reinvest early effort and money in better market knowledge.
4. Treat distribution as part of the venture, not as an afterthought.
5. Figure out what people are really buying.

That last point becomes a short theory of benefits. Nobody specifically wants the product. People want what the product does. Jones’s example is vision: we may hire glasses, contact lenses, or a LASIK surgeon, but what we really want is to see better. The same principle reaches backward

through the entire lecture. Nobody wanted “guitar lessons” in the narrow sense. They wanted social life, competence, romance, purpose, performance, collaboration. Nobody wanted “glucose” in the narrow sense. They wanted safety, sleep, and the ability to function without public stigma.

Jones then asks what market segmentation really means. Conventional marketers segment by geography, income, or gender. He argues that the entrepreneur, especially the undercapitalized entrepreneur, must segment by motivation. The transcript’s cleanest example is the product that prevents a second heart attack. The person who has just survived the first one sits at the top of the motivation pile. Nearby family members may listen next. Social peers may pay some attention after that. The broad public is the low-motivation block at the bottom:

$$\text{most motivated} > \text{related observers} > \text{adjacent social group} > \text{general public.} \quad (2.24)$$

That ordering is qualitative, not measured. Its meaning is strategic. A startup cannot afford to spend scarce resources trying to convert the least motivated part of the world. Jones calls that big bottom block a graveyard.

The comparison between Regain and Night Bite makes the principle vivid. Regain improved clinical markers, but the patient felt no immediate difference. Every six months the doctor might say, nice going, your blood chemistries are better. Night Bite produced immediate morning feedback: I think it worked. In Jones’s joking form, the user did not wake up dead. The line is facetious, but the point is exact. Motivation rises when the benefit is not only real but felt.

He then compresses his own view of marketing almost to absurdity. After acknowledging all the sophisticated mechanics, he says he can save us two years of graduate school with two bullet points, and the two bullet points are the same: find out what your customers want; then find out what your customers want. The repetition is not carelessness. It is emphasis.

The very last quick story before the closing homework lands the lesson one final time. Jones says a group of them once invented a sleep-support product and, in a foolish top-down move, decided the market was working moms. They assembled a huge list. They heard nothing. Then a few repeat purchasers emerged, and when Jones called them personally and asked why they were buying, every answer turned out to be some version of the same thing: I run marathons, I need better recovery, better sleep helps me recover, and therefore your product helps me perform. Jones thought he was selling sensitivity and care. He was selling to warriors. He says flatly that he would have died of old age before learning that from a sequence of Google ads. Talk to them.

That is why the lecture ends by narrowing, not widening. The next session will be about pitching. The homework is to pitch the business in three minutes, then in thirty seconds, then, if possible, in three sentences and about ten seconds. The compression is not a new topic. It is the compressed form of everything in this lecture.

2.9 Summary

The lecture begins with attendance codes and team formation, but its real structure is mathematical and strategic. We begin at the end of the story, ask how much money we need, estimate what one customer is worth, and derive how many customers are required. Then we ask the harder questions: who is the user, who is the payer, who is the champion, who else wins if we win, and what benefit is actually being purchased?

The Regain case shows how completely a business can fail even when the clinical result is real

and the arithmetic looks encouraging. The Night Bite case shows how the same kind of medically grounded business can succeed once fear, identity, channel, and reorder behavior are taken seriously. The final principle is severe and practical: do not aim at everybody. Find the people who want the offering most, learn what they are really buying, and build the venture around them.

Chapter 3

How Are You Going to Make Money? The Business/Venture Model

This lecture begins with a clean handoff. Bob Jones had already shown, in the preceding session, that customer choice and pricing choice are not afterthoughts; they are already the beginning of a business model. Rich Kivel now takes the next step in Joseph Hadzima's MIT OpenCourseWare course and asks the harder question: once we have an invention, a service, or a technology, how does it become revenue, investor return, and finally a sustainable business? The lecture answers this in a deliberate order. We are first given an operating story, then a definition, then a set of recognizable business-model types, then two major platform examples, and only after that a return to the industry-specific pressure of med tech. The notes should therefore unfold as the lecture unfolds: step by step, with each example motivating the next.

3.1 From Customer Discovery to Business Model

Kivel opens by making the scope of the evening unmistakable. We are no longer only asking who the customer is. We are asking how a company makes money, how it brings a product to market, and how it eventually produces returns for investors or strategic partners. In the lecture this is presented as the “nitty-gritty” of company building, and that phrase matters. The business model is not a slogan placed after the invention. It is the operating structure that carries the invention into economic life.

A compact reconstruction of the lecture's practical target is

$$\text{technology} \rightarrow \text{distribution} \rightarrow \text{adoption} \rightarrow \text{revenue} \rightarrow \text{sustainable business.} \quad (3.1)$$

Kivel does not write this in symbols, but it is the underlying motion of the lecture. A technology by itself is inert from the standpoint of a company. Between invention and revenue there lie channels, buyers, behavior, and scale.

His autobiographical setup is also part of the argument. He presents himself as someone who spent the first part of his career on the operating side of companies: sales, marketing, partnerships, then CEO work, then investing. This is why the lecture does not begin with a framework diagram. It begins with operating experience. We are meant to see that a business model is discovered, stressed, and revised in contact with actual buyers.

3.2 MolecularWare and the First Operating Model

The first sustained example is the MIT spinout MolecularWare, a bioinformatics company from the era of the Human Genome Project. The technical scene is vividly drawn. A lab is full of autonomous devices: scanners, robots, micro-arrayers, barcode refrigerators, each with its own computer and its own software. The consequence is not elegance but spreadsheet chaos. Scientists and engineers become the unwilling clerks of disconnected systems.

Kivel's first important sequence comes here:

$$\text{capture} \rightarrow \text{visualize} \rightarrow \text{data mine} \rightarrow \text{create value.} \quad (3.2)$$

This deserves to be preserved because it is the lecture's first clean reduction of a business problem to a usable structure. One first captures information from the instruments. One then visualizes it. Once the information is visible, one can interpret overexpression, underexpression, patterns, or anomalies. Only then does the downstream problem of data mining become meaningful, and only there does value begin to emerge in the commercial sense.




Figure 3.1: Transcript-derived reconstruction of the MolecularWare value chain.

Kivel then makes a point that is easy to lose if we over-academicize the story. He had no background in biotech or pharma. He jokes that he did not know how even to spell bioinformatics. But the point is serious: what mattered was not disciplinary identity; what mattered was whether the company was creating value out of the invention. That is why he broadens the example from gene-expression labs to drones, satellites, emissions sensing, and other domains. The lecture's refrain is that the value often comes from the data.

From there he moves, with almost no pause, from product structure to channel structure. When he arrived as CEO, MolecularWare had one business model: it sold to academics. That made sense because the founders sold to the people they knew. It did not scale well, and it was not very attractive to the investment community. So the business model had to change:

$$\text{direct to academics} \rightarrow \text{distributors} \rightarrow \text{OEM.} \quad (3.3)$$

The first extension was geographic. Building a sales force in Asia was impossible for a fifteen-person Kendall Square company, so distributors became the second part of the model. The second

extension was even more decisive. The company had strong software and no reason to build instruments. Other firms had strong instruments and weak software. OEM bundling then becomes the elegant move: put the software inside the instrument and let the instrument company sell the complete system.




Figure 3.2: Transcript-derived reconstruction of the channel ladder in the MolecularWare case.

3.2.1 Question & Answer

Question. What is OEM, and why did it change the company?

Answer. OEM means original equipment manufacturer. In the lecture it is not merely a vocabulary check. It is the moment at which the company changes how it reaches the buyer. MolecularWare stops trying to sell every installation itself and instead places its software inside another firm's instrument. The manufacturer now bears much of the burden of selling, distribution, and customer reach. MolecularWare receives value when the bundled instrument is sold. The route to market changes, and with it the economics of growth.

Kivel then generalizes from this example to his larger operating experience. He has taken companies

$$8 \rightarrow 85 \rightarrow 250 \rightarrow 1,000 \text{ people,} \quad (3.4)$$

and he attributes that kind of growth not only to funding but to a business model flexible enough to support the growth. That claim prepares the ground for the formal definition.

3.3 Definition, Flexibility, and the Nine-Part Canvas

Only after the operating story does the lecture pause to define its subject. This order should be preserved. We are not given a framework and then asked to decorate it with examples. We are shown a company under pressure, and then we are given the abstraction that explains what we have just seen.

Definition 3.1. A business model describes the rationale by which an organization creates, delivers, and captures value.

A faithful symbolic compression of the lecture's definitional core is

$$\text{Business model} = \text{create value} + \text{deliver value} + \text{capture value.} \quad (3.5)$$

Remark 3.2. This displayed equation is a transcript-backed compression rather than a board-written formula. It is useful because it prevents us from treating creation, delivery, and capture as separable afterthoughts.

Kivel then turns to the *Business Model Generation* canvas and immediately refuses to let it become static. The nine components are not equal-sized slices of a permanent pie. Their weight changes with stage, financing, product maturity, and strategic pressure. They blend together. A manager may oversee key resources or key activities; a marketing team may be shaping the value proposition and the segmentation of customers; but the underlying claim is that these pieces must move together.

This is why the orchestra analogy appears. If the pieces of the model move harmoniously, the result is coherent. If one part lags or drifts, there are second- and third-order consequences. Cost structure is the lecturer's favored example. If the cost structure is not designed and revised properly, the firm starts to lose money; support weakens; repeat business fades; the model begins to damage itself.

There is a more general principle here, stated almost explicitly in the lecture: a successful business must be willing to modify its model quickly in response to the market. The model is not the frozen form of success. It is the mechanism by which a company remains able to move.

3.4 A Menu of Canonical Models

From the operating case and the formal definition, Kivel deliberately broadens the field. He reminds the room that none of this should feel exotic. We are all consumers. We live inside business models every day, whether or not we name them.

The lecture names a familiar cluster:

- subscription,
- freemium,
- marketplace,
- advertising,
- direct sales,
- reseller,
- OEM.

But the point is not taxonomy for its own sake. Kivel is training recognition. Years earlier, he says, business-model talk could feel almost arcane: are you going to sell online or retail, are you going to customize like Dell or stock inventory like everyone else? Then ubiquitous connectivity arrives, and suddenly models appear that would scarcely have made sense a short time before.

Subscription is introduced through ordinary life. We used to drive to Blockbuster, rent a movie, bring it home, and return it. There was no meaningful subscription in that world. The deeper shift, as Kivel phrases it, is from owning to experiencing. Freemium extends the same logic in another register: give access now, charge later, keep the user inside the system. Marketplace models let many sellers reach buyers through a shared platform. Advertising transforms attention into revenue rather than requiring equal direct payment from each user.

The older structures remain alive. Direct sales, reseller arrangements, and OEM still matter because the route to market is often the heart of the problem. In pharmaceutical biotech, direct sales forces survive because doctors are busy and physical education of the buyer still matters. The Intel example sharpens the point even further. Intel does not care very much about selling directly to end users. It sells to the manufacturer. In Kivel's improvised phrase, this is almost business-to-manufacturer.

The lecture therefore uses familiar models not to simplify the topic, but to show that the topic has been with us all along.

3.5 Netflix as a Moving Target

The lecture's first large platform example is Netflix. The narrative begins with Blockbuster, because Blockbuster makes the problem concrete. A retail chain owns stores, staffs them, secures them, manages inventory, handles returns, and charges late fees. Netflix appears first not as a studio or global streaming platform but as a company that mails DVDs.

That first move depends on the technological moment. Kivel stresses that at the beginning there was not enough bandwidth to make downloading a feature-length film practical. One could wait hours, consume memory, watch, delete, and repeat the process. The mailed-DVD model was not primitive. It was suited to the state of the world.

This first business model already contains several important features: convenience, no late fees, centralized inventory, and later a subscription layer that lets a household hold several DVDs out at once. The lecture is careful to keep the stages distinct. First there is mailing. Then there is the subscription tiering laid on top of mailing. Then the world changes.

Netflix does not survive by a theatrical leap of reinvention. It survives because it modifies its model before the old model becomes useless. Streaming enters gradually. Later, original content. Later, live events. Later still, small but powerful adjustments in pricing and account rules.

3.5.1 Question & Answer

Question. How did Netflix avoid dying with the DVD business?

Answer. It did not mistake its first successful form for its permanent identity. Mailing DVDs solved the immediate distribution problem in a low-bandwidth world. Streaming solved the next problem. Original programming solved the problem of dependence on others' content. Live events open yet another line of relevance. The company survives not by discarding the old model in panic but by modifying it before the old logic is exhausted.

The lecture then invites arithmetic. In a subscription business the natural first approximation is

$$R \approx pN, \tag{3.6}$$

with p the average subscription price and N the subscriber count.

If both price and subscriber count change, the first-order update is

$$R' \approx (p + \Delta p)(N + \Delta N) \tag{3.7}$$

$$= pN + N\Delta p + p\Delta N + \Delta p \Delta N. \tag{3.8}$$




Figure 3.3: Transcript-derived reconstruction of the staged evolution of Netflix’s business model.

This algebra is not written in the lecture, but it formalizes exactly the point Kivel is making verbally. In a system with a very large recurring base, tiny movements in price and nontrivial movements in the paying population are economically powerful. The lecture-era scale markers are

$$\Delta N \approx 18,000,000, \quad N \approx 300,000,000. \quad (3.9)$$

So even a small Δp is amplified by a huge N , while even without a large price increase the term $p\Delta N$ can be material. This is why Kivel treats the password-sharing restriction not as a wholly new business model, but as a profitable fine-tuning of an existing one.

He then adds one more layer: personalization. Recommendation, search, and feedback make the platform sticky. The company does not merely stream a library; it uses data about our behavior to make the next interaction more likely.

Remark 3.3. The subscriber counts are clear enough in the transcript to use as scale markers. The later market-cap discussion around Netflix is garbled and should not be sharpened into a precise numerical claim.

3.6 Google, Data, and Predictive Value

If Netflix shows business-model revision across delivery format and content ownership, Google shows something different: the migration of value from visible service into data exhaust and prediction.

The opening move is search. Google begins, in Kivel’s telling, as a better answer to a frustrating problem. Search engines existed, but search still felt bad. Humans should not have to master Boolean logic merely to find ordinary things. Better search, then, is already a value proposition.

But the lecture does not let Google remain “just search.” Docs and Drive come next. Collaborative documents eliminate version chaos. Teams in different places can work in real time. Google Maps appears free and ordinary, but the analytics behind Maps are where the deeper value lives. The company knows geography, movement, and patterns at a scale that far exceeds the visible app.

A compact reconstruction is

$$\text{search/maps/docs} \rightarrow \text{data exhaust} \rightarrow \text{predictive signals} \rightarrow \text{monetizable value.} \quad (3.10)$$




Figure 3.4: Transcript-derived reconstruction of how Google turns visible services into predictive value.

Kivel uses Google’s size rhetorically rather than analytically:

$$\text{market cap} \approx \$2.4 \text{ trillion.} \quad (3.11)$$

The force of the number is not in its exact financial interpretation. It is in the contrast with the humble origin point: a company that began by fixing search became one of the most consequential data-and-infrastructure firms in the world.

The health example clarifies the mechanism. Regional search spikes for symptoms become biomarkers, then data points, then public-health signals. Those signals can inform health organizations, clinicians, and even pharmaceutical or over-the-counter product planning. The company that began by indexing web pages now participates in the practical anticipation of disease patterns.

3.6.1 Question & Answer

Question. Why is trailing data nearly useless, and what replaces it?

Answer. Because trailing data arrives after the strategic advantage has mostly evaporated. Quarterly reports tell us what happened. They do not tell us early enough what is about to happen. Kivel’s substitute is the leading indicator: shipping activity, search behavior, satellite observation, lot inventory, and any other present signal that predicts future movement.

The lecture’s asymmetry can be written as

$$\text{decision value(leading indicator)} > \text{decision value(trailing report).} \quad (3.12)$$

This is not a theorem but an operating principle. If official data tell us how many cars were sold last quarter, that information is already old. If current shipping patterns tell us that fewer cars

are leaving port now, then a prediction about next quarter becomes actionable before the official report exists.

Kivel folds this back into a more general enterprise example. A logistics product may promise

$$\Delta \text{productivity} \approx 15\%. \quad (3.13)$$

That number is not presented as a measured law. It is the sort of concrete business claim by which an offering becomes meaningful to a client. Yet even here the same warning returns: the first model may be simple, but the company must keep becoming more valuable, more relevant, and more sticky.

At this point the lecture pauses, briefly, for questions, and then turns. We have seen two major platform examples. Now Kivel narrows the field again and says, in effect: let us talk about industries.

3.7 Industry-Specific Models: MedTech, Diagnostics, and Pharma

The lecture's pivot back to industry specificity is crucial. Kivel explicitly says that pharmaceutical, diagnostic, and med-tech businesses are not interchangeable. Even inside life sciences, each sector carries a different business model.

Med tech becomes the sharpest example because it compresses technical success and commercial failure into the same space. A lab may achieve proof of concept. A device may work. A device may be safe. Yet the market question remains: who will adopt it, who will prescribe it, who will pay for it, and what service structure surrounds it?

The central tension can be written cleanly as

$$(\text{efficacious} \wedge \text{safe}) \not\Rightarrow \text{adopted}. \quad (3.14)$$

This is the lecture's most important formal reconstruction. It captures exactly the obstacle Kivel is building toward. In med tech, technical success is necessary, but it is not sufficient.

The lecture's validated visual evidence belongs here.

The slide itself is not a workflow. It is a radial framework of simultaneous pressures:

$$\{\text{efficacy, data-driven, device validation, remote service}\}. \quad (3.15)$$

3.7.1 Question & Answer

Question. Why is “safe and effective” still not enough?

Answer. Because the commercial system is not the same as the technical system. One must ask which clinicians adopt first, which specialists are influential, how fast the specialty changes its behavior, whether reimbursement exists, whether the product is direct-to-consumer or prescription-mediated, and whether the service structure supports remote or repeated use. Kivel's examples are pointed. Surgeons may be eager for new instruments; psychiatrists may be slower to adopt new




Figure 3.5: MedTech framework slide with efficacy, validation, data, and remote-service dimensions.




Figure 3.6: Pocket-safe redraw of the MedTech framework visible in the lecture slide.

technologies. A patient may see an advertisement, but if the doctor does not prescribe, the product does not sell.

That is the decisive lesson of the MedTech beat. In a simple consumer product the buyer and the user may be the same person. In med tech, the chain from innovation to payment is often multi-layered and institutionally constrained.

3.7.2 Switching Costs, Segmentation, and the Late Examples

Once this sector-specific pressure has been made vivid, the lecture widens again. The remaining examples may appear miscellaneous if we flatten them, but they are not miscellaneous. They are a final sweep of heuristics.

First comes switching cost:

$$\text{low switching cost for buyer} \Rightarrow \text{high churn risk for seller.} \quad (3.16)$$

The lecture illustrates this with mobile carriers and gym memberships. Buyers like easy exit. Sellers dislike it because it creates turnover. Yet competitive pressure may force the seller to permit exactly that low-friction movement.

This is why Mint Mobile matters in the lecture. Mint uses T-Mobile's network, but reaches a different segment: bring your own device, no long contract, cheaper monthly access, less bandwidth privilege. The core idea is not merely telecom. It is segmentation. An incumbent model may work well and still miss an entire part of the market.

The same principle reappears in brand ladders. Toyota and Lexus, Nissan and Infiniti, Ralph Lauren and its internal hierarchy of labels: these examples all serve the same lesson. A firm can create adjacent brands or adjacent offers to reach people whom the original identity could not reach. The lecture's live exchange on exact automotive genealogy wanders somewhat, but the conceptual point remains sharp.

Digital commerce then supplies a new cluster of examples. Uber coordinates rides without owning cars. Airbnb coordinates lodging without owning rooms. Rent the Runway is not merely "sharing" but a personalized access model. The replenishment economy removes the need to remember to reorder household goods. The do-it-for-me service economy expands conveniences that once seemed reserved for the wealthy.

Travel becomes the last broad example before the close. What used to require a travel agent can now be built through search or generated through AI prompts. Here again Kivel's point is not that technology changes everything in the abstract. It is that changed technology permits new models of reaching the user.

The AI coda makes the same point one final time from the infrastructure side. A data center is a business. A data center in Norway, using a particular climate advantage, a particular energy profile, a particular funding structure, and a particular compute contract, begins to look like a business model. On the application side, tools such as Otter turn transcription into summary and summary into action items. The visible app is the surface; the model lies in where it sits in the value chain and what workflow it displaces.

The closing advice is therefore earned rather than generic. Do not overcomplicate the model. Understand the market dynamics. Be flexible. The companies that survive are the ones that see the future early enough to alter themselves.

3.8 Summary

This lecture begins with a handoff from customer discovery and ends with a warning about adaptability. Between those two points it builds a coherent sequence. First, a business model is the practical structure by which value is created, delivered, and captured. Second, channel design is not secondary to product design; the MolecularWare case makes that unmistakable. Third, successful companies survive by modifying their model before the old one collapses. Netflix does this across format, content, and access rules. Google does it by turning visible services into predictive signals. Med tech shows the sharpest tension of all: technical success does not guarantee adoption. The final lesson is therefore sober and operational. We should ask of any company not merely what it makes, but how it reaches the buyer, how it gets paid, and how it will change when its first successful arrangement is no longer enough.

Chapter 4

Legal Issues

These notes follow Joseph Hadzima's MIT OpenCourseWare lecture on legal issues in new ventures, with curation in this volume by LazyingArt LLC. The lecture is practical from the first minute. We are not being invited to become amateur lawyers. We are being taught to notice the places where a venture can lose ownership, lose leverage, or create a tax problem simply by doing the right thing too late.

Hadzima therefore organizes the hour as a sequence of timing problems. He begins with the life cycle of the venture, moves through intellectual property and ownership, tightens around patents as exclusion rights and strategic assets, and then closes with entity choice, tax elections, and the Section 83 trap. The unifying question is always the same: what has to be done early, before success itself makes the legal position more expensive?

4.1 Framing Legal Issues Along the Venture Life Cycle

Hadzima lowers the temperature before he raises it. This is background information, not legal advice. The reason is twofold and worth preserving exactly in spirit: legal outcomes are fact-dependent, and the law itself changes. The lecture is thus not a catalogue of eternal rules. It is a hazard map. If we see a particular pattern, we should know enough to stop, ask questions, and get help before stumbling into a pothole.

He also tells us, almost apologetically, that he is going to overload the room. That warning matters. It establishes the pedagogical contract of the lecture. He has one shot to put the structure in our heads; the slides and video can be revisited later. So we should expect density, but not randomness.

The first structural device is the venture life cycle. Across the top sit stages: idea, pre-outside-financing, first financing, second financing, and beyond. Down the left run recurring topic rows: intellectual property, legal entity, people, and financing. Even without the slide itself, the logic is clear. We are being asked to read legal issues in temporal order. What matters at the idea stage is not identical to what matters after a priced round.

That is why he begins where founders usually postpone thinking: "I have an idea," and "I am getting ready for outside financing." At that stage we are already forced to ask who owns the idea, what kind of protection exists, what legal body should hold the rights, and whether the company has been formed soon enough to avoid avoidable trouble later. The investor will ask those questions

in diligence. Hadzima wants the founders to ask them while the answers are still cheap.

4.2 Why Intellectual Property Matters, and What Kind of Protection We Actually Have

The lecture's first quantitative move is motivational. Why worry about intellectual property at all? Because the value of firms has moved away from plant and equipment and toward intangibles. Hadzima uses a 2015 slide in which the intangible share of S&P 500 market capitalization is about 84%. The mathematical skeleton of his hierarchy is

$$\text{Intellectual property} \subset \text{Intellectual capital} \subset \text{Intangible assets}, \quad \text{Intangible share of S\&P 500 market cap} \approx 84\% \quad (4.1)$$

This nested relation matters. Intangible assets include people, processes, ideas, and organizational capability. Intellectual capital is the narrower set of ideas and processes. Intellectual property is narrower still: it is the subset that the law lets us protect, license, assign, or lose.

Definition 4.1. Intellectual property is the legally protectable subset of a venture's intellectual capital. In Hadzima's practical register, it is the part of intangible value that can be fenced, transferred, and tested in financing.

He then walks through the ladder of protection in a very deliberate order. The lecture starts with "none," which is important. We are meant to remember that legal protection is not automatic. We may simply have ideas that we choose not to protect. From there he moves to trade secret, trademark or service mark, copyright, and patent. Each step is motivated by what it protects, what it does not protect, how long it lasts, and what it tends to cost.

A concise summary of the lecture's timing structure is

$$\text{Term}_{\text{trade secret}} = \infty \quad \text{while secrecy is maintained}, \quad (4.2)$$

$$\text{Term}_{\text{patent, U.S.}} = 20 \text{ years from filing}, \quad (4.3)$$

$$\text{Intent-to-use window} \leq 3 \text{ years}, \quad (4.4)$$

$$\text{Copyright duration} \approx 70 + \text{ years}, \quad (4.5)$$

$$\text{Copyright registration cost} \approx \$100 - \$500. \quad (4.6)$$

Trade secret is the simplest to state and one of the hardest to manage well. It lasts as long as we keep the information secret. That is why Hadzima immediately connects it to non-disclosure agreements. The trade secret is not magic; it is a secret that gives market advantage, and the law only helps if we behave as though secrecy matters. The lecture's practical rule is not to explain the secret sauce unless there is a specific reason, a specific recipient, and usually an agreement around that disclosure.

Trademark and service mark are different. They identify source, not technical function. A trademark is for goods, a service mark for services. Rights arise through use, and federal registration strengthens them. Hadzima keeps the founder on practical ground here. Before registration, we

mark the claimed brand with TM or SM; only after federal registration does the registered-mark symbol become available. The mark can be a word, symbol, phrase, sound, or even color. The lecture's red-sole-shoe example is meant to make that breadth memorable.

Two practical consequences follow. First, a mark should be fanciful rather than merely descriptive. "Apple" for computers is protectable in a way that "micro dose" for low-dose aspirin is not. Second, one should search before investing heavily in branding. Rights arise through use, but Hadzima emphasizes the *intent-to-use* filing precisely because founders may spend a great deal on design and launch preparation only to learn too late that the mark is unavailable.

He also preserves a subtlety that matters to founders who think federal registration erases everything that came before. It does not. The senior unregistered user can survive in the geographic area where use began earlier. His McDonald's example makes the point vividly: a small earlier user may survive locally, even after the national chain registers federally.

Copyright is introduced with equal care. It arises upon creation, not registration. If Hadzima writes a letter, he already has a copyright in the expression. But the right protects expression, not idea or function. That is why Shakespeare's treatment of a love triangle is protectable while the abstract idea of a love triangle is not. The distinction is not philosophical decoration. It explains why copyright is a natural fit for literature and music, and a more awkward fit for software.

The lecture then becomes founder-facing again. Registration is cheap relative to patents, but if one wants to sue, registration matters. The work must be deposited, and in software that raises obvious anxiety: must we deposit source code? Hadzima's answer is procedural rather than doctrinal. There are ways to deposit redacted portions so that authorship can be proved without giving away everything.

His Aero Maps stories serve the same didactic function. Phantom roads are inserted so copying becomes detectable. The Boston directory story is even sharper: the city stopped buying updated directories because it had copied the old one. That anecdote does not merely illustrate infringement. It demonstrates why a venture should put notices on what it wants to protect, and why one should not confuse "I made it" with "I can prove I own and can enforce it."

The most important copyright lesson for founders is ownership. The author of a work owns it unless the work falls within employment or a valid work-for-hire arrangement. Hadzima's import-export software story is painful for exactly that reason: a company can spend roughly half a million dollars building software and still not own it if the contract never assigned ownership. The companion lesson is symmetric. When a startup consults for a large company, the services agreement will often attempt to transfer everything created. Founders must negotiate the difference between the deliverable and the preexisting tools, modules, and know-how used to produce it.

The lecture adds two modern edges to this section. One is open-source compliance: if we modify open-source code in ways that trigger license obligations, that becomes part of diligence. The other is artificial intelligence. Hadzima's lecture-era caution is simple enough to keep: copyright requires human authorship, AI alone is not the author, and the boundary of "sufficient human authorship" remains unsettled.

4.2.1 Question & Answer

Why worry about intellectual property if we only have an idea?

Because at the earliest stage the venture is mostly intangible. The people, the process, the know-

how, the technical insight, and the first sketches of product-market fit are the company. If value has shifted toward intangibles, then ownership and protection are not side issues. They are the first structural questions. Before we ask how fast the venture can grow, we must ask what exactly it owns, what it can protect, and what later investors will be able to underwrite.

The lecture's rhythm also matters here. Hadzima does not begin with patent doctrine. He first makes the room care about intangible value. Only then does the taxonomy of protections make economic sense.

4.3 Patents as Exclusion Rights, Timing Constraints, and Strategic Assets

Once patent appears in the taxonomy, the lecture changes pace. The earlier categories are compared and illustrated; patent is analyzed. Hadzima calls it a limited-time monopoly, but immediately sharpens that phrase into its legal meaning. A patent is a federally granted exclusion right. It lets us prevent others from making, using, selling, or importing what falls within the claims. It does not, by itself, guarantee that we are free to practice every commercially useful embodiment of our invention.

Proposition 4.2. *A patent is a right to exclude, not a complete affirmative permission to practice an invention.*

Proof. Take Hadzima's cup example. Inventor *A* patents a vessel to hold a liquid. Inventor *B* later patents a handle. A commercial cup with a handle practices both inventions. Inventor *A* cannot commercialize the handled cup without infringing *B*'s handle patent. Inventor *B* cannot commercialize the handled cup without infringing *A*'s vessel patent. Each can exclude; neither has unilateral freedom to practice the combined product. A cross-license is required. $\square$

The real-estate analogy now becomes exact enough for the lecture's purposes. The claims are the fence. They define what is inside the property right. But a fence is not a road. It marks exclusion, not universal freedom of use.

We should also preserve the deeper social logic Hadzima mentions. Patent law is a trade: society grants a limited monopoly in exchange for disclosure. That is why patent "cuts both ways." It gives a powerful exclusion right, but only after we explain the invention well enough that others can learn from it. The system is justified not because monopolies are generally loved, but because the disclosure permits later cumulative invention.

The lecture's checklist for patentability is standard but should remain explicit: the invention must be new, non-obvious, useful, and within patentable subject matter. It must also survive the timing traps created by disclosure and sale. Hadzima's compact timing rule is

$$\text{U.S. patent filing window after an enabling disclosure or an offer for sale} = 1 \text{ year.} \quad (4.7)$$

The word "enabling" does real work. "I have anti-gravity boots" is not enabling; it tells us nothing about how the device works. A detailed public explanation may be enabling. Once the disclosure is enabling, or once the invention has been offered for sale, the one-year U.S. clock begins. In many other jurisdictions the rule is harsher: pre-filing disclosure may destroy foreign patent rights

immediately. Hadzima's story of the Egyptian team that published its result in a journal before considering patent consequences is included precisely to show how a scientific success can become a commercialization loss.

Non-obviousness receives a similar practical treatment. Hadzima does not turn it into doctrinal abstraction. He points instead to arguments of the form: the literature said this could not be done, and the market is now buying it, so it could not have been obvious to a person of ordinary skill in the art. The point is not that commercial success automatically proves non-obviousness. The point is that the analysis sits inside a real evidentiary world.

The lecture also marks a historical shift: first-inventor-to-file. The timing pressure is no longer merely "invent first." It is "file first, provided you are an inventor." That is why collaborations, joint ventures, and ambiguous middle-ground inventions become dangerous. Each side may think the jointly developed improvement is still in discussion while the other side is already racing to the office.

Before leaving the doctrinal core, Hadzima inserts another essential founder reminder: just as the author owns the copyright, the inventor owns the patent unless rights have been assigned. Companies therefore insist on invention disclosure and assignment agreements. The practical rule is simple: if the company needs to own the invention, the assignment must not be left to implication.

This leads directly to freedom to operate. A founder can spend time and money building a product only to discover that someone else's patent blocks launch. The freedom-to-operate opinion is the tool Hadzima mentions here. It is not magic protection and not a judicial decree of safety. It is a law firm's reasoned view that the proposed activity does not obviously infringe. That opinion also matters because willful infringement can produce treble damages, and so documented caution changes litigation posture.

4.3.1 Question & Answer

If we have a patent, are we free to practice the invention?

No. We are free only with respect to the claim boundary we own. Other patents may still cover indispensable components, adjacent improvements, manufacturing methods, or commercial embodiments. That is why Hadzima insists on keeping two questions distinct. "Can we get rights?" is a patentability question. "Can we launch without being blocked?" is a freedom-to-operate question.

This is one of the lecture's most important separations of mechanism from recommendation. The mechanism is exclusion by claim. The recommendation is to ask about infringement before market entry, not after.

4.4 Patent Process, Prior Art Mapping, and Cost Discipline for Startups

From here the lecture pivots from doctrine to operating strategy. If patents are exclusion rights in a crowded field, then the founder's real problem is no longer just "what can be patented?" but "what should be patented, when, at what cost, and in what competitive terrain?"

Patent mapping is introduced through an arithmetic failure. Hadzima describes sending out for an outside prior-art search. It took about six months. During those six months roughly 5000 new U.S.

patents were issuing per week. The rough consequence is

$$5000 \text{ new U.S. patents/week} \times 26 \text{ weeks} \approx 1.3 \times 10^5 \text{ new patents.} \quad (4.8)$$

The precise number is only a back-of-the-envelope estimate, but the strategic conclusion is solid: a static search can be obsolete by the time it returns.

That motivates the patent map. Hadzima describes it carefully. A box is a patent or patent application. The left edge of the box marks issue or publication date. A tail extending left marks filing date. The lines connecting boxes are citations. These citations matter in a patent-specific way. They are not merely academic references. Prior art must be disclosed to the patent office, and willful failure to disclose relevant prior art can threaten the validity of an otherwise strong patent because it may be treated as a fraud on the office.

Tesla is the lecture's main illustration. One can look at a single Tesla patent and see both the prior art behind it and the later patents that cite it. One can also look at the portfolio over time and infer where the company is extending a core invention. A cluster of later patents citing a core patent may tell us where the company is building a continuation or extension strategy.

The deeper lesson is business-facing. Hadzima gives two opposite stories from SBIR-style commercialization work. In one case a cluster of patents suggested a thicket to avoid; a sparser region suggested a better entry point. In another case the same kind of cluster looked attractive, because the founders believed they had solved the very problem that those clustered filings were circling. In both cases the map functioned not as a legal ornament but as a tool for choosing a beachhead market and even for deciding whether acquisition-by-portfolio might be the right strategy.

That is the point at which the lecture asks what to patent. Large firms may reward sheer filing activity. Startups cannot afford that luxury. The founder's question is not "what claims could possibly be obtained?" but "which filings add value, create leverage, support licensing, or protect a market worth entering?"

When should we file? Hadzima's answer is layered. We file before losing rights through public disclosure or sale. We file in time to win the first-inventor-to-file race. We file in time to support a licensing or acquisition conversation with actual issued or pending rights. And if full filing is too expensive too early, we use the provisional system as a disciplined placeholder.

$$\text{Provisional priority period} = 1 \text{ year,} \quad (4.9)$$

$$\text{Provisional filing fee}_{\text{micro}} \approx \$64. \quad (4.10)$$

The lecture is precise about what a provisional does and does not do. It establishes a filing date. It need not contain claims. But it must contain a meaningful description of the invention. Only what is described gets the earlier priority. This is why Hadzima's fruits-and-vegetables example matters. If protein becomes part of the invention later, and protein was not described in the provisional, the original filing date does not magically expand to cover the new matter.

That makes the provisional less like a complete patent and more like a disciplined temporal marker. Hadzima even mentions the startup version of this: a conference paper can be turned into a provisional by filing it with a cover sheet, provided it truly describes the invention well enough. For a cash-constrained startup, that can be the difference between preserving rights and losing them.

His own example from IP Vision makes the discipline concrete. Development ran in six-week cycles. Before deployment, the team asked what had been built, what had been thought through but not yet implemented, and what was important enough to file. Provisional applications became a routine timing tool rather than a rare heroic event.

The structure of a full application also enters here. Hadzima describes it almost as a term paper: field of invention, background, summary, detailed description, and then claims. The crucial doctrinal point returns: the detailed description must disclose the invention and its best mode of practice. The claims then define the boundary. Because claim drafting is the most technical part of the process, this is where he wants founders to spend money on qualified patent counsel rather than improvising.

Costs now become explicit:

$$\text{Patent preparation cost} \approx \$5,000 - \$15,000, \quad (4.11)$$

$$\text{Outlier patent-preparation bill} \approx \$35,000, \quad (4.12)$$

$$\text{Foreign patent cost}_{10 \text{ major countries}} \approx \$330,000 - \$500,000. \quad (4.13)$$

The \$35,000 anecdote is not a statistical estimate; it is a governance lesson. The co-founder kept calling the patent lawyer directly, adding changes, and running the meter. Hadzima's response was managerial, not theoretical: no one talks to the patent lawyer without coming through a single gatekeeper. Patent cost is partly doctrinal, partly procedural, and partly behavioral.

Prosecution cost is open-ended in the same way. The examiner often rejects first claims, and Hadzima even says that can be a good sign, because immediate allowance may mean the applicant tried to claim too little. But every back-and-forth round costs money. Fixed-fee patent packages may therefore yield claims that are formally allowed and strategically weak.

Foreign filing raises the economic question again. If an exclusive licensee can be brought in early, part of the bargain may be that the licensee carries later patent costs. This turns timing of licensing into a financing strategy for IP itself.

4.4.1 Question & Answer

Should we patent this idea, or disclose it to block others?

Sometimes the right move is not to patent everything, but to patent the core and disclose the periphery. Hadzima's "picket fence" example explains why. Suppose we hold a strong core patent on the vessel to hold a liquid. Others might patent the top, the sleeve, the handle, or the insulating shell and thereby block commercial use of the core invention. If the core patent is strong, a cheap counterstrategy is to disclose those peripheral possibilities without claiming them. The disclosure becomes prior art and makes those later blocking patents harder or impossible to obtain.

This is not a universal rule. It depends on confidence in the strength of the core patent. But it is exactly the kind of startup arithmetic the lecture wants us to learn: spend carefully, disclose strategically, and think in terms of market control rather than filing count.

Tesla is the larger version of the same point. Hadzima uses Tesla not to preach "open patents," but to broaden the notion of patent strategy. If the goal is to sell electric vehicles, complementary

infrastructure matters. Giving some patent access away may stimulate the ecosystem the company needs. The legal instrument is still patenting; the business objective is now ecosystem design.

4.5 Working With Lawyers and Moving University IP Into a Startup

At this point the lecture compresses the IP discussion into founder checklists. Employees should sign invention disclosure and assignment agreements. Consultants should be governed by proper work-for-hire or assignment language. Non-disclosure agreements should be used where appropriate, but with eyes open: venture investors often refuse them because they see many similar companies and do not want later claims that they misused one founder's disclosure when funding another.

The operational principle remains the same. Do not disclose the secret sauce until the last responsible moment. Make sure the company owns what it thinks it owns. Understand the patent landscape over time, not once. Preserve rights by filing on time.

Hadzima uses even the provisional-application discussion to make a cultural point. Large companies sometimes resist provisional discipline because it is hard to track. Startups cannot afford that luxury. Their problem is not internal paperwork elegance. Their problem is losing rights while moving too slowly.

Then comes the lecture's comic but serious interlude on lawyers. Legal expertise is bought by the hour. If we dump a lab notebook on the desk and say "write me a patent," we are buying expensive confusion. If we wait until the one-year deadline is nearly gone and demand a filing by Monday, we are buying emergency pricing. If instead we prepare invention disclosures, read patents in the area, draft the technical description, and let counsel focus on claims and legal framing, we turn legal spend into leverage.

The next founder question then arrives naturally: what happens when the IP comes from a university? Hadzima had been asked about this on a founders' panel, and he answers by inserting a quick institutional history. Before Bayh-Dole, federally funded university inventions belonged to the federal government, were managed from Washington, and were typically licensed non-exclusively. For a startup that wants to build a business around a specific technology, this was a poor commercialization regime.

$$\text{Pre-Bayh-Dole commercially licensed patents} \approx 5\%. \quad (4.14)$$

Bayh-Dole changes the geometry. Universities retain ownership, file patents, license them, prefer small businesses, give the government a royalty-free license, build commercialization programs, and share royalties with inventors. The modern technology licensing office is the institutional consequence of that change.

MIT appears in the lecture as a comparatively founder-friendly case. Hadzima describes the Technology Licensing Office as easier to work with than many counterparts because its first priority is getting inventions into the world, not maximizing immediate revenue. That difference in mission shapes terms and process.

The ownership rules matter. Significant use of MIT facilities or MIT-administered funds typically means MIT owns the invention. MIT does not assign the IP outright; it licenses. Sponsors may have first licensing rights. The practical moral is plain: if we are building around university technology, we must know whether the invention belongs to the company, the founder, the sponsor, or the

university. Investors will eventually force the issue if the founders do not.

Hadzima's Ali example and his Kendall Square explanation show how this becomes operational. If the work was not done with significant use of MIT facilities or funds, the founder may obtain a letter clarifying non-ownership by MIT. If the invention does belong to MIT, the startup typically develops it across the street, in commercial space, not in university labs. That physical move mirrors the legal move.

4.5.1 Question & Answer

How do we take university IP into a startup without quietly mis-owning it?

By bringing it through the front door. Hadzima's recommended device is often an option to obtain a license. The option gives the startup time to explore markets, talk to investors, and raise capital without pretending that the technology has already been transferred into the company. It is the institutional analogue of the provisional patent application: a way to preserve position while uncertainty is still being resolved.

The lecture preserves several useful lecture-era financial ranges:

$$\text{University license issue fee} \approx \$50\text{K} - \$150\text{K}, \quad (4.15)$$

$$\text{University running royalty} \approx 3\% - 5\%, \quad (4.16)$$

$$\text{MIT startup equity stake} \approx \text{single-digit } \%. \quad (4.17)$$

He also describes a specifically startup-oriented equity practice in which the university may take a small equity position and keep that percentage fixed through the first outside financing round. In his example, if MIT holds 5% and an investor later buys 20%, MIT still sits at 5% through that first round rather than being diluted immediately with the founders. That is not a universal law of licensing. It is a lecture-era practice point, and it changes how founders should think about early capitalization.

4.6 Entity Choice, Tax Elections, and the Section 83 Trap

The lecture now tightens again. After asking who owns the technology, Hadzima asks what legal body owns the company. If founders do nothing, they may already have a general partnership: one or more persons carrying on a business for profit. That requires no written agreement, which is precisely why it is dangerous. Liability is joint and several. Ownership, absent a contrary agreement, follows money rather than founder effort. Neither feature is hospitable to a venture-backed startup.

The practical recommendation is therefore immediate incorporation, usually as a corporation, and often in Delaware if repeated outside financing is likely. Delaware matters here not as brand mystique but as a jurisdiction with well-developed corporate law and comparatively predictable outcomes. Hadzima also insists on corporate formalities. A corporation only protects if it is actually treated as a separate legal person in records, signatures, and operations.

Tax then enters through that corporate personhood. A corporation taxed under Subchapter C bears tax at the corporate level, and distributions to shareholders are taxed again. A Subchapter S election changes the structure by making the income pass through to the owners.

C-corp asset sale : corporate-level tax + shareholder-level tax, (4.18)

S-corp asset sale : pass-through to the owners without corporate-level tax. (4.19)

The eligibility conditions are narrow enough that Hadzima states them explicitly:

S-corp shareholder cap < 100, classes of stock = 1, S-election filing window $\approx$ 3 months after incorporation (4.20)

A venture capital investor usually destroys S status because the shareholder rules are restrictive and only one class of stock is permitted. But Hadzima does not let the audience dismiss the election as irrelevant. Founders say: we will be loss-making for years, so why worry now? His answer is the exit. Small acquisitions are often structured as asset sales. In that case, the C corporation pays tax on asset gain, while the S corporation does not bear corporate-level tax in the same way.

The lecture gives two vivid numbers here: an illustrative exit-tax difference on the order of 13%, and the DIVA story in which an early S election reportedly saved about \$2,000,000. These are example numbers, not timeless constants, but they preserve the lecture's logic: early entity and tax choices can matter only at the end, which is exactly why founders neglect them.

Then Hadzima marks the climax audibly: if we are falling asleep, wake up. This is the big trap. Section 83 says that if we receive property in connection with services, we recognize ordinary income equal to the fair market value of the property minus what we paid for it. In the lecture's notation bank, the core formula is

$$I_{\text{ordinary}} = \text{FMV} - P. \quad (4.21)$$

To make the tax consequence explicit, we may write

$$T = \tau(\text{FMV} - P), \quad (4.22)$$

where T is the tax owed and τ is the marginal tax rate used for the example. Hadzima gestures at "37% or whatever" and then computes with 40%, so we keep the symbolic formula and treat 40% as a rough example rate.

Now the lecture turns the formula onto the founders themselves. Suppose an investor pays \$1,000,000 for 50% of the company. What is the founder's 50% worth? Hadzima first lets the room infer a simple answer: perhaps about \$500,000. If the founder paid essentially nothing for those shares, then

$$I_{\text{ordinary}} \approx \$500,000, \quad (4.23)$$

$$T = 0.40 \times \$500,000 = \$200,000. \quad (4.24)$$

He then pushes the room toward an even harsher inference. The wording is damaged in the transcript, but the intended arithmetic is clear enough: if the founder's block is instead treated as worth about \$1,000,000, then

$$I_{\text{ordinary}} \approx \$1,000,000, \quad (4.25)$$

$$T = 0.40 \times \$1,000,000 = \$400,000. \quad (4.26)$$

The economic violence of the example is the point. Founder stock can create a personal tax bill long before the founder has cash liquidity.

Then comes the second blow: vesting. If the stock vests over time, the income may be measured when the stock vests rather than when it is first issued. So the tax base grows with success. The lecture's timing structure is

$$T(t_v) = \tau(\text{FMV}(t_v) - P), \quad \text{FMV}(t_v) \gg \text{FMV}(t_0), \quad (4.27)$$

where t_0 is the early issuance date and t_v the later vesting date. Hadzima's plain-English version is better than any abstraction: when the company begins, the idea is a dime a dozen; when the stock vests two years later and the company has done well, the shares may be worth \$100 each. That is exactly when the tax becomes lethal.

His California CFO story is included to show that this is not theoretical. A later executive asked for the 83(b) election paperwork, and the company suddenly realized it had issued a large amount of vesting stock without understanding the tax consequences. What looked like a successful company was also a ticking time bomb.




Figure 4.1: Section 83 trap slide with timing-versus-priority chart. The slide itself gives the operational recommendation: separate founder stock issuance from later outside investment and make the 83(b) election.

The screenshot matters because it preserves the lecturer's own wording. The top bullet says, in essence, separate in time the issuance of founder stock from the later investment by others:

incorporate earlier, issue stock, and make the 83(b) election. The second bullet asks why stock does not get issued in time. The visible answers are embarrassingly human: the founders are too busy, and they are not yet sure who should get what.




Figure 4.2: A pocket-scale reconstruction of the visible timing sketch. The frame supports only a qualitative claim: technical work is early and highly important, and that very urgency helps crowd out timely stock allocation.

4.6.1 Question & Answer

Why does stock not get issued in time, and how does an 83(b) election avoid the trap?

Because early ventures postpone the awkward split conversation while technical work, customer work, and financing consume attention. By the time the founders resolve who gets what, an outside investment or other valuation event may already have made the stock expensive for tax purposes. The 83(b) election changes the measuring time. It aims to tax the stock at the earlier, lower-value moment of transfer rather than the later, higher-value moment of vesting. But that only works if the stock has actually been issued, and if the issuance can be separated in time from the later event that makes valuation obvious.

That is why the lecture's recommendation is temporal, not merely formal: incorporate early, issue the stock early, let later cash investment occur later, and then file the 83(b) election promptly.

4.7 Founder Equity as Negotiated Structure, Not Myth

The lecture does not end by closing the tax folder and declaring victory. It spills, deliberately, into founder culture. How much equity should each founder have? Hadzima refuses to turn this into mythology. The amount of equity and who gets it is, in his words, about culture and negotiation.

His Keenan Systems example serves as a counter-myth. One founder kept all of the stock, paid employees very well, and later distributed bonuses when the company was acquired. That was a culture, not a universal fairness theorem. The lecture uses it to break the habit of thinking that there is a single morally canonical founder split.

The variables Hadzima wants on the table are concrete: value of past contribution, value of future

contribution over a specified time horizon, ownership of IP, sacrifice, commitment, reputation, opportunity cost, and the signaling effect of who is actually all-in. His professor-and-postdoc story preserves that method. Each person writes down what they contribute and what they give up. The professor brings name recognition, perhaps access to capital, perhaps one consulting day a week under MIT policy. The postdoc gives up a salary path and perhaps an academic career. The conversation becomes fairer not because someone solved it algebraically, but because the actual variables became visible.

The lecture then returns to numerical structure in the option-pool discussion. The equity compensation pool is not a decorative cap-table entry. It is a reserve sized for the next hiring horizon, typically two to three years and tied to the headcount plan. Hadzima's lecture-era range is

$$\text{Equity compensation pool after the first outside round} \approx 12\% - 18\%. \quad (4.28)$$

The lower end corresponds to a company that already has more of the needed senior team. The upper end corresponds to a company that still must recruit experience. Once again the lecture translates culture into timed structure. Who gets the pie and how large are the slices? We are supposed to have reasons, not slogans.

The ending remains intentionally unfinished. Hadzima notes that later lectures and slides will treat ownership examples, dilution, and financing instruments in more detail. That incompleteness is part of the chapter's rhythm. Founder equity is introduced here not as a closed calculus, but as the bridge from legal formation into financing.

4.8 Summary

The lecture moves from a high-level claim about intangible value to a founder-level tax trap, but the thread tying the whole argument together is timing. We ask when an intangible becomes property, when a disclosure becomes enabling, when a provisional must become a full application, when university ownership must be formalized by option or license, when a corporate election matters only because an exit may later occur, and when founder stock becomes taxable income.

If one lesson remains after the details blur, it is this: ventures are often injured not by exotic legal doctrine, but by ordinary legal rules encountered too late. Ownership, disclosure, assignment, filing, incorporation, election, vesting, and valuation all sit on a clock. Hadzima's practical achievement in this lecture is to make us see that clock before it starts billing us.