

Financial Markets

Robert J. Shiller

Lazyearn track, Yale lecture series

Transcript-derived notes curated by [LazyingArt LLC](#) with [Video2Book](#)

Original lectures by Robert J. Shiller. Transcript-derived course notes curated by [LazyingArt LLC](#).

Contents

1	Introduction and What This Course Will Do for You and Your Purposes	1
1.1	Finance as a Pillar of Civilized Society	1
1.2	Philosophy, Detail, and a World Perspective	2
1.3	Mathematics, Diversification, and the Real World	3
1.4	Why Everyone Needs Finance	4
1.5	Wealth, Scale, and Moral Purpose	5
1.5.1	Question & Answer	5
1.6	Finance in Action: Speakers, Regulation, and Behavioral Finance	6
1.7	Preview of the Mathematical and Institutional Spine	7
1.8	Summary	8
2	Risk and Financial Crises	9
2.1	Probability, Narrative, and the Crisis	9
2.2	Return and Measures of Central Tendency	10
2.2.1	Question & Answer	12
2.3	Variance and the First Language of Risk	13
2.4	Covariance, Correlation, and Independence	14
2.4.1	Question & Answer	15
2.5	The Law of Large Numbers, VaR, and CoVaR	16
2.6	Market Return, Beta, and Idiosyncratic Risk	17
2.7	Normality, Outliers, and Fat Tails	18
2.7.1	Question & Answer	19
2.8	Summary	19
3	Technology and Invention in Finance	20
3.1	Finance as Engineering	20
3.2	Probability Review and the Central Limit Theorem	21

3.2.1	Question & Answer	23
3.3	From 1970 to the Next Forty Years	23
3.4	Risk, Inequality, Framing, and the Slow Acceptance of Invention	24
3.5	Limited Liability and Organizational Invention	25
3.5.1	Question & Answer	26
3.6	Inflation Indexation and the Chilean UF	26
3.7	Swaps, Credit Default Swaps, and the Ambivalence of Innovation	28
3.7.1	Question & Answer	28
3.8	Summary	29
4	Portfolio Diversification and Supporting Financial Institutions	30
4.1	Finance as Institutional Invention	30
4.2	The Equity Premium Puzzle	31
4.2.1	Question & Answer	31
4.3	Markowitz's Breakthrough and the One-Risky-One-Riskless Case	32
4.3.1	Question & Answer	32
4.4	Two Risky Assets and the Efficient Frontier	34
4.5	Three Assets, Diversification, and Oil	36
4.6	The Riskless Asset Returns: Tangency, Mutual Funds, and Market Risk	38
4.7	Summary	40
5	Insurance, the Archetypal Risk Management Institution, its Opportunities and Vulnerabilities	41
5.1	Insurance as a Risk-Management Institution	41
5.2	From Ancient Risk Pooling to the Seventeenth-Century Insurance Idea	42
5.2.1	Question & Answer	42
5.3	Probability, Aristotle, and the Mathematical Logic of Pooling	43
5.3.1	Question & Answer	44
5.4	Insurance as Institutional Invention: Contract Design, Moral Hazard, and Selection Bias	44
5.4.1	Question & Answer	45
5.5	AIG and the Failure of Independence	46
5.6	Guarantee Funds, State Regulation, and the American Regulatory Structure	47
5.6.1	Question & Answer	47
5.7	Types of Insurance, Health Insurance, and the Unfinished Technology of Risk Management	48

5.8	Summary	49
6	Guest Speaker David Swensen: The Yale Model and Long-Horizon Investing	51
6.1	Opening challenge: the Yale model under criticism	51
6.2	1985 endowments and the two failed tests	52
6.2.1	Question & Answer	53
6.3	Three tools of return	53
6.3.1	Question & Answer	54
6.4	Long horizons, equity risk, and the trap of literalism	54
6.4.1	Question & Answer	55
6.5	Market timing and security selection as negative-sum behavior	55
6.5.1	Question & Answer	56
6.6	Where active management matters	57
6.7	Returning to Barron's: diversification, alternatives, and the institutional coda	57
6.7.1	Question & Answer	58
6.8	Summary	59
7	Efficient Markets	61
7.1	Opening paradox: efficient markets and David Swensen	61
7.1.1	Question & Answer	61
7.2	Sharpe ratios, hidden risk, and the temptation to game the statistics	62
7.2.1	Question & Answer	63
7.3	Integral Investment Management, disclosure, and the moral of the Swensen recap . .	64
7.4	Before the phrase: Gibson, Reuters, and Conant	64
7.4.1	Question & Answer	66
7.5	From doctrine to revision: Roberts, Fama, CRSP, and the softening of orthodoxy . .	66
7.6	Technical analysis as the foil	67
7.6.1	Question & Answer	67
7.7	Random walk, AR(1), and the half-truth conclusion	67
7.7.1	Question & Answer	70
7.8	Summary	71
8	Theory of Debt, Its Proper Role, and Leverage Cycles	72
8.1	Why interest rates need an explanation	72
8.2	From Bohm-Bawerk to Fisher's two-period framework	73

8.3	Robinson Crusoe, first alone and then in a loan market	74
8.3.1	Question & Answer	75
8.4	Debt arithmetic: discount bonds, compounding, and present discounted value	76
8.5	Perpetuities, annuities, and conventional bonds	77
8.6	Forward rates from the term structure	79
8.6.1	Question & Answer	79
8.7	From the theory of interest to the problem of usury	81
8.7.1	Question & Answer	81
8.8	Summary	82
9	Corporate Stocks	83
9.1	Corporate Stocks as an Organizational Form	83
9.2	The Scale of Listed Equity in the World Economy	84
9.3	Corporation, Equity, and Governance	85
9.4	Shares, Market Capitalization, and the Logic of Dividends	86
9.4.1	Question & Answer	87
9.5	Dividend Pricing and Financing Mechanics	88
9.5.1	Question & Answer	89
9.6	Do Stock Markets Matter for Big Firms, and How Do Firms Choose Dividends? . .	90
9.6.1	Question & Answer	90
9.7	Xerox and Microsoft: Book Equity and Market Value	92
9.7.1	Question & Answer	92
9.8	Summary	93
10	Real Estate	94
10.1	Mortgage Lending in Long Historical Perspective	94
10.2	Property Rights as the Precondition for Mortgage Markets	95
10.2.1	Question & Answer	95
10.3	Commercial Real Estate Before Democratization	96
10.4	REITs and the Democratization of Commercial Real Estate	96
10.4.1	Question & Answer	97
10.5	The Great Depression and the Redesign of the Home Mortgage	97
10.6	The Mathematics of the Fixed-Rate Amortizing Mortgage	98
10.6.1	Question & Answer	101
10.7	From Fixed-Rate Mortgages to Securitization, Crisis, and Reform	101

10.7.1 Question & Answer	102
10.8 Summary	103
11 Behavioral Finance and the Role of Psychology	104
11.1 Behavioral Finance and the Limits of Rationality	104
11.2 Morality, Exploitation, and Character	105
11.2.1 Question & Answer	105
11.2.2 Adam Smith and Praiseworthiness	105
11.2.3 Personality Differences and Manipulability	106
11.3 Prospect Theory: Setup and Board Architecture	106
11.4 The Value Function	107
11.4.1 Question & Answer	108
11.5 Weighting Probabilities and Rare Events	109
11.6 Other Behavioral Mechanisms	111
11.6.1 Regret and Gambling	111
11.6.2 Overconfidence: a Worked Example	111
11.6.3 Cognitive Dissonance	112
11.6.4 Question & Answer	113
11.6.5 Anchoring, Representativeness, and Social Contagion	113
11.7 Institutions, Shared Value, and Integrity	113
11.8 Summary	114
12 Misbehavior, Crises, Regulation and Self Regulation	115
12.1 Regulation After Behavioral Finance	115
12.2 Too Big to Fail, Microprudential, and Macroprudential Regulation	116
12.2.1 Question & Answer	116
12.3 Why Markets Ask for Referees	116
12.3.1 Question & Answer	117
12.4 Five Levels of Regulation	117
12.5 Within-Firm Regulation: Boards, Character, and Tunneling	118
12.5.1 Question & Answer	119
12.6 Trade Groups, Local Regulation, and National Disclosure	120
12.6.1 Question & Answer	120
12.7 From Dodd–Frank to Europe and the International Layer	122
12.8 Summary	124

13 Banks	125
13.1 What a Bank Is: Spread, Notes, and Liquidity	125
13.1.1 Question & Answer	126
13.2 Origins of Interest and Early Banking	126
13.3 Modern Bank Types and Scale	127
13.4 Diamond–Dybvig, Liquidity, and Multiple Equilibria	128
13.4.1 Question & Answer	129
13.5 Banks as Information and Monitoring Institutions	129
13.5.1 Question & Answer	129
13.6 Deposit Insurance in Practice: FDIC, FSLIC, Northern Rock, and IKB	130
13.6.1 Question & Answer	131
13.7 Basel Regulation and the Arithmetic of Capital Requirements	131
13.7.1 Question & Answer	133
13.8 Global Banking Crises, Contagion, and Shadow Banking	134
13.8.1 Summary	134
14 Guest Speaker Maurice “Hank” Greenberg	135
14.1 Opening Frame: Finance Beyond Formulas	135
14.2 From Soldier to Underwriter: Learning to Price Risk	135
14.2.1 Question & Answer	136
14.3 Turning Around American Home and Building AIG	136
14.4 Innovation, Diversification, and Opening Markets	137
14.4.1 Question & Answer	138
14.5 Organization, Incentives, and the Internal Logic of AIG	138
14.5.1 Question & Answer	139
14.6 AIG Financial Products, Risk Monitoring, and the Crisis Mechanism	140
14.6.1 Question & Answer	141
14.7 From AIG to the Wider Crisis: Diagnosis and Policy Lessons	142
14.7.1 Question & Answer	143
14.8 Summary	144

Chapter 1

Introduction and What This Course Will Do for You and Your Purposes

These notes follow Robert J. Shiller’s opening lecture on financial markets, prepared for this series with curation by LazyingArt LLC. The lecture does not begin with trading rules, valuation formulas, or market lore. It begins by enlarging the meaning of finance. We are asked to see finance as one of the basic coordinating mechanisms of a modern society: it allocates resources across time, organizes risky ventures, rewards contribution, and makes large plans possible. From that opening point the lecture unfolds in a deliberate sequence: a broad philosophical claim, an insistence on institutional detail, a restrained but real mathematical preview, and finally a moral question about what finance is for.

1.1 Finance as a Pillar of Civilized Society

Shiller begins by taking the title *Financial Markets* away from the narrow picture of trading and speculation. The word “markets” might suggest buying and selling securities, but the lecture immediately says that this is too small. Finance is presented instead as a pillar of civilized society. It is the structure through which large things are done.

Definition 1.1. In the sense of this lecture, finance is the set of institutions and arrangements through which society allocates scarce resources across time and uncertainty, sponsors collective ventures, incentivizes contribution, rewards constructive participation, and manages risk.

The opening definition already carries a mathematical flavor, even though the lecture does not formalize it. If we want a minimal symbolic picture, we may think of a social plan as specifying a state-contingent allocation $c_t(s)$, where t denotes time and s denotes a state of the world. Shiller does not develop a full equilibrium theory here, and we should not force one onto the lecture. The point is simpler and more foundational: finance concerns intertemporal and risky organization rather than mere exchange.

The lecture then identifies four functions that remain active throughout the course:

- allocation of resources across time and uncertainty,
- incentive design,

- sponsorship of ventures that require many participants,
- management of uncertainty and risk.

Risk enters at once. Anything large or important that we do is uncertain. That is why financial markets matter. The opening claim is therefore not simply that finance moves money around. It is that finance lets us act under uncertainty without being paralyzed by uncertainty.




Figure 1.1: An explanatory reconstruction of Shiller’s opening picture of finance as a social technology.

That large opening claim is not left hanging. Shiller’s next move is to say how such a subject must be studied.

1.2 Philosophy, Detail, and a World Perspective

Having made a broad civilizational claim, Shiller immediately narrows the method. The course will have a philosophical underpinning, but it will also be focused on details. This is one of the decisive sentences of the lecture. Finance is not to be discussed at the level of slogans. It is to be studied through institutions, contracts, and mechanisms, because, as he says, it is in the details that things happen.

The list of institutions arrives quickly: banking, insurance, securities, futures markets, derivatives markets, crises, and the future itself. The insistence on insurance is revealing. Some people treat insurance as if it were adjacent to finance rather than internal to it. The lecture refuses that separation. If finance is about risk-bearing and the organization of uncertain life, insurance belongs at the center.

The audience is then widened. The course has a U.S. bias because that is the institutional world Shiller knows best, but it is not meant to be provincially American. Many students will work outside the United States, and many listeners are not in the classroom at all. The Open Yale setting therefore matters intellectually, not just administratively. The lecture’s international perspective is motivated twice over: by the future careers of the students, and by the fact that the course itself is being offered to the world.

This global orientation is reinforced by the historical moment. The course has been updated because finance changes quickly, and the recent financial crisis has forced governments around the world to reconsider regulation and institutional design. Shiller makes a sharp comparison here: an undergraduate physics course may need only modest updating over a few years, but finance cannot

be treated that way. The crisis, the G20, and worldwide reform efforts are part of the subject itself.

The methodological principle is therefore clear. Large claims about finance must be tested and clarified at the level of institutions, and those institutions now operate in a global and rapidly changing setting. From there the lecture makes its next turn: if details matter, what role will mathematics play?

1.3 Mathematics, Diversification, and the Real World

Shiller places this course beside John Giannakopoulos's Econ 251. That other course is more theoretical and more mathematical. It includes state pricing, bond pricing, dynamic present value, uncertainty, hedging, the capital asset pricing model, and the leverage cycle. This course will not ignore mathematics, but it will use mathematics sparingly, intuitively, and always in the service of understanding the real world.

That restriction is not anti-mathematical. On the contrary, the lecture says mathematics is useful and cannot be avoided completely. Since many listeners will not take the more formal course, Shiller wants to provide the intuitive mathematical skeleton here. What he rejects is the idea that mathematics alone is the subject. The first mathematical spine of the course is already visible: independent risks, diversification, the law of large numbers, and the failure of the independence assumption in crisis.

Worked derivation. Let $X_1, \dots, X_n$ denote many separate risk exposures, and consider the equally weighted pooled exposure

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i. \quad (1.1)$$

A direct variance calculation gives

$$\text{Var}(\bar{X}_n) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \quad (1.2)$$

$$= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) + \frac{2}{n^2} \sum_{i < j} \text{Cov}(X_i, X_j). \quad (1.3)$$

If the risks are independent and each has variance σ^2 , then the covariance terms vanish and we obtain

$$\text{Var}(\bar{X}_n) = \frac{\sigma^2}{n}. \quad (1.4)$$

This is the simplest mathematical expression of the intuition Shiller emphasizes: if risks are independent, a large portfolio can diversify them away. In the same spirit, the law of large numbers may be written schematically as

$$\bar{X}_n \rightarrow \mathbb{E}[X] \quad \text{as } n \rightarrow \infty, \quad (1.5)$$

which is just a formal way of saying that many independent fluctuations tend to average out.

At this point the lecture makes its crucial turn. The attractive theory becomes dangerous when it is trusted too much. The financial crisis is narrated as a failure of the independence assumption. Portfolio theory had encouraged people to think that perhaps one investment might go bad, but not

all of them together. The crisis showed that this was false. In the language above, the neglected covariance terms return precisely in the bad states:

$$\text{Cov}(X_i, X_j) > 0 \quad \text{in crisis states} \quad \implies \quad \text{losses do not average out.} \quad (1.6)$$

This is all Shiller needs at this stage. He does not derive modern portfolio theory in full. He isolates its intuitive core, shows why it is attractive, and then shows why overconfidence in that core contributed to a worldwide crisis. The Great Depression enters here as the historical benchmark: by 1933 U.S. unemployment was about 25%, and the recent crisis is presented as a near miss rather than a trivial disturbance.

An important philosophical point follows. The crisis does not, in Shiller's view, refute finance. Airplanes can crash without refuting aviation. Likewise, financial institutions can be fragile without ceasing to be socially indispensable. Once mathematics has been put into that proper place, the lecture can ask a more personal question: what is this course actually for?

1.4 Why Everyone Needs Finance

The lecture now shifts from method to purpose. Shiller's answer is strikingly broad: this is not merely a course for the subset of students who will enter the finance industry. Almost everyone should know finance, because finance is woven into the structure of ordinary and extraordinary life.

He even says, with some deliberate self-consciousness, that this may be one of the most useful courses in Yale College. That claim is not based on salaries alone. It is based on action. Finance prepares people to do things in the world. The lecture reinforces the point with a personal image: when Shiller gives talks on Wall Street, he likes to ask for a show of hands from former students who are now there. The anecdote is light in tone, but its function is serious. The course is meant to enter real careers and real institutions.

This is where the lecture raises a local tension. Is the course vocational? Shiller resists both extremes. It is not vocational in the vulgar sense of narrow job training. Yet it is also not detached from life. Yale may be a liberal-arts institution, but that does not mean it should disdain practical knowledge about how large things are actually done. He even jokes that perhaps a hermit does not need finance. Everyone else does.

The key analogy is memorable: finance is a kind of engineering, but it is an engineering that works with people rather than machines. That is why institutional detail matters. A course that refuses detail would be like a language course that refuses vocabulary. If we are to act in the world, we need to know the words of finance.

This is also why the lecture lingers over textbooks, readings, and concrete institutional categories. Shiller says he wants students to become interested in details, and he compares learning finance to learning a language with many thousands of words. The point is not pedantry. The point is that one cannot make things happen in a modern economy while remaining illiterate in its institutional grammar.

The lecture reinforces this with an occupational comparison reconstructed from a Bureau of Labor Statistics chart. Shiller emphasizes the projected 2018 bars because those are the careers his students are approaching. The point of the chart is not contempt for other fields. The point is relevance. Finance is not a tiny niche. It is a large and expanding zone of actual work and institutional responsibility.

Occupation mentioned in the lecture	Approximate magnitude
Financial analysts	300,000
Financial managers	> 500,000
Personal financial advisors	250,000
Economists	20,000

Table 1.1: Approximate occupational magnitudes emphasized in the lecture, reconstructed from Shiller’s spoken Bureau of Labor Statistics comparison.

The deeper point is larger than employment statistics. Finance should be understood as a general-purpose technology for action. If we want to organize households, firms, charities, cities, universities, or long-term projects, we cannot remain mystified by the financial language in which those things are arranged. That broad claim of usefulness now gives way to a harder question, because finance is also associated with great fortunes.

1.5 Wealth, Scale, and Moral Purpose

At this point the lecture introduces another text: Shiller’s own unfinished manuscript, at that moment titled *Finance and the Good Society*. He notes that the title may sound like an oxymoron. That is precisely the problem. Finance is under moral suspicion, especially after the crisis. People see bailouts, lobbying, and enormous private rewards. Why should this be thought of as a noble profession at all?

The Forbes 400 discussion is the lecture’s answer-in-progress. It is not merely anecdotal color. It is the bridge from usefulness to moral tension. The point is that the richest people are typically not athletes or movie stars. Their wealth is usually tied to organizations, capital raising, ownership, acquisition, scaling, and deal-making. In other words, their power is institutional. They build and control arrangements through which many people work together.

This is why Shiller insists that finance is about making things happen on a large scale. It builds organizations. It marshals capital. It creates durable structures of coordination. The quiet people behind the scenes may matter more, in this institutional sense, than the visible celebrities.

The lecture then sharpens the question with arithmetic:

$$1 \text{ billion dollars} = 1000 \text{ million dollars.} \quad (1.7)$$

This elementary equality is used rhetorically, but it matters. Once one writes a billion dollars as a thousand million dollars, the ordinary consumer imagination begins to fail. One can buy cars and houses, but the scale quickly outruns ordinary consumption. The arithmetic forces the moral problem into view.

1.5.1 Question & Answer

Question. What should immense financial success be for, once private consumption is no longer the real constraint?

Answer. The lecture answers this in stages. First, it rejects conspicuous consumption as an adequate answer. Sumptuary laws in many societies are cited as evidence that open display of wealth has long been felt to be socially suspect. The problem is not only envy. It is that mere display seems too small for fortunes of such magnitude.

Second, the lecture turns to Andrew Carnegie's *Gospel of Wealth*. Carnegie's thesis is scandalous and interesting at the same time: those who are especially gifted at organizing business affairs have a moral obligation to redirect their fortunes toward public purposes rather than merely consume them or leave them to heirs. The fortune should not terminate in private luxury. It should be converted into institutions and public goods.

Shiller does not endorse this doctrine without qualification. He explicitly says there is only an element of truth in it. The strong claim that the rich are simply the most enlightened or deserving people is rejected. Yet the weaker claim is preserved: large-scale financial success becomes morally serious only when it is connected to purposes beyond personal display.

That is why Carnegie matters in the lecture. He is not introduced as a saint or as a theorem. He is introduced because he makes the problem vivid. If finance can create immense private wealth, then the moral question is what social form that wealth should eventually take. Carnegie's universities, halls, and public foundations are concrete answers, even if not a complete moral philosophy.

The lecture extends the same thought to the present, mentioning the campaign by Warren Buffett and Bill Gates to persuade billionaires to give most of their wealth away while still alive. At the same time, Shiller adds another qualification. There may even be a bubble in finance careers: as more people crowd toward a lucrative field, average outcomes may disappoint. Yet he still thinks that the top fortunes of the world will remain finance-related, precisely because the underlying technology scales so powerfully. The point is not that every financier is virtuous. The point is that finance, when it succeeds, creates a scale of action that immediately raises moral obligations.

From there the lecture moves back from moral philosophy into living institutions and living people.

1.6 Finance in Action: Speakers, Regulation, and Behavioral Finance

The outside speakers are not mere course logistics. They are embodiments of the lecture's thesis that finance is socially embedded and morally charged.

David Swensen is the clearest case. The lecture gives an explicit quantitative marker for the Yale endowment:

$$E_{1985} < \$1 \text{ billion}, \quad E_{2008} = \$22.9 \text{ billion}, \quad E_{2010.06} = \$16.7 \text{ billion}. \quad (1.8)$$

These numbers are used for more than admiration. They show the scale on which financial judgment can transform an institution. Swensen could have taken larger private compensation on Wall Street; the lecture presents his continued service at Yale as an example of finance directed toward a public purpose.

Maurice Hank Greenberg extends the point in a different register. Insurance, corporate power, philanthropy, and controversy are all brought together in one career. The lecture does not hide the scandal surrounding AIG, but it refuses to reduce Greenberg to the later collapse of a unit he no longer controlled. Again the pattern is the same: finance operates at a scale where institutional achievement and public scrutiny are inseparable.

Laura Cha supplies the missing regulatory pole. If finance involves incentives, risk, and the possibility of manipulation, then regulation is not external to finance. It is part of its moral and institutional architecture. The lecture is explicit on this point, and that is why a regulator belongs among the exemplary figures students are asked to hear.

The same expansion occurs in the teaching assistant discussion. Behavioral finance is introduced not as decorative psychology but as a change in how finance is understood. Mutual fund fee schedules, default choices, and consumer steering are cited as examples. A fund company may offer a menu of choices that looks benign on the surface but in practice steers clients toward arrangements that are not in their interest. This is exactly why finance cannot be treated as a purely mathematical subject. It is also about real people making choices under pressure, confusion, persuasion, competition, and imperfect information.

This movement from Swensen to Greenberg to Cha to behavioral finance completes a large arc of the lecture. Finance appears as portfolio management, insurance, regulation, and psychology all at once. Only now does Shiller open the subject into a full course roadmap.

1.7 Preview of the Mathematical and Institutional Spine

At the end of the lecture the outline of the course is given explicitly, but it no longer feels like administration. It feels like the consequence of everything that has already been argued.

Risk, crisis, and diversification. The next lectures begin with risk and financial crisis. Probability is mentioned, but only lightly. The mathematical core is exactly the one already isolated above: independent risks, diversification, law of large numbers, and the disaster that occurs when bad states produce common failure rather than cancellation. The capital asset pricing model is named as an important mathematical theory of diversification, but it is deferred rather than developed.

Technology, insurance, and efficient markets. Finance is then described as a technology. Securities, options, derivatives, and cash-flow partitions are treated as inventions rather than as arbitrary abstractions. Insurance becomes one of the central examples. It pools risk, prices uncertainty, and can improve safety by giving insurers an incentive to monitor the underlying environment. Efficient markets theory is previewed in the same restrained way: markets may often incorporate information rapidly, but the lecture refuses to turn that insight into a dogma.

Debt, equity, and contingent claims. Debt markets are introduced through the life cycle. A household may want a house before it has accumulated the cash to pay for one, so borrowing transfers purchasing power across time. Equity is introduced through a joint-enterprise example. Friends start a firm, but they contribute different amounts of capital, labor, and managerial effort. Shares are needed to divide claims and incentives:

$$\alpha_1 + \alpha_2 + \cdots + \alpha_m = 1. \quad (1.9)$$

This is not yet valuation theory. It is the simpler and more primitive problem of institutional design. The stock market, on this telling, is not primarily a casino. It is a mechanism for making collective enterprise stable and scalable. Options and other derivatives then appear as higher-order claims written on top of the underlying ownership structure.

Real estate, psychology, and banking. Real estate is flagged as both economically central and psychologically volatile. The recent crisis is linked to a bubble in home prices and to the collapse of that bubble. Behavioral finance will later return to the role of narratives, euphoria, and human judgment. Banking, multiple expansion of credit, and bank regulation are previewed as essential because the banking system itself nearly failed.

Later institutions. The remainder of the roadmap includes forwards and futures, options, monetary policy, investment banking, professional money management, exchanges, public finance, and nonprofit finance. The list is long, but the lecture's principle remains simple: each of these institutions is part of the machinery by which modern societies organize risk, capital, and purpose.

The final lecture of the course is said in advance to return to the question of purpose. This is the right ending because it is already the beginning. The entire roadmap is governed by the proposition that finance is not ultimately about making money. It is about giving shape to plans that would otherwise remain unrealized.

1.8 Summary

The first lecture does not yet teach a body of technical finance. It does something prior to that. It tells us what kind of subject finance is. It is a practical and moral technology of coordination under uncertainty. The mathematical content is real but deliberately modest: independence, averaging, diversification, and the failure of those ideas when correlations rise in crisis. Around that mathematical core the lecture builds a wider picture of institutions, incentives, regulation, wealth, philanthropy, psychology, and global change. The closing thought is therefore not an ornament. It is the governing thesis of the chapter: finance matters because it is one of the main ways in which we make our purposes happen.

Chapter 2

Risk and Financial Crises

We begin, as the lecture begins, with probability; but we are not asked to study probability in the abstract. Robert J. Shiller places it immediately inside the financial crisis since 2007 and uses that crisis to motivate the mathematics. This is only one way to think about the crisis, and not even necessarily his preferred way, but it is a revealing one. It lets us see how finance moves from stories about bubbles and failures toward models of returns, risks, co-movements, and extreme events. The thread of the lecture is deliberate: first crisis as a problem of uncertainty, then return as the basic financial object, then measures of central tendency and variability, then dependence, and finally the two breakdowns that matter most for crisis — failure of independence and fat-tailed shocks.

2.1 Probability, Narrative, and the Crisis

Shiller begins by saying that probability theory is fundamental to the way we think about finance, even though he does not take it as a prerequisite. That opening has a definite pedagogical shape. He does not start with a formal definition or a theorem. He starts with the crisis and asks us to see probability as one lens through which that crisis can be interpreted.

The first contrast is between narrative history and probabilistic modeling. The familiar narrative is quickly sketched: bubbles in stock and housing markets, the breaking of those bubbles, institutional failures, bank runs, emergency cooperation, bailouts, and eventual rebound. It is an intelligible story, and he does not reject it. Indeed, he says plainly that he likes the narrative story and will come back to it.

But then he pivots. Financial theorists often prefer to ask a different question. They do not look only at a few dramatic events. They ask how a very large number of small shocks accumulate. The crisis, in that frame, is not merely a sequence of memorable episodes; it is the visible outcome of many smaller disturbances whose cumulative behavior is governed by probabilistic laws. The stories remain, but they no longer carry the whole explanatory burden.

That is why the lecture broadens, for a moment, into a defense of probabilistic thinking itself. Probability, in its modern sense, is historically recent. That historical remark is not a curiosity. It is meant to underscore that thinking in terms of likelihoods and distributions was a genuine advance in human understanding. Once we think that way, we can model systems too complex to narrate event by event. Shiller's comparison is weather forecasting. We do not follow each molecule in the air, but we do build models for their aggregate motion. The probabilistic imagination in finance

aims at something similar.

The lecture also announces, already at this early stage, the two assumptions whose failure will matter later. One is independence. The other is the assumption of thin-tailed or non-outlier-prone distributions. The mathematics that follows is not detached technique. It is the machinery needed to say clearly what those assumptions mean and how their breakdown can help us interpret crisis.

2.2 Return and Measures of Central Tendency

With that broader motivation in place, the lecture turns to the first technical object. Finance begins with return, because finance begins with the question: if we invest over an interval of time, what happened to the value of the investment?

Definition 2.1. Let p_t be the asset price at the beginning of period t , let p_{t+1} be the price at the end of the period, and let d_t be the dividend paid during the period. The one-period simple return is

$$r_t = \frac{p_{t+1} - p_t + d_t}{p_t}. \quad (2.1)$$

The corresponding gross return is

$$1 + r_t. \quad (2.2)$$

The time index t may refer to a year, a month, or a day; the lecture often speaks as though we are thinking in monthly intervals. What matters is that return is attached to a period. We compare beginning-of-period price, end-of-period price, and any cash flow received in between.

Shiller immediately adds the limited-liability constraint. In the world he is assuming, one cannot lose more than the amount invested. Therefore

$$r_t \in [-1, \infty), \quad 1 + r_t \in [0, \infty). \quad (2.3)$$

This apparently simple bound becomes important later. Gross return is always nonnegative, and that is what makes the geometric mean economically meaningful in an investment setting.

He now moves to the first slide, titled *Expected Value, Mean, Average*. The structure of the slide itself is part of the lecture's rhythm: first expectation for a discrete random variable, then expectation for a continuous one, then sample mean, and finally geometric mean.

In cleaned form, while preserving the visible notation of the slide, the displayed equations are

$$E(x) = \mu_x = \sum_{i=1}^{\infty} \text{prob}(x = x_i) x_i, \quad (2.4)$$

$$E(x) = \mu_x = \int_{-\infty}^{\infty} f(x) x \, dx, \quad (2.5)$$

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i. \quad (2.6)$$

The slide visibly writes the discrete sum to ∞ , though the verbal explanation is more general: x is a random variable with countably many possible values. The point is that expectation is a probability-weighted average. For a continuous variable, the weighted sum becomes an integral against the density $f(x)$. For a sample, the estimate of expectation is the ordinary average $\bar{x}$.




Figure 2.1: Expected value, mean, and average. The lecture first emphasizes expectation as a measure of central tendency.

The lecture is careful about why these formulas are appearing. They are not introduced as an abstract statistics interlude. They are introduced because we want to evaluate an investor. If we observe n annual returns, the first and most obvious question is what that investor did on average. The arithmetic mean is therefore our first rough measure of success.

But the slide does not stop there. It places, in the lower-right corner, a second way of averaging:

$$G(x) = \left(\prod_{i=1}^n x_i \right)^{1/n}. \quad (2.7)$$

The layout matters. The lecture is about to make us feel the difference between adding and multiplying.

The slide's lower row can be read as a contrast:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i, \quad G(x) = \left(\prod_{i=1}^n x_i \right)^{1/n}. \quad (2.8)$$

Shiller's point is that for investment evaluation the geometric mean should be applied to gross returns, not to simple returns:

$$G(1+r) = \left(\prod_{i=1}^n (1+r_i) \right)^{1/n}. \quad (2.9)$$

This is the first important local payoff of the lecture. The arithmetic mean answers the question, "What was the average observed return?" The geometric mean answers the more financially relevant question, "What multiplicative growth rate is consistent with what happened to the portfolio through time?"




Figure 2.2: The same slide, now with the geometric mean explicitly indicated. This is the lecture's turn from average return to compounded investment performance.

2.2.1 Question & Answer

Question. Why is the arithmetic mean not enough when we evaluate an investment manager?

Answer. Because an investment portfolio compounds multiplicatively, and that means survival matters. A sequence of returns can have a moderate arithmetic average and still contain a wipeout. The arithmetic mean is therefore not a disciplined summary of compounded performance. The geometric mean, applied to gross returns, forces us to respect the possibility that one disastrous period can dominate everything that came before.

Worked example. Shiller's own example is extreme, and it is meant to be. Suppose a manager reports returns of 50%, then 30%, and then -100% . The arithmetic mean of the simple returns is

$$\bar{r} = \frac{0.5 + 0.3 - 1}{3} \approx -0.067. \quad (2.10)$$

That sounds merely bad. But the investment is not merely bad; it is dead. In gross-return form the same sequence is

$$1 + r_1 = 1.5, \quad 1 + r_2 = 1.3, \quad 1 + r_3 = 0. \quad (2.11)$$

So the geometric mean of gross returns is

$$G(1 + r) = (1.5 \cdot 1.3 \cdot 0)^{1/3} = 0. \quad (2.12)$$

This is precisely the discipline the lecture wants. If the portfolio is wiped out in one year, then whatever happened in the other years no longer rescues the compounded outcome. That is why Shiller recommends the geometric mean for investment evaluation.

At this point the lecture gathers expectation, arithmetic mean, and geometric mean under the common heading of *central tendency*. We now know how to talk about the center of a return distribution. But finance cares about more than the center. It cares, perhaps even more urgently, about the spread around that center.

2.3 Variance and the First Language of Risk

The lecture now makes a very natural move: from average outcome to variability. If the mean tells us what happens in the center, then risk asks what happens around the center.




Figure 2.3: Variance and standard deviation. The upper formula is introduced as the basic measure of variability.

The slide presents variance as

$$\text{var}(x) = \sum_{i=1}^n \text{prob}(x = x_i)(x_i - \mu_x)^2. \quad (2.13)$$

The idea is simple and fundamental. We first measure how far x is from its mean μ_x , then square the deviation so that positive and negative departures both contribute positively, and finally take the probability-weighted average of those squared deviations. Variance is therefore a measure of how far the random variable tends to wander from its center.

The standard deviation is then defined as the square root of variance:

$$\sigma_x = \sqrt{\text{var}(x)}. \quad (2.14)$$

The square root is what returns us to the units of the original variable. That is why standard deviation is easier to talk about in practice. If the mean return is 8% and a typical deviation from that mean is about 1%, then the standard deviation is naturally read as about 1%. Variance, by contrast, is measured in squared units.

A small numerical illustration helps fix the point. If we measure returns in percentage points and the return is often 1 percentage point above or below its mean, then the squared deviations are often near 1. In that unit system, the variance is near 1 and the standard deviation is again near 1. The precise value is not the point; the point is the role of squaring. Variance converts signed deviations into a positive measure of spread.

The lecture then moves, very explicitly, to the empirical estimator on the lower part of the same slide.




Figure 2.4: The sample-variance formula. The lecture now turns from population variability to an empirical estimator.

In cleaned form the estimator is

$$S_x^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2. \quad (2.15)$$

The symbol on the slide appears to be S_x^2 , and the lecture keeps the denominator n , noting only in passing that some people use $n - 1$. The order of thought is important: first the abstract variance of a random variable, then the practical formula one computes from observed returns.

So far, however, we have still been speaking about a single variable. The lecture's next question is unavoidable: financial risks are rarely isolated. How do two risky quantities move together? That question leads directly to covariance and correlation.

2.4 Covariance, Correlation, and Independence

Shiller introduces covariance with a stock-pair example. Let x be the return on IBM and y the return on General Motors. If IBM tends to be above its mean when GM is above its mean, then the two move together. If one is typically above its mean when the other is below, then they move in opposite directions.

A cautious reconstructed formula for the lecture's verbal definition is

$$\text{Cov}(x, y) = E[(x - \mu_x)(y - \mu_y)]. \quad (2.16)$$

The logic is straightforward. The product $(x - \mu_x)(y - \mu_y)$ is positive when the deviations have the same sign and negative when they have opposite signs. Averaging that product tells us whether co-movement tends to be positive, negative, or absent.

Correlation is then introduced as scaled covariance:

$$\text{Corr}(x, y) = \frac{\text{Cov}(x, y)}{\sigma_x \sigma_y}. \quad (2.17)$$

This scaling matters because it produces a dimensionless number between -1 and 1 :

$$-1 \leq \text{Corr}(x, y) \leq 1. \quad (2.18)$$

A correlation of 1 means the variables move exactly together, a correlation of -1 means they move exactly opposite one another, and a correlation of 0 means there is no linear tendency to move together.

Now comes the hidden assumption that lies beneath much of ordinary financial risk management: independence. In the lecture's simplified treatment, independence implies zero covariance, hence zero correlation. More importantly, it gives the basic aggregation formula:

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2 \text{Cov}(X, Y). \quad (2.19)$$

If X and Y are independent, then $\text{Cov}(X, Y) = 0$, and so

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y). \quad (2.20)$$

This is more than an algebraic identity. It is the mathematical reason diversification sounds reassuring. Some idea of unrelatedness, as Shiller puts it, underlies our ordinary thinking about risk.

2.4.1 Question & Answer

Question. Why did diversification, value at risk, and related risk-management ideas fail to protect institutions in the crisis?

Answer. Because the reassuring formulas were being used under an assumption of independence, or at least relative independence, that failed when it mattered most. If risks are genuinely unrelated, covariance terms disappear and aggregation becomes stabilizing. But if adverse shocks become connected, covariance rises exactly in the bad states. Positions that seemed diversified in normal times suddenly begin to fall together. The mathematics does not collapse on its own terms; rather, the dependence structure assumed by the model turns out to be wrong.

This is a decisive beat in the lecture. A notion that can sound technical and harmless — covariance near zero — becomes the key to understanding why apparently prudent institutions were far more vulnerable than they thought.

2.5 The Law of Large Numbers, VaR, and CoVaR

Once independence is on the table, the lecture turns to the promise it makes possible. If many shocks are independent, then averaging should reduce uncertainty. That is the intuition behind diversification, insurance, and a great deal of modern risk measurement.

Shiller frames this first through *value at risk*, or VaR, a concept firms emphasized after the 1987 stock-market crash. The lecture does not develop a full formal theory of VaR, but it gives its characteristic bottom line. A firm might say: there is a 5% probability that we will lose \$10 million in a year. In shorthand:

$$P(L \geq \$10 \text{ million in one year}) = 0.05. \quad (2.21)$$

That kind of statement requires a probability model. And, as the lecture insists, it quietly relies on assumptions about the relation among the shocks that generate losses.

The law of large numbers is then introduced through a coin-flip example. A single toss, with payoff +1 for heads and −1 for tails, has substantial uncertainty. But many independent tosses, averaged together, should display much less uncertainty. The lecture uses this as the intuitive bridge to finance and insurance.

The derivation can be written cleanly. Let $X_1, \dots, X_n$ be independent, identically distributed random variables with common variance σ^2 . Then

$$\text{Var}\left(\sum_{i=1}^n X_i\right) = \sum_{i=1}^n \text{Var}(X_i) + 2 \sum_{1 \leq i < j \leq n} \text{Cov}(X_i, X_j) \quad (2.22)$$

$$= n\sigma^2, \quad (2.23)$$

because independence makes every covariance term vanish. Dividing by n^2 , we obtain the variance of the average:

$$\text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{\sigma^2}{n}. \quad (2.24)$$

Hence the standard deviation of the average is

$$\text{sd}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{\sigma}{\sqrt{n}}. \quad (2.25)$$

As n grows, the uncertainty of the average shrinks toward zero.

This is the mathematical heart of the lecture's reassurance. It is also why Shiller invokes insurance. The idea that many independent risks can be pooled so that the average outcome becomes predictable is ancient in intuition, even if probability theory gave it a precise form only much later. Insurance relies on exactly that principle.

But the lecture does not let us linger in reassurance. It turns immediately toward the problem. The crisis suggests that the relevant shocks were not independent enough. That is why firms that computed VaR-style loss bounds turned out to be too optimistic. The model assumed that many things would not fail together; history showed that they could.

This is also the point at which Shiller introduces *CoVaR*. The name matters less than the adjustment in viewpoint. Once we recognize that portfolios may co-vary much more strongly in bad states than in ordinary ones, we need a risk measure that takes that conditional co-movement seriously. CoVaR is presented as a post-crisis response to exactly that problem.

2.6 Market Return, Beta, and Idiosyncratic Risk

Having developed the abstract language, the lecture turns back to market data. Shiller first looks at the aggregate stock market from 2000 onward and then overlays Apple on the same time period. The market line is already dramatic: large declines from 2000 to 2002 and again from 2007 to 2009. But once Apple is put on the same plot, the market must be visually compressed. Apple's behavior is vastly more extreme.

He pauses to explain that the Apple series is split-adjusted. This is not a change in underlying economics, only a correction for units. A stock split should not appear as a genuine collapse in value. That little institutional remark is characteristic of the lecture: even while developing abstract ideas, it keeps concrete market conventions in view.

The next move is crucial. Shiller replots the same underlying data as monthly returns rather than price levels. This changes the story. What looked, in one representation, like a long arc of extraordinary success now appears as a sequence of sharp ups and downs. The point is not merely graphical. The same data can look simple in levels and bewilderingly complex in returns. That is why probabilistic thinking is needed.

The human side of this issue appears in the Yale class-of-1954 story. A risky pool of \$375,000, given to Joe McNay, becomes \$90 million by the fiftieth reunion. Is that genius? Is it luck? The lecture never resolves the question completely, and that is the point. Historical success tempts us to narrate talent where probabilistic reasoning reminds us to consider realized luck and unobserved failures.

Shiller then turns from time series to a scatter plot of Apple returns against market returns. Each point is a month. The lecture's conceptual decomposition is that a stock return can be seen as the sum of a market component and an idiosyncratic component. In cautious reconstructed notation,

$$r_i = \beta_i r_m + \varepsilon_i, \quad (2.26)$$

where r_m is market return and ε_i is the idiosyncratic term.

Because no validated screenshot survives for this later plot, the following figure is an explanatory reconstruction rather than a reproduction.

The slope of the regression line is about 1.45, so

$$\beta_{\text{Apple}} \approx 1.45. \quad (2.27)$$

Shiller interprets this directly: Apple tends to move roughly one and a half times as much as the market. But he immediately adds the second half of the story. The scatter around the line is large. Apple still has substantial idiosyncratic risk, and that risk is tied to company-specific events — his vivid example is the rumor about Steve Jobs's health.

One detail from the lecture is especially revealing. Around the Lehman Brothers collapse, the broad market falls violently, yet Apple does not fall as much as one might have expected from market beta alone, because a separate company-specific news cycle is moving in the opposite direction. This is a concrete reminder that market and idiosyncratic components coexist rather than replacing one another.

Only after this empirical and institutional detour does the lecture turn to its final topic. So far we have worried that shocks may be too dependent. The last step is to worry that even the marginal distribution of shocks may be less tame than ordinary models suppose.




Figure 2.5: Explanatory reconstruction of the market and idiosyncratic decomposition. The fitted line represents the market component; deviations from the line represent idiosyncratic risk.

2.7 Normality, Outliers, and Fat Tails

Shiller now turns from dependence to distributional shape. Another common assumption in finance is that shocks are normally distributed. The lecture briefly reviews the attraction of the normal distribution: it is the familiar bell-shaped curve, the area under it is one, and different standard deviations change the scale without changing the general form.

The transcript becomes somewhat rough around the Mandelbrot discussion, but the analytic point is clear. If we look only at a limited sample of ordinary observations, a thin-tailed normal model can seem entirely plausible. Most days are not extreme days. The middle of the distribution can look orderly. The problem is in the tails.

That is where Shiller invokes Benoit Mandelbrot. The central claim is not that finance never looks normal in the middle. It is that rare, large events occur too often to be dismissed as negligible. In other words, the tails are heavier than the simple thin-tail picture suggests.

The lecture makes this vivid with market history. On October 30, 1929, the stock market rose 12.53% in a single day — the largest one-day increase in the historical sample he is discussing. But this was the rebound after the crash, not a calm market fluctuation. The days immediately before it had already produced enormous declines. The sequence itself makes independence look doubtful.

The second example is even more forceful. On October 19, 1987, the market fell 20.47% in one day on the S&P measure. If one fits a normal model to the central part of the historical distribution and asks for the probability of a decline this severe, the lecture reports a probability on the order of

$$P\left(\frac{\Delta p}{p} \leq -0.2047\right) \approx 10^{-71}. \quad (2.28)$$

The exact number is less important than the meaning. Under that fitted normal model, the event is effectively impossible. Yet it happened.

2.7.1 Question & Answer

Question. How can a financial event that a normal model treats as essentially impossible still appear in history?

Answer. Because the thin-tailed model is being fit to the center of the distribution and then trusted too far into the tails. In ordinary times the fit may look harmless. But the tails govern the probabilities of the crashes that matter most. If the true distribution has more mass in the tails than the normal model allows, then the model will systematically understate the frequency of extreme events. The event looks impossible only because the model is too thin-tailed, not because history has violated logic.

This is the lecture's second major breakdown. The first was failure of independence. The second is the failure of the thin-tail assumption. Put differently: diversification may fail because shocks become linked, and ordinary probabilistic reassurance may also fail because the shocks themselves are more extreme than the standard model admits.

The lecture closes by explicitly tying these ideas together. If independence held comfortably through time or across assets, then diversification would create safety. If thin-tailed normality held comfortably in the tails, then crises of the largest sort would be fantastically rare. But the crisis taught otherwise. Shocks can line up, and the tail behavior can be much more dangerous than the ordinary model suggests.

2.8 Summary

The lecture unfolds as a guided escalation. We begin with crisis narrative, then step behind the narrative into a probabilistic way of thinking about many small shocks. We introduce return because finance must first say what random variable it studies. We introduce expectation, arithmetic mean, and geometric mean because we need measures of central tendency; and we discover, through the wipeout example, that compounding forces us to care about multiplicative rather than merely additive averages. We then move to variance and standard deviation because finance is about risk, not just central outcome, and from there to covariance, correlation, and independence because risk in finance is never purely one-dimensional.

The deeper message of the lecture is that familiar mathematical comforts are conditional. The law of large numbers, diversification, insurance, and value at risk all lean on some form of independence. Normal-model reasoning leans on thin tails. The crisis reveals what happens when those assumptions are used too casually. Independence can fail, so shocks that appeared separate collapse together. Tails can be fat, so events that seemed impossible arrive with disturbing regularity. Probability therefore remains fundamental to finance, but not as a source of complacency. It is fundamental because it tells us exactly which assumptions have to be watched, and exactly how dangerous it is when they break.

Chapter 3

Technology and Invention in Finance

This lecture begins with a deliberate change of classroom style. Robert J. Shiller sets aside PowerPoint, takes up index cards and chalk, and makes finance feel less like canned presentation and more like worked engineering. That is the right entry point. The lecture is about invention in finance, but not invention in the thin sense of a clever formula. It is about devices that solve problems under uncertainty, devices that must preserve incentives, and devices that must also make sense to ordinary human beings. We therefore begin where Shiller begins: with a review of probability. Only after we see both the power and the limits of probabilistic reasoning do we move into the institutional inventions by which finance manages risk in practice. These notes follow Shiller's Yale lecture and are curated for this volume by LazyingArt LLC.

3.1 Finance as Engineering

Shiller's governing metaphor is announced almost immediately: finance is a form of engineering. That claim does several things at once. It tells us that finance is practical rather than merely doctrinal. It tells us that contracts are devices rather than static legal texts. And it tells us that small details matter. An airplane works only because thousands of parts and subassemblies fit together well enough to do something that at first seems implausible. Financial arrangements, in this lecture, should be understood in much the same way.

The lecture does not stop there. Engineering, Shiller reminds us, always has a human side. Devices are used by people, and people are imperfect. This is why engineering schools teach human factors engineering: one designs the machine so that ordinary users are less likely to misuse it. The move into psychology is therefore not a diversion from finance. It belongs to the inventive side of finance. If financial contracts are to help people pursue their purposes in life, they must be designed with actual human behavior in mind.

That is the opening stance of the whole lecture. Finance is not presented as a detached theory of prices. It is presented as a device-making discipline concerned with the risk problem, and in particular with the problem of maintaining incentives in the face of risks.

3.2 Probability Review and the Central Limit Theorem

Before turning to today's inventions, Shiller pauses for what is really a substantial reconstruction of the previous lecture. This is an important structural beat. He does not abandon theory and then move to institutions; he starts inside probability theory and carries its strengths and weaknesses forward into the new topic.

We begin with return. In the lecture, return is recalled as having two components, capital gain and dividend. In the standard notation used in companion notes, we may summarize the earlier derivation as

$$R_t = \frac{P_t - P_{t-1} + D_t}{P_{t-1}}. \quad (3.1)$$

The present lecture does not linger over the algebra. It uses the formula only to reopen the probabilistic vocabulary on which finance depends.

Definition 3.1 (Random variable). A random variable is a quantity created by an uncertain event or experiment: it is unknown in advance and becomes known later.

From there the review rebuilds the core list of objects: probability distributions, central tendency, the mean, the geometric mean, variance, correlation, regression, and finally the normal distribution. It is worth preserving that order, because it shows how the lecture narrows toward the central mathematical issue rather than jumping straight to it.

The regression example from the prior lecture is also brought back into view. Shiller recalls regressing Apple stock returns on market returns. In a compact notation one may write

$$\hat{R}_{\text{Apple},t} = \beta R_{M,t}, \quad (3.2)$$

$$R_{\text{Apple},t} = \beta R_{M,t} + \varepsilon_t, \quad (3.3)$$

$$\text{Corr}(\varepsilon_t, R_{M,t}) = 0. \quad (3.4)$$

The fitted value $\beta R_{M,t}$ is the market component of Apple risk, while ε_t is the idiosyncratic component. The lecture is careful to remind us that this residual is uncorrelated with the market component by construction. That caution matters. Finance often works with decompositions that are illuminating, but part of their neatness comes from how we have set up the model.

Shiller then returns to the intuition of independence. A coin toss appears independent of the toss before it. Daily stock returns may also look roughly independent if we think of them as responses to news, and news by definition has to be new. But this intuition is only partly right. The Apple example remains vivid because firm-specific news can still arrive in a way that breaks any lazy assumption of smoothness. The discussion of Steve Jobs's leave from Apple serves exactly this purpose: the world looks regular until a concrete event reminds us that the regularity was only approximate.

At this point the lecture leans hard on the central limit theorem. The blackboard frame is valuable evidence here because it shows the theorem heading and the assumption structure written out as abbreviated chalk reminders rather than as a full formal derivation.

The board supports the heading and the assumptions, especially the ideas of independent and identically distributed random variables and finite variance. The full formula is not visible, so what follows is best read as a cautious standard reconstruction of the theorem that Shiller is paraphrasing.




Figure 3.1: Central limit theorem assumptions

Theorem 3.2 (Central limit theorem). *Let $X_1, \dots, X_n$ be independent and identically distributed random variables with mean μ and finite variance σ^2 , and let*

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i.$$

Then

$$\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \Rightarrow N(0, 1)$$

as $n \rightarrow \infty$.

Remark 3.3. The blackboard image supports the theorem title and assumption list, not a fully written symbolic conclusion. The normalized display above is therefore a standard mathematical sharpening of the spoken statement rather than a verbatim board transcription.

The lecture-level intuition is clear and worth stating in the same order as Shiller states it. Averages are approximately normally distributed. The bell-shaped curve works as well as it does because so many things we observe are averages or sums of many effects. Large historical events, too, are often the joint outcome of many forces rather than the direct expression of one isolated cause.

The theorem also gives a clean expression of the pooling intuition that will matter for the rest of the lecture.

A worked derivation. If $X_1, \dots, X_n$ are independent and each has variance σ^2 , then the average $\bar{X}_n$ has variance

$$\text{Var}(\bar{X}_n) = \text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \quad (3.5)$$

$$= \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) \quad (3.6)$$

$$= \frac{1}{n^2} \left(\sum_{i=1}^n \text{Var}(X_i) + 2 \sum_{i < j} \text{Cov}(X_i, X_j) \right) \quad (3.7)$$

$$= \frac{1}{n^2} \left(n\sigma^2 + 2 \sum_{i < j} 0 \right) \quad (3.8)$$

$$= \frac{\sigma^2}{n}. \quad (3.9)$$

Under independence, the variance of the average shrinks like $1/n$. This is the clean mathematical form of the risk-pooling idea.

But the lecture does not leave the theorem in triumph. It immediately cuts back against it. The normal distribution has thin tails. The probability of extreme events never literally becomes zero, but the tails die away so quickly that, after a few standard deviations, the events look essentially impossible. That is where finance gets into trouble.

3.2.1 Question & Answer

If the bell curve is so powerful, why did the last crisis still defeat it?

Because the theorem is only as good as its assumptions. It assumes independence, and it assumes finite variance. If the underlying variables themselves are fat-tailed enough, or if apparent independence breaks down under stress, then the normal approximation ceases to be a trustworthy guide. This is why Shiller can say, in effect, that the central limit theorem is both important and in an important sense wrong for finance. It remains a fundamental piece of theory, but it does not command the entire field.

This is also where he corrects the intellectual genealogy of fat-tailed distributions, moving from Mandelbrot back to Paul Pierre Lévy. The correction is not incidental. It marks a larger habit of mind: one should know not only the theorem that makes the world look well-behaved, but also the tradition that explains why the world sometimes is not.

So much for the review. It gives us a mathematically disciplined way to think about pooling and averages, but it also shows why theory alone will not solve the risk problem. That is exactly the bridge into the rest of the lecture.

3.3 From 1970 to the Next Forty Years

Having restored both the power and the fragility of probabilistic reasoning, Shiller turns to invention proper. His first move is historical. Finance, he says, has changed astonishingly quickly.

The year 1970 serves as the benchmark. In 1970 there were no options exchanges, no financial futures, no swaps, and no electronic trading. There were options, but not exchange-traded options. Trading itself still relied heavily on voice, paper, and the physical floor of the exchange. This historical reminder has a clear purpose: it keeps us from treating today's financial menu as ancient, natural, or complete.

Shiller wants us to feel the distance between 1970 and the present because that distance opens a further question. If finance changed this much in forty years, what will the next forty years look like? This is where the lecture briefly becomes predictive. Technical progress, in his telling, has no obvious reason to stop; if anything, invention in finance is likely to continue and perhaps accelerate.

The lecture also injects a caution at just this point. Recent financial inventions have, in the crisis, "blown up" in front of the public. People therefore become suspicious of financial engineers in the same way they become suspicious of airplane designers after a crash or boiler-makers after an explosion. Shiller's answer is not that progress is harmless. It is that progress and regulation have to move together. Airplanes and boilers are regulated, not abolished. Finance, if it is indeed an engineering discipline, must be treated the same way.

3.4 Risk, Inequality, Framing, and the Slow Acceptance of Invention

At this point the lecture broadens sharply. The new question is not merely what instruments exist, but what finance is for. Shiller's answer is framed around risk and inequality.

He says, in effect, that perhaps the most important concept in finance is risk. A large part of human resentment toward inequality concerns arbitrary or gratuitous inequality, the kind produced by luck rather than effort or insight. Finance, at least in one of its nobler forms, is supposed to reduce the purely random component in people's lives. But Shiller is equally clear that finance also creates opportunity, and opportunity can itself widen inequality. That is why the lecture does not moralize simplistically. Finance can reduce arbitrary dispersion and yet increase the dispersion associated with differential success.

This social framing then leads to a more general point about risk-sharing. Financial inventions are one family within a much older set of human responses to uncertainty. Socialism, in Shiller's telling, can be read as one attempt to pool risks through common ownership and common life. Robert Owen's New Harmony becomes a striking but unsuccessful experiment in total sharing. Kibbutzim provide another example: they can work for some people, but they are not a general solution for an entire society. The American frontier custom of rebuilding a neighbor's burned house is an even more primitive form of insurance. It works only because the underlying risks are not perfectly correlated. If every house burned at once, the arrangement would be useless.

That is why Shiller says that mathematical theory motivates the discussion but is not the whole subject of the course. The old intuition remains the same: if risks are independent, we may pool them. But the actual problem is how to build devices that let that intuition operate reliably in society.

This is also the point at which he places finance beside tax and welfare institutions rather than in opposition to them. Progressive taxation, welfare, and public education are all treated as large-scale risk-management devices. They are among the things that make life tolerable in modern society. Finance is not meant to replace them. It is an add-on: a further set of inventions, often created by

private agents for their own purposes, which may nonetheless make the social world work better.

From here the lecture introduces its next pair of themes. One is framing. The other is devices. Framing matters because people do not spend their lives deriving utility-maximizing conclusions from first principles. They respond to associations, names, formats, and habits. Devices matter because finance, like engineering more broadly, builds structures with many moving parts for a definite purpose.

That is why the lecture dwells so long on nonfinancial inventions. The gimlet is not a digression. It is a reminder that useful tools can disappear from ordinary life even when they still solve certain problems better than newer devices do. The wheeled suitcase is not a joke. It illustrates how a useful invention may arrive late, then require a second refinement before it becomes truly compelling: Bernard Sadow's 1972 wheeled suitcase is followed only in 1991 by Robert Plath's roller board, which solves the problems of flopping and awkward handling. Even the red-cap objection matters, because it shows how a social expectation can block the obvious. The same is true of wheels in the pre-Columbian Americas and of subtitles in film. The point of all these examples is the same: invention is not just about logical possibility. It is about timing, complements, status, familiarity, and the slow acquisition of cultural legitimacy.

By the time Shiller turns back to finance, we are meant to have learned a general lesson. A financial innovation can be useful and still spread slowly. If it is not framed correctly, or if its surrounding institutions are missing, it may fail to take hold.

3.5 Limited Liability and Organizational Invention

The lecture now returns directly to financial invention, beginning with the corporation. The simple act of dividing a business into shares is old, almost elementary. We and our partners undertake a project, allocate shares according to contribution, and then each of us has an incentive to make the enterprise succeed. Shiller treats that as an ancient invention.

But he wants to focus on a later and much more decisive refinement: limited liability. In a compact modern notation, one may summarize the shareholder's position as

$$\Pi_{\text{shareholder}} \geq -I, \quad (3.10)$$

where I is the amount invested. The most the shareholder can lose is the capital already committed. The upside, however, is not capped in the same way.

That asymmetry is both legal and psychological. Legally, it means that if the company cannot satisfy its debts, creditors cannot simply continue through the firm and seize the shareholder's entire household. Psychologically, it changes the experience of investing.

The 1811 New York corporate law is the lecture's central case. Shiller emphasizes two distinct features. First, starting a corporation became comparatively easy; one filed papers rather than seeking a special legislative act. Second, the law made limited liability a clear right of the shareholder. Around the same time, Massachusetts moved in the opposite direction and kept shareholders exposed. That institutional contrast is the real drama of the episode.

In Massachusetts, open-ended liability made broad stockholding unattractive. A person might buy only one share and yet remain fearful that some future lawsuit against the company would come after his house and everything else he owned. In New York, by contrast, companies could proliferate and investors could participate without placing their entire lives on the line.

3.5.1 Question & Answer

Why did limited liability make New York financially dynamic while Massachusetts lagged?

Because it changed both the economics and the psychology of investment. Economically, limited liability bounded downside risk and thus made dispersed ownership feasible. Psychologically, it framed stockholding in a much more appealing way. Shiller, following David Moss, stresses exactly this point. Under limited liability, one pays once and knows that no further legal calamity can arrive from that decision. The contract begins to look like an attractive gamble: fixed downside, potentially very large upside. That is why Moss compares it to a lottery ticket. If the same share came with a small chance of unlimited ruin, it would no longer be entertaining and would no longer draw in capital so readily.

The lecture does not romanticize the result. Critics feared that easy incorporation and limited liability would produce irresponsible, fly-by-night companies. Many firms did fail. But enough firms succeeded that the overall system became transformative. New York attracted capital, innovation, and eventually national financial centrality. The broader lesson is exactly Shiller's: financial invention changes the economy when it solves both an incentive problem and a framing problem at the same time.

The next example, the Township and Village Enterprise in China, makes the same point in a different institutional key. Here the problem is not investor liability but weak legal enforcement in an economy emerging from strict communism. If a Chinese entrepreneur started to prosper privately, the surrounding village might appropriate the gains through taxes or direct interference. The solution was to design the enterprise so that the town itself was embedded in the organization. The firm was not purely private in the American sense; it was structurally tied to the local community. The lecture's numbers are striking: millions of such enterprises by the mid-1980s, and by the mid-1990s a dominant role in Chinese industrial production. The lesson is not that this form is universally superior. It is that invention in finance and organization is highly local. One invents around the actual obstacle one faces.

3.6 Inflation Indexation and the Chilean UF

The next problem is the instability of money. Most debts are written in nominal terms, that is, in units of currency. But inflation makes the purchasing power of currency uncertain through time, and long-term nominal contracts therefore carry inflation risk.

The natural response is indexation. In a standard and very simple reconstruction, an indexed payment rule takes the form

$$P_t = P_0 \frac{\text{CPI}_t}{\text{CPI}_0}. \quad (3.11)$$

If the price level rises, the nominal payment rises with it. The point is not mathematical sophistication. The point is to preserve real value.

Shiller's first historical example is the Massachusetts indexed bond of 1780. The setting is wartime inflation during the Revolutionary War. The state effectively defined a price basket, though not with the later language of the Consumer Price Index, and corrected bond payments according to changes in that basket. The lecture is careful on another detail here: the instrument was expressed in pounds, not dollars, and tied to the cost of actual commodities. In other words, it was already

grasping the essential problem. Money is a poor unit for a long horizon if its purchasing power is unstable.

Yet the invention disappeared. After the war, the United States largely forgot the indexed bond until 1997. That disappearance is one of the clearest places where the lecture's framing theme comes back into view. People find indexation conceptually awkward. They want money, not an unfamiliar formula. The fact that the instrument is useful does not guarantee adoption.

The 1997 reintroduction of indexed U.S. debt, associated in the lecture with Larry Summers, is therefore treated less as the arrival of a wholly new theory than as the difficult reappearance of an old and sensible invention. Even then, adoption remained partial. Indexed debt became a modest but not dominant fraction of U.S. public debt, and Shiller plainly thinks that is too timid.

Chile is the lecture's richer case because in Chile the invention did not remain confined to one security. It became a general social habit. The background is persistent inflation and repeated currency embarrassment. In the lecture's numerical sketch,

$$1 \text{ escudo} = 1000 \text{ old pesos}, \quad (3.12)$$

$$1 \text{ new peso} = 1000 \text{ escudos}, \quad (3.13)$$

so that

$$1 \text{ new peso} = 10^6 \text{ old pesos}. \quad (3.14)$$

The arithmetic is simple, but the institutional meaning is severe: the standard money unit has failed badly enough that it has to be repeatedly renamed and rescaled.

Chile's response, beginning in 1967, was the Unidad de Fomento, or UF. This is the important conceptual point: the UF is not merely an indexed bond. It is an indexed unit of account. Contracts are written in UFs and then settled in pesos at whatever the current peso value of the UF happens to be. In a schematic form, if a contract calls for q UFs at time t , then the payment in pesos is

$$\text{payment in pesos at time } t = q \times (\text{peso value of 1 UF at time } t). \quad (3.15)$$

The real promise remains stable even while the peso drifts.

The transcript is hesitant about the precise initial UF conversion in 1967, so one should not overstate it. But the later figures are clear enough for the lecture's purpose:

$$1 \text{ UF}_{1977} \approx 450 \text{ new pesos}, \quad 1 \text{ UF}_{\text{lecture date}} \approx 21,468 \text{ pesos}. \quad (3.16)$$

Thus the peso value attached to one UF rose by roughly

$$\frac{21,468}{450} \approx 47.7, \quad (3.17)$$

that is, about fiftyfold. A nominal peso contract would have been ravaged by that inflation. A UF contract would not.

The lecture is especially good here because it does not stop at financial securities. It asks how a society writes ordinary long-term promises. Rent, alimony, tuition, housing prices: these too should be stabilized against inflation if the underlying real obligation is supposed to remain constant. Chile, in Shiller's telling, made this natural. People there became accustomed to quoting obligations in UFs and then converting into pesos only at the moment of payment.

The contrast with the United States is revealing. Since 1913, the U.S. price level has risen by about a factor of twenty-two:

$$\frac{\text{CPI}_{\text{lecture date}}}{\text{CPI}_{1913}} \approx 22. \quad (3.18)$$

That is a dramatic cumulative change. Yet Americans still overwhelmingly contract in nominal dollars. Here again framing and habit explain persistence more than strict logic does. People think in money, so they write in money.

3.7 Swaps, Credit Default Swaps, and the Ambivalence of Innovation

The lecture's final examples are swaps and credit default swaps. These are placed at the end for a reason. They exhibit the full double edge of financial invention: genuine increases in the ability to manage risk, combined with the creation of new and poorly understood risks.

A swap is introduced simply as a long-term contract to exchange cash flows. In the lecture's simplest example, one side has dollar income and wants euros, while the other has euro income and wants dollars. A very bare schematic is enough:

$$\text{Party A} \rightarrow \text{Party B} : \text{dollar cash flows over time}, \quad (3.19)$$

$$\text{Party B} \rightarrow \text{Party A} : \text{euro cash flows over time}. \quad (3.20)$$

The swap rate is fixed in advance. The contract therefore converts uncertain future exchange opportunities into a planned exchange rule. One may denote that preset conversion abstractly by a swap rate k , but the lecture does not develop pricing formulas; it stays with cash-flow structure and use.

Shiller attributes the invention of the swap, in this lecture, to David Swensen during his Wall Street period in the early 1980s. Whatever one concludes from broader historical sources, the lecture's analytic emphasis is clear. The swap looks obvious in retrospect, but it required a legal environment in which such contracts could be documented and enforced. That is why ISDA enters the story. The International Swaps and Derivatives Association represents the institutional side of invention: lobbying for legal clarity, standardizing documentation, and helping make the device usable at scale.

The credit default swap then appears as a more specialized contract. One side, the protection buyer, makes regular payments; the other side, the protection seller, owes money if a specified credit event occurs. In the most compact form,

$$\text{protection buyer} \rightarrow \text{protection seller} : \text{regular premium payments}, \quad (3.21)$$

$$\text{protection seller} \rightarrow \text{protection buyer} : \text{payment if a credit event occurs}. \quad (3.22)$$

The event may be bankruptcy or some other contractually defined default-type event. The lecture insists on the obvious economic intuition: the buyer is purchasing protection against credit risk.

3.7.1 Question & Answer

Isn't a credit default swap just insurance?

Economically it looks very much like insurance, and Shiller explicitly stages that question. The answer is that the surrounding institutional world is different. Credit insurance had existed since the nineteenth century, but it grew inside a different regulatory and conceptual framework. Traditional credit insurers were often more limited in scale and more involved in monitoring or advising their clients. Credit default swaps, by contrast, were built inside the derivatives infrastructure, supported by different documentation practices and by organizations such as ISDA. The difference is therefore not merely in the payoff pattern. It lies in legal framing, market culture, and scalability.

This is the right place for the lecture to close, because it returns us to the crisis. AIG's near-failure is presented as a failure connected to credit default swaps, but not as a final indictment of innovation itself. The larger point is more subtle. Swaps and CDS contracts increased the financial system's capacity to manage risk. They were important inventions. But they also created new chains of exposure that were not well enough understood by firms, regulators, or governments. Innovation moved faster than comprehension.

3.8 Summary

The lecture unfolds in a deliberate arc. We begin with probability because the mathematical intuition of independence and averaging lies underneath modern finance. We then learn why the central limit theorem is powerful and why it fails when tails are fat or dependence reappears under stress. That failure is not the end of finance; it is the reason finance becomes inventive. The lecture then shows invention at several levels: the social level of risk-sharing and inequality, the psychological level of framing, the institutional level of incorporation and limited liability, the monetary level of indexation and the UF, and the contractual level of swaps and credit default swaps. The final balance is unmistakable. Innovation is necessary and often beneficial, but it is never innocent. Finance remains an engineering discipline, and engineering requires design, human factors, and regulation all at once.

Chapter 4

Portfolio Diversification and Supporting Financial Institutions

In Robert J. Shiller's fourth lecture on financial markets, as curated by LazyingArt LLC, the mathematics is deliberately delayed. We begin with a historical invention, not with an optimization problem. The VOC, the Amsterdam Stock Exchange, brokers, street-name ownership, short selling, and limited liability all appear before any frontier is drawn. Only after that world is in place do we ask the lecture's central quantitative question: if expected returns, variances, and covariances are given, how should we think about a portfolio? The answer unfolds step by step: first leverage, then hyperbolas, then diversification, then tangency, and finally the CAPM intuition that the risk which matters is covariance with the market.

4.1 Finance as Institutional Invention

The lecture opens by returning to the previous lecture's theme that finance is a branch of engineering. Financial arrangements are invented to perform functions; when they work, they spread. That is the spirit in which the lecture turns to 1602 and to the Verenigde Oostindische Compagnie, together with the Amsterdam Stock Exchange built to trade its shares.

Definition 4.1. A portfolio is a collection of investments held over a specified horizon.

This historical prelude matters because the lecture does not want to treat portfolio theory as a detached theorem. The VOC was not a one-voyage merchant partnership that dissolved when the ships came home. It was a long-term company, meant to continue indefinitely, to build trading posts across the world, and to attract capital from a broad public. Once such a company exists, the lecture argues, a whole institutional framework begins to arise almost automatically.

Shares become tradable every day rather than only when the company reopens its books. Brokers hold inventories and intermediate between buyers and sellers. Ownership is effectively held in what we would now call street name: the broker's records tell the investor that he owns the shares, even before the company itself updates its own books. And once that arrangement exists, a broker can sell more shares than he currently owns and plan to acquire them later. Short interest is not an exotic later invention. It appears naturally once the market is liquid enough.

The lecture lingers over this point because it is one of the main institutional morals. A large, durable,

tradable company does not produce calm. It produces disagreement, speculation, and volatility. Some traders become wildly optimistic and bid the stock up; others become pessimistic and short it. The market becomes active because the value is uncertain, and the value is felt to be uncertain because the market is active.

The same is true of limited liability. If shareholders can lose no more than the funds they invest, then a much wider public can take part in risky enterprise. One need not inspect every ship captain or every distant trading post personally. One can buy the stock and know that the downside is capped at the original investment. In that sense the stock market is a social technology. It channels appetite for risk into the financing of large productive ventures.

This is why the opening history is not decorative. The lecture is showing us the institutional soil out of which portfolio theory grows. Once tradable risky claims exist and once they fluctuate in public view, the question of how to hold them becomes unavoidable.

4.2 The Equity Premium Puzzle

From that institutional prelude the lecture pivots into a puzzle. One can already hear it in the VOC story. Suppose an equity claim appears to earn astonishing returns year after year. Why does capital not flood into it until the extraordinary advantage disappears?

The lecture broadens that question beyond the VOC itself. The striking point is not merely that one famous stock once did well. The point is that common stocks seem to have done very well over very long periods. In the U.S. data cited in the lecture, the real geometric average return on stocks is about 6.8% per year, whereas the real return on short-term government securities is only about 2.8%. The difference is therefore about

$$6.8\% - 2.8\% = 4.0\%. \quad (4.1)$$

That is already a large gap. The lecture then adds that in the long U.S. sample under discussion there was no rolling thirty-year period in which stocks underperformed safe short-term government claims or long-term bonds.

The same broad phenomenon appears internationally. The lecture cites evidence from Dimson, Marsh, and Staunton: across the countries they study, stocks outperform bonds, with equity premia running from about 3% to about 6%. The puzzle is therefore not parochial. It is not just a peculiarity of the United States.

This is where the lecture deliberately slows down. If equities have performed so well for so long, why does everyone not simply learn the lesson and hold them?

4.2.1 Question & Answer

Question. If stocks do so well, why does everyone not simply hold them?

Answer. The lecture's first answer is the standard one: risk. Stocks are riskier; the extra return is a risk premium. But the lecture does not allow that answer to stand in vague form. What exactly is the relevant risk? Is it just price fluctuation taken by itself? Or is the matter subtler once assets are placed inside a portfolio? That is the point of transition. The lecture is not satisfied with the

slogan that stocks are volatile. It wants a theory of why that volatility matters and of how portfolio structure changes the answer.

At this point the lecture introduces Harry Markowitz. The move is methodological. Before 1952, investors knew the proverb about not putting all their eggs in one basket. But the lecture insists that this was not yet a theory. It was good advice without a mathematical problem attached to it. Markowitz changed the subject by asking: if we know expected returns, variances, and covariances, what portfolio should we choose?

4.3 Markowitz's Breakthrough and the One-Risky-One-Riskless Case

The lecture treats Markowitz's 1952 paper as a genuine conceptual break. In retrospect the problem sounds obvious. But that is part of the point: it sounded obvious only after someone finally asked it sharply. What had existed before was a loose culture of investing and advice literature. What Markowitz supplied was a well-posed problem in which the inputs are expected returns, variances, and covariances, and the output is a portfolio choice.

The lecture dramatizes this by imagining a mathematically inclined portfolio manager. We are given a one-year horizon. We gather historical data on every possible investment we can think of. We compute average returns, variances, covariances, and correlations. Then we step back from judgment and prediction and ask the stripped-down mathematical question: if these statistics are taken as given, what is the best portfolio? The lecture's astonishment is that such a question was not standard long before 1952.

A small historical aside reinforces the point. The lecture notes that older investment manuals did say "do not put all your eggs in one basket." But they stopped there. They did not say how to diversify, how much to hold, or why a particular mix is optimal. That is exactly the gap Markowitz filled.

The lecture begins the formal analysis with the simplest nontrivial case: one risky asset and one riskless asset. The risky asset may first be imagined as the VOC itself. The riskless asset may be taken, approximately, to be a short government claim whose maturity matches the investor's horizon.

4.3.1 Question & Answer

Question. If leverage lets us reach many target returns from one risky asset, what could an "optimal portfolio" even mean?

Answer. It cannot mean "the best asset" in isolation. Once borrowing, lending, and shorting are allowed, a single risky asset already generates a whole menu of feasible portfolios. So the right question is not "What is the best investment?" but rather "What tradeoff between expected return and risk do we want?" That is the first genuinely Markowitzian turn in the lecture.

The slide visible at this point in the lecture is worth keeping because its order of presentation is itself instructive: first the allocation, then expected return, then variance, then standard deviation, and only then the reduced relation between standard deviation and expected return. The slide says




Figure 4.1: Lecture slide: risky and riskless asset portfolio formulas.

“portfolio expected value,” but the lecture is plainly using expected return. We standardize the notation while preserving the lecture’s sequence.

If a fraction x of wealth is placed in the risky asset and $1 - x$ in the riskless asset, then

$$r_p = xr_1 + (1 - x)r_f, \quad (4.2)$$

$$\text{Var}(r_p) = x^2 \text{Var}(r_1), \quad (4.3)$$

$$\sigma_p = |x| \sigma_1. \quad (4.4)$$

The variance expression is simple because the riskless return is known at the horizon. The lecture then solves the expected-return equation for x :

$$x = \frac{r_p - r_f}{r_1 - r_f}. \quad (4.5)$$

Substituting this into the expression for standard deviation gives the cleaned version of the bottom slide formula:

$$\sigma_p = \left| \frac{r_p - r_f}{r_1 - r_f} \right| \sigma_1. \quad (4.6)$$

This is the lecture’s reduced formula for the broken straight line in (σ_p, r_p) -space.

Worked example. The lecture’s numerical example chooses a riskless rate of 5%, a risky-asset expected return of 20%, and a risky-asset standard deviation of 40%. It begins not with an abstract weight x , but with guilders.

1. Suppose we have 100 guilders and put them all into the risky asset. Then the expected profit is 20 guilders and the standard deviation is 40 guilders.
2. Suppose instead we borrow another 100 guilders and invest 200 guilders in the risky asset. Then the asset side has expected gain 40 guilders, but the loan costs 5 guilders in interest. So

the expected return on our original wealth is

$$2(20\%) - 5\% = 35\%. \quad (4.7)$$

The standard deviation doubles with the position size:

$$2(40\%) = 80\%. \quad (4.8)$$

3. Now go the other way and short 200 guilders of the risky asset. The short position has expected loss 40 guilders, while the 300 guilders held in cash earn 5%, or 15 guilders. So the expected return on the original wealth is

$$-2(20\%) + 3(5\%) = -25\%. \quad (4.9)$$

The lecture then turns the bookkeeping into geometry. If we plot standard deviation on the horizontal axis and expected return on the vertical axis, the feasible set becomes a broken straight line issuing from the riskless point: an upward branch for lending and leveraged long positions, and a downward branch for sufficiently large short positions. The lecture explicitly calls this a degenerate hyperbola.




Figure 4.2: Transcript-based reconstruction of the one-risky-one-riskless opportunity set.

This is where the lecture makes its first big philosophical move. If a client asks for 100% expected return, one can produce it by leverage. That does not make the portfolio manager a genius. It only means that the menu is already wide. There is no single best investment here, only a tradeoff between return and standard deviation. That is why the lecture says that before we have multiple risky assets, we have not yet fully gotten into Markowitz.

4.4 Two Risky Assets and the Efficient Frontier

Now the lecture leaves the degenerate case behind and moves to the first true Markowitz curve. We temporarily set the riskless asset aside and imagine two risky assets. In the numerical example used




Figure 4.3: Lecture slide: two-asset frontier variance substitution.

in the lecture, these are U.S. stocks and long-term Treasury bonds held over a one-year horizon. The latter are called “bonds,” but the lecture is careful to say that they are not riskless here, because their market price changes over the year.

If x_1 is the weight in asset 1 and $1 - x_1$ is the weight in asset 2, then the portfolio expected return is

$$r_p = x_1 r_1 + (1 - x_1) r_2. \quad (4.10)$$

The lecture then repeats the same logical move as before: solve for the weight in terms of the target return. The slide shows

$$x_1 = \frac{r_p - r_2}{r_1 - r_2}. \quad (4.11)$$

The underlying two-asset portfolio variance formula, standard in this setting, is

$$\sigma_p^2 = x_1^2 \sigma_1^2 + (1 - x_1)^2 \sigma_2^2 + 2x_1(1 - x_1) \sigma_{12}. \quad (4.12)$$

Substituting the solved expression for x_1 produces the visible frontier formula:

$$\sigma_p^2 = \left(\frac{r_p - r_2}{r_1 - r_2} \right)^2 \sigma_1^2 + \left(\frac{r_1 - r_p}{r_1 - r_2} \right)^2 \sigma_2^2 + 2 \frac{(r_p - r_2)(r_1 - r_p)}{(r_1 - r_2)^2} \sigma_{12}. \quad (4.13)$$

This is the lecture’s first genuine hyperbola. The covariance term is not an afterthought. It is the reason the curve bends inward. If the two assets do not move perfectly together, then combining them can reduce risk in a way that no single-asset comparison could have revealed.

The lecture now asks, in its own voice, “What do we learn from this?” The first answer is that the whole curve is not efficient. Only the upper branch above the minimum-variance point belongs to the efficient frontier. The second answer is much sharper: a 100% bond portfolio is dominated. There exists a stock-bond mixture with the same standard deviation and a higher expected return.

This is one of the lecture’s most memorable practical claims. A conservative portfolio is not automatically rational just because it carries a familiar label such as “all bonds.” Once the geometry




Figure 4.4: Simplified redraw of the two-asset opportunity set. The efficient frontier is the upper branch above the minimum-variance point.

is drawn, we can see that some seemingly prudent choices lie strictly below the efficient set. The lecture even remarks that this now looks obvious and yet once was not. That is why it treats Markowitz as such a fundamental breakthrough.

The lecture’s fondness for geometry is not ornamental here. It briefly invokes Apollonius and conic sections to make the point that the hyperbola is doing real conceptual work. The curve is teaching us something about investment choice that was not obvious before we drew it.

4.5 Three Assets, Diversification, and Oil

Having gone from one risky asset to two, the lecture now goes one step further and adds oil. The transition is important: the two-asset case gives us the frontier geometry, but the three-asset case gives us the lecture’s main application of diversification.

Remark 4.2. At this point the transcript garbles some of the third-asset subscripts. The formulas below are standard reconstructions that match the lecturer’s intended meaning and the visible frontier chart. They should be read as careful normalization rather than as a literal transcription of every spoken symbol.

With three risky assets, expected return remains a weighted average:

$$r_p = x_1 r_1 + x_2 r_2 + x_3 r_3, \quad x_1 + x_2 + x_3 = 1. \quad (4.14)$$

The portfolio variance becomes

$$\sigma_p^2 = x_1^2 \sigma_1^2 + x_2^2 \sigma_2^2 + x_3^2 \sigma_3^2 + 2x_1 x_2 \sigma_{12} + 2x_1 x_3 \sigma_{13} + 2x_2 x_3 \sigma_{23}. \quad (4.15)$$

The lecture’s chart should then be read as follows: for each target expected return, solve for the minimum-variance mixture of stocks, bonds, and oil. The resulting set of minimizing portfolios traces a new frontier.

The slide is visually dense and worth keeping as evidence. The old stock-bond frontier remains on it as the weaker comparison set, labeled “W/O Oil.” The new frontier, labeled “With Oil,” lies to the




Figure 4.5: Lecture slide: efficient frontier with and without oil.

left over the relevant range. That leftward shift is the lecture's main diversification result: once we add an asset that offers return and is not highly correlated with the stock market, we can attain the same expected return with less standard deviation.

The lecture likes to say this in deliberately ordinary language: more eggs in the basket, not fewer. That is the inversion of the old proverb. The point is not to spread wealth thinly at random. The point is that, once covariances enter the picture, more asset classes can make the whole portfolio safer for a given target return.

The visible labels on the slide support the narrative. We can still see benchmark portfolios such as 100% bonds, 50% stocks and 50% bonds, and 100% stocks. But once oil is available, the frontier includes mixtures such as 15% oil, 53% stocks, 32% bonds, and 21% oil, 79% stocks, 0% bonds. At the upper end one labeled point even uses a negative bond weight, reminding us that the frontier can combine diversification with leverage.

The lecture then turns back from mathematics to institutions. Norway is described as having far too much implicit oil exposure, not because its statisticians cannot do the arithmetic, but because oil is politically and symbolically charged. The lecturer's argument is that Norway need not literally sell the oil in the ground in order to diversify; it could reduce exposure with derivative transactions, for example by shorting oil futures. Yet politics makes such hedging difficult. Mexico appears as a similar example. The lesson is that correct geometry does not automatically become actual policy.

The lecture also takes care to sharpen the frontier concept once more. If we ask for the minimum-variance portfolio that delivers, say, a 9% return, we may land at 100% bonds. But that point is still not efficient if a higher-return point exists at the same standard deviation. So the efficient frontier is only the upper branch above the minimum-variance point. Minimum variance by itself is not the same thing as best.

This second screenshot of the oil chart matters for a different reason. Here the lecturer points directly to the left end of the horizontal axis and remarks that the diagram does not show zero.




Figure 4.6: Simplified redraw of the frontier comparison, preserving the lecture's truncated horizontal axis beginning at 5%.

That observation is not incidental. It prepares the next step. Once we add the riskless asset back into the picture, the relevant line begins at zero standard deviation on the vertical axis at the riskless return. The oil chart is a picture of risky portfolios only; it is not yet the full geometry.

4.6 The Riskless Asset Returns: Tangency, Mutual Funds, and Market Risk

Now the lecture returns to the riskless asset. The move is deliberately parallel to the earlier one-risky-one-riskless example. We take any efficient risky portfolio on the frontier and treat it as a single composite risky asset. Then we ask what happens when we mix that composite with the riskless asset.

Geometrically the answer is simple. Any mixture of a given risky portfolio and the riskless asset lies on a straight line. So the problem becomes: which straight line through the riskless point lies highest? The lecture phrases it in exactly that way. Start from the riskless return on the vertical axis and draw the highest feasible straight line. The portfolio at which that line touches the efficient risky frontier is the tangency portfolio.

If r_T and σ_T denote the expected return and standard deviation of that tangency portfolio, then every mixture of the riskless asset and that portfolio lies on the capital market line

$$E[r_p] = r_f + \frac{E[r_T] - r_f}{\sigma_T} \sigma_p. \quad (4.16)$$

The lecture does not derive this equation algebraically, but it does derive the geometry that the equation expresses.

This restores, in a subtler form, the idea of an optimal portfolio. Earlier the lecture argued that there is no single best investment in the crude sense. With only one risky asset and a riskless asset, we had a whole broken line of possibilities. With many risky assets and no riskless asset, we had a hyperbolic frontier of efficient but taste-dependent choices. But once the riskless asset is available again, a special risky portfolio emerges: the tangency portfolio. Everyone who is efficient wants




Figure 4.7: Lecture slide: the horizontal axis begins at 5%, not at zero.

to lie on the same line and therefore wants to hold the same risky fund, scaled up or down by borrowing or lending.

Proposition 4.3. *Under the lecture's Markowitz assumptions—shared expected returns, shared variances and covariances, and the ability to borrow or lend at the riskless rate—every efficient portfolio is a mixture of the riskless asset and the same tangency portfolio.*

Proof. Fix any efficient risky portfolio on the frontier. Mixing it with the riskless asset produces a straight line through the riskless point. Among all such lines, only the one with maximal slope delivers the highest expected return for every attainable standard deviation. That maximal-slope line is tangent to the risky frontier. Therefore any efficient investor chooses a point on that same line and differs from any other efficient investor only in the scale of the holding in the riskless asset versus the tangency portfolio. $\square$

This is the mutual fund theorem in the lecture's form. One need not search endlessly among fundamentally different efficient risky funds. Under the theory's assumptions, there is one. Investors differ not in the identity of the risky portfolio, but in how much leverage or de-leverage they want around it.

The lecture then takes the market-clearing step. If all investors want to hold the same risky portfolio, that portfolio must be the aggregate portfolio of risky assets outstanding. Otherwise supply and demand would not match. Hence

$$\text{market portfolio} = \text{tangency portfolio}. \quad (4.17)$$

This is the bridge from Markowitz geometry to equilibrium asset pricing.

The lecture does not give a full derivation of the capital asset pricing model. It states the canonical relation and then interprets it:

$$E[r_i] = r_f + \beta_i(E[r_M] - r_f), \quad (4.18)$$

with

$$\beta_i = \frac{\text{Cov}(r_i, r_M)}{\text{Var}(r_M)}. \quad (4.19)$$

The central interpretive point is that investors do not care about standalone variance once an asset can be embedded inside a diversified portfolio. Idiosyncratic fluctuations can be averaged away. What remains costly is covariance with the market, because that risk survives diversification. In the lecture's own language, risk is no longer mere uncertainty. Risk is co-movement with the market.

That is why beta matters. A high-beta stock is not merely a stock with a jumpy price. It is a stock whose return moves strongly with the market's return. The lecture gives Apple, with beta around 1.5, as an example of an asset that responds to market movements in an amplified way. That market exposure is what commands extra expected return.

The final practical coda is the Sharpe ratio:

$$S_p = \frac{E[r_p] - r_f}{\sigma_p}. \quad (4.20)$$

Along the tangency line, this ratio is constant. Geometrically it is just the slope of the line from the riskless point to the portfolio. Economically it is a correction for leverage. A manager who advertises a 15% average return may simply be leveraging up an ordinary risky strategy. The Sharpe ratio asks instead how much excess return was earned per unit of standard deviation.

This is how the lecture ends: with a theory that begins by redefining risk and closes by telling us how to evaluate portfolio managers. The mathematics is not an ornament. It is a way of refusing to be impressed by raw return without asking what leverage and what market exposure produced it.

4.7 Summary

The lecture begins with institutions and ends with equilibrium asset pricing, and the order is the argument. We first see finance as invention: the VOC, the Amsterdam exchange, brokers, street-name ownership, short interest, limited liability, and volatility. That institutional world then raises the equity premium puzzle. Markowitz enters by turning a proverb about diversification into a mathematical problem stated in terms of expected returns, variances, and covariances. With one risky asset and one riskless asset, we discover that there is no single best investment, only a risk-return tradeoff. With two risky assets, the feasible set becomes a hyperbola and dominated portfolios become visible. With three risky assets, diversification becomes a geometry of leftward shifts in the frontier, even if institutions and politics resist the prescription. Finally, when the riskless asset is added back, the tangency portfolio, the mutual fund theorem, the market portfolio, the CAPM intuition, and the Sharpe ratio all fall into place. The lecture's deeper lesson is that risk is not a property of an asset in isolation. It is a property of the asset inside the portfolio and, in the end, inside the market.

Chapter 5

Insurance, the Archetypal Risk Management Institution, its Opportunities and Vulnerabilities

These notes follow Robert J. Shiller’s fifth lecture in the *Financial Markets* course and are curated by LazyingArt LLC. The lecture begins with a brief conversational aside about Davos and then pivots immediately to its real theme: insurance is one of the central institutions of risk management, even if history and regulation have taught us to place it outside finance. Shiller’s route is deliberate. We begin with the intuition of risk pooling, move back through early forms of insurance and forward into probability, then turn from theorem to institution: contracts, incentives, data, organizational form, regulation, AIG, health insurance, catastrophe risk, and the unfinished work of making insurance serve human welfare.

5.1 Insurance as a Risk-Management Institution

We should hear this lecture as a continuation of the preceding one, not as a detour. In the previous lecture, risk management appeared through mean–variance analysis and the capital asset pricing model. Here the same underlying problem reappears in a different institutional language. Instead of portfolios and traded securities, we now speak of policies, claims, reserves, and regulation. But the objective is unchanged: people want protection against outcomes that would be intolerable if borne alone.

Shiller begins by pushing back against the ordinary emotional image of insurance. Insurance sounds dull; worse, it sounds intrusive. One imagines the life-insurance salesman at the door, turning an ordinary afternoon into a meditation on death. Shiller wants to reverse that mood. Insurance matters because it helps make life workable. It prevents a bad state of nature from becoming a permanent family catastrophe. That, and not the sales ritual, is the real substance.

Definition 5.1. Risk pooling is the transfer of many individual exposures into a common pool so that an extreme loss for one person becomes a smaller aggregate fluctuation for the institution that bears the pool.

This is why Shiller insists that insurance belongs inside finance. The separation between “finance”

and “insurance” is, in his telling, partly an accident of history and partly a consequence of regulation. Conceptually, however, both are branches of risk management. The lecture’s refrain is simple and worth keeping in front of us: the fundamental idea is risk pooling.

A second refrain appears almost as early. Financial institutions are inventions. They are not natural objects that merely appear in society. Someone has to design them, specify what they do, and discover by trial, error, and regulation how to make them reliable. Insurance is one of the oldest such inventions, but it is still an invention.

5.2 From Ancient Risk Pooling to the Seventeenth-Century Insurance Idea

The lecture then slows down historically. Shiller does not say that insurance sprang fully formed into existence in the seventeenth century. People had been pooling risks much earlier. Ancient Rome had funeral insurance, organized through guilds and commercial associations, and this already shows the basic intuition: a group can spread an irregular personal expense across many contributors. The need was not trivial. In the ancient world, proper burial carried religious and social significance, so funeral coverage answered a real demand.

But the historical distinction matters. Early forms of risk sharing are not yet modern insurance in a clear conceptual sense. Shiller emphasizes how blurred many ancient and medieval examples remain. He recalls a purported Renaissance insurance contract that did not even sound, in its language and structure, like a modern insurance policy. The practice was there, but the concept had not yet been stated cleanly.

5.2.1 Question & Answer

Question. If people had pooled risks for centuries, what was still missing before modern insurance?

Answer. What was missing was not collective help, but conceptual clarity. People had risk sharing, guild protection, and mutual aid. What they did not yet have was a clean statement that many regular payments, collected into a fund and analyzed over a long enough horizon, could finance a relatively small number of large losses.

That clearer statement appears, in Shiller’s telling, in the seventeenth century, just as probability becomes a recognizable mathematical language. He singles out an anonymous 1609 letter to Count Oldenburg as the earliest clear description of the insurance idea. The proposal is plain: each homeowner pays roughly one percent of the value of the house each year into a common fund, and the fund replaces houses destroyed by fire. In compact notation,

$$\text{annual premium} \approx 0.01 \times \text{house value.} \tag{5.1}$$

The letter looks over a long window—Shiller mentions thirty years—and argues that the amount collected should exceed, or at least cover, the houses destroyed by fire over that same period. This is not yet a formal theorem. It is an institutional intuition: there will not be so many fires that the common fund cannot pay.

That is the step the lecture wants us to notice. Ancient risk sharing is one thing. A clear statement of insurance as a calculable social fund is another. Once probability becomes available, the second step begins to separate itself from the first.

5.3 Probability, Aristotle, and the Mathematical Logic of Pooling

Only after the historical intuition is in place does Shiller make the mathematical turn. He first reaches back to Aristotle, not to claim that Aristotle had modern probability theory, but to show that the intuition is older than the formalism. Aristotle remarks that repeating the same throw of the dice once or twice is easy, but repeating it ten thousand times is effectively impossible. The point is not a theorem; the point is that persistent exact repetition becomes fantastically unlikely as the number of trials grows.

Now the lecture makes its staircase upward. The Oldenburg letter gives us the institutional idea. Aristotle gives us the pre-probabilistic intuition. Modern probability gives us the explicit formula that turns the intuition into quantitative reasoning.

Let X be the number of insured events among n independent exposure units, each with event probability p . Writing the sample proportion as

$$X \sim \text{Binomial}(n, p), \quad \hat{p} = \frac{X}{n}, \quad (5.2)$$

we obtain the standard reconstruction of the formula Shiller states verbally:

$$E[\hat{p}] = p, \quad (5.3)$$

$$\text{Var}(\hat{p}) = \frac{p(1-p)}{n}, \quad (5.4)$$

$$\sigma(\hat{p}) = \sqrt{\frac{p(1-p)}{n}}. \quad (5.5)$$

The derivation is short and worth preserving because it shows exactly where the pooling effect comes from:

$$E[X] = np, \quad \text{Var}(X) = np(1-p), \quad (5.6)$$

$$\text{Var}(\hat{p}) = \text{Var}\left(\frac{X}{n}\right) = \frac{1}{n^2} \text{Var}(X) = \frac{1}{n^2} np(1-p) = \frac{p(1-p)}{n}. \quad (5.7)$$

Taking square roots gives the formula for $\sigma(\hat{p})$.

The key comparative-static point in the lecture is therefore

$$\sigma(\hat{p}) \propto \frac{1}{\sqrt{n}}. \quad (5.8)$$

As the number of independent policies rises, the dispersion of the realized claim proportion shrinks. The expected proportion stays fixed at p , but the realized proportion becomes more tightly concentrated around that expectation.

Worked example. Suppose $p = 0.01$. With $n = 100$ independent policies,

$$\sigma(\hat{p}) = \sqrt{\frac{0.01 \cdot 0.99}{100}} \approx 0.00995. \quad (5.9)$$

With $n = 10,000$,

$$\sigma(\hat{p}) = \sqrt{\frac{0.01 \cdot 0.99}{10,000}} \approx 0.000995. \quad (5.10)$$

The pool is one hundred times larger, but the standard deviation of the claim proportion is only one tenth as large. That is the whole force of the law-of-large-numbers intuition in this setting.

5.3.1 Question & Answer

Question. Why does writing many insurance policies make the insurer safer rather than riskier?

Answer. Because what matters is not the absolute number of claims, which does rise with the size of the pool, but the claim proportion relative to the size of that pool. Under independence, the uncertainty of that proportion falls like $1/\sqrt{n}$. So the insurer writes more policies, but its aggregate experience becomes more predictable in proportional terms.

This is exactly how Shiller glosses the formula. If the insurer writes many policies, then deviations from the mean by one or two standard deviations become increasingly unlikely in relative terms. That is the mathematical core of insurance. But the lecture does not stop there. It turns immediately and sharply: the theorem is not the institution.

5.4 Insurance as Institutional Invention: Contract Design, Moral Hazard, and Selection Bias

Here the lecture makes one of its decisive moves. Pooling works in principle. But making insurance work as an actual institution is hard. The problem is no longer purely probabilistic. It is architectural.

Shiller lists, in effect, the pieces that have to be built around the pooling theorem. First, the contract must define what is covered and what is not. Second, the contract must control incentive problems. Third, the insurer must have data. Fourth, it must have an organizational form capable of absorbing error. Fifth, it must be regulated, because the buyer may pay for years before learning whether the firm is sound.

Definition 5.2. Moral hazard is the change in behavior induced by insurance when coverage creates incentives that raise the probability or severity of loss.

The classic example in the lecture is fire insurance. If an owner can gain by destroying the insured house, then the contract undermines itself. In simple notation, if V_{house} is the value of the house and I_{house} is the insurance indemnity, then one basic anti-moral-hazard principle is

$$I_{\text{house}} < V_{\text{house}}. \quad (5.11)$$

If full overcompensation is impossible, then burning the house down ceases to be an economically attractive act. The same logic lies behind exclusions for suspicious causes of death in life insurance.

Definition 5.3. Selection bias, which in modern language overlaps with adverse selection, is the tendency for people who privately know they are high risk to seek insurance more aggressively than others.

The lecture's mechanism is straightforward. If the people who buy the policy are disproportionately likely to claim on it, then premiums must rise. Once premiums rise, healthier or lower-risk people withdraw. In compressed form,

$$\text{selection bias} \implies \text{higher expected losses} \implies \text{higher premiums} \implies \text{exit of lower-risk buyers.} \quad (5.12)$$

That is why underwriting, exclusions, and policy definitions are not minor technicalities. They are part of the economic logic of the institution.

Shiller then broadens the design problem further. Real insurance requires:

- precise definitions of the insured loss and of acceptable proof of loss;
- a mathematical model of risk pooling, beginning with independence but not ending there;
- statistics on risk, such as mortality tables, and judgment about the quality of those data;
- a company form, whether corporate or mutual, that determines who bears residual modeling and contract risk;
- government supervision, because policyholders often pay in long before they find out whether the company will perform.

The mortality-table point is especially important historically. Shiller notes that the seventeenth century did not give us probability alone. It also gave us the first systematic collection of mortality data. The insurance industry needed those data, and so the data became part of the invention.

Likewise, the distinction between corporate and mutual insurers is not merely legal. In the corporate form, shareholders absorb some of the error when the company misprices risk. In the mutual form, the institution is run more directly for the benefit of policyholders. Different ownership forms allocate the remaining uncertainty differently.

5.4.1 Question & Answer

Question. If risk pooling works mathematically, why is real-world insurance so hard to make reliable?

Answer. Because pooling solves only one layer of the problem. It reduces random fluctuation under the assumption that the modeled risks are sufficiently independent. It does not by itself solve moral hazard, selection bias, ambiguous definitions of loss, weak data, poor governance, or lack of trust. Those all have to be designed around the theorem.

This is why regulation enters the lecture so early and so naturally. If a family buys life insurance, it may pay for decades before the decisive event ever occurs. A promise so long-dated cannot rest on casual trust. Regulation is part of the machinery that makes the promise credible.

5.5 AIG and the Failure of Independence

At this point Shiller makes the discussion concrete by turning to AIG. The choice is not arbitrary. It was, until recently, the largest insurance company in the world; it became the largest bailout of the crisis; and Maurice “Hank” Greenberg was due to visit the class. So the lecture lingers on AIG’s history before it reaches the collapse.

The secure chronology is this. AIG traces back to 1919 in Shanghai, where Cornelius Vander Starr founded American Asiatic Underwriters. Shanghai was then a world business center, and the company grew into a major international institution. Starr led it for decades; Greenberg then led it for decades more. Shiller even pauses over Greenberg’s wartime biography, including the liberation of Dachau, as part of the company’s human history before the crisis. The effect is deliberate: AIG is made to feel like a long-built institution before it is shown to fail.

The economic failure comes after Greenberg’s departure in 2005. Here Shiller states the lesson with unusual bluntness: AIG failed largely because of a failure of the independence assumption. The company accumulated huge exposures to real-estate risk, both through credit default swaps and through mortgage-related securities. The comforting idea was that house prices might fall in one city or region, but not everywhere at once.

That reasoning is exactly the reasoning of risk pooling—and exactly its vulnerability. If the shocks are local, the pool diversifies them. If the shocks become national, the pool ceases to pool.

We can summarize the contrast as

$$\text{independent local housing shocks} \implies \text{diversification}, \quad (5.13)$$

but

$$\text{national housing collapse} \implies \text{correlated losses} \implies \text{pooling breaks down}. \quad (5.14)$$

That is what happened. The event AIG’s models treated as essentially impossible—a broad decline in home prices across the country—occurred. Once the losses moved together, the insurer faced not a smoothed aggregate but a concentrated shock.

The federal government then committed roughly \$182 billion to the rescue. Shiller is careful about what this did and did not mean. It did not preserve shareholder wealth. The shareholders lost more than 90% of their value, and the company later executed a reverse split,

$$20 \text{ old shares} \mapsto 1 \text{ new share}. \quad (5.15)$$

The company survived, but the old equity was devastated.

Why, then, the public anger? Shiller’s answer is subtle. The anger was not really about shareholders, who were largely wiped out. It was about counterparties. Firms on the other side of AIG’s contracts, including major financial institutions, did not suffer the same losses because the government chose to keep AIG alive rather than risk a cascading collapse. The rescue was therefore a systemic-risk intervention, not a gift to equity holders.

That is the deeper lesson. AIG is not simply a story about bad bets. It is the lecture’s concrete example of what happens when an insurance institution mistakes correlated risk for poolable risk. Once that error becomes large enough, insurance turns into a threat to the financial system itself.

5.6 Guarantee Funds, State Regulation, and the American Regulatory Structure

The lecture then asks the natural next question. If insurance companies are regulated, and if there are guarantee funds for policyholders, why did AIG require a federal intervention at all? Shiller uses that question to explain the architecture of U.S. insurance regulation.

He begins with an analogy to banking. Bank deposits in the United States are insured by the FDIC up to about \$250,000. Insurance also has backstop institutions, but they are state-based guarantee funds, not one large federal scheme. The oldest such fund in the lecture is New York's, dating to 1941; Connecticut's begins in 1972. By now, virtually all states have some version.

5.6.1 Question & Answer

Question. If insurance already has guarantee funds, why did AIG still require a federal bailout?

Answer. Because guarantee funds are designed to protect ordinary policyholders up to limited amounts. They are not designed to absorb the collapse of a giant, globally connected insurer whose failure could damage banks, counterparties, and the broader financial system.

Shiller gives the relevant magnitudes. New York and Connecticut provide limits around \$500,000. Many states are closer to \$300,000. That is enough to protect many households against the failure of an insurer, but it is nowhere near enough for a firm the size of AIG. It is not even necessarily enough for an ordinary life-insurance need once we remember tuition, housing, and family support.

The lecture also notes two instructive differences from deposit insurance. First, one may be able to spread bank deposits across multiple institutions and obtain repeated FDIC protection, but insurance guarantee funds may cap protection across policies in a stricter way. Second, at least in the state example Shiller discusses, insurers are not supposed to advertise the guarantee fund in the way banks advertise FDIC coverage. The moral is that the insurance buyer is still expected to care about the quality of the insurer.

That leads directly to the American regulatory structure. Under the McCarran–Ferguson Act of 1945, insurance regulation remains primarily a state function. So an insurer operating nationally must deal with a fragmented regulatory map. The National Association of Insurance Commissioners helps coordinate model rules and standardization, but it is not a single sovereign national regulator.

Dodd–Frank introduces a limited federal overlay after the crisis. Shiller's point is not that the United States nationalized ordinary insurance regulation. It did not. Rather, it created a federal insurance office to collect information and watch for systemic risk, and linked that process to the post-crisis systemic-risk framework centered on the Financial Stability Oversight Council. The federal role becomes active precisely when an insurer begins to look like more than an insurer—when it begins to resemble a threat to the stability of the wider system.

Here again the mathematics returns in institutional dress. The state system can handle ordinary insurance regulation. The federal overlay appears when dependence across risks becomes large enough to matter system-wide. In other words, AIG forced the government to take the independence assumption seriously as a regulatory problem, not only as a classroom formula.

Shiller briefly widens the comparison internationally. He mentions China's newer insurance protection

fund, with coverage around 50,000 yuan, or roughly \$6,000 in the lecture's conversion. The contrast underscores how limited and uneven insurance backstops remain across countries.

5.7 Types of Insurance, Health Insurance, and the Unfinished Technology of Risk Management

Only after the crisis and regulation arc does the lecture return to what looks, at first sight, like textbook classification. Even here the sequencing matters. Having seen how insurance can fail, we now return to the ordinary forms through which it still does essential work.

Types of insurance. Shiller sketches the major categories. Life insurance is the largest privately offered category in his 2009 figures, at nearly \$5 trillion. Property and casualty insurance is much smaller, at roughly \$1.3 trillion, though still enormous. Health insurance sits in between in conceptual importance if not in that particular ordering. He briefly mentions term insurance, whole life, variable life, and annuities, not to build a taxonomy for its own sake, but to remind us that every one of these is an invented contract form.

Life insurance also receives a historical comment that matters. Its social importance was once even greater than it is now, because early death was far more common. What life insurance insures, strictly speaking, is not life itself but the family against the loss of a breadwinner.

Health insurance. The lecture then settles on health insurance as the clearest modern case where moral hazard, selection bias, and institutional invention are all visible at once. Many countries adopted national systems. The United States kept a more private tradition, and therefore repeatedly ran into a structural problem: if healthy people do not enter the pool, the pool becomes too sick and premiums become too high.

That is the same mechanism we have already seen:

$$\text{healthy people stay out} \implies \text{sicker insurance pool} \implies \text{higher premiums.} \quad (5.16)$$

The lecture then adds the provider side. If doctors are paid per procedure, they may have an incentive to over-treat rather than to preserve long-run health. So the health-insurance problem is not only who enters the pool, but also how treatment incentives are arranged.

This is why Shiller presents the HMO as an institutional invention. In a health maintenance organization, doctors are salaried rather than paid per procedure. The design aims to weaken the incentive to order unnecessary interventions. The HMO Act of 1973, requiring sufficiently large employers to offer a federally certified HMO option, appears as one milestone in that direction. The Yale Health Plan, founded even earlier, serves as a local example of the same idea: align the institution with prevention and long-run care rather than procedure volume.

EMTALA, by contrast, solves a different and narrower problem. It requires emergency rooms to treat people in urgent need, whether or not they are insured. Shiller calls this, in effect, a form of national health insurance at the emergency margin, but only in a distorted sense. It is an unfunded mandate. The hospital must treat, but payment remains uncertain, so debt and strained incentives remain. The law therefore alleviates abandonment, not the underlying insurance design problem.

That brings the lecture to the 2010 reforms. Their logic is explicitly anti-selection bias. Insurance exchanges widen access. Penalties for going uninsured push healthier people into the pool. The lecture's rough annual example is about \$700 for remaining uninsured. In the intended mechanism,

$$\text{broader enrollment} \implies \text{healthier average pool} \implies \text{lower cost per insured} \implies \text{lower premiums.} \quad (5.17)$$

Restrictions on refusing pre-existing conditions then become feasible only because the pool is being broadened on the other side.

Catastrophe risk and new instruments. The closing movement of the lecture widens the horizon again. Haiti illustrates an insurance problem in a raw form. The earthquake losses were largely uninsured. That meant, first, that owners could not collect after the collapse. But it also meant something deeper: insurance companies were not present beforehand to impose underwriting discipline, building standards, and risk-based scrutiny on the construction process itself.

Katrina presents a more developed but still imperfect system. Insurance existed, and Shiller cites figures suggesting that many homes received payments. But the case exposed another difficulty central to the lecture: definitions matter. In a hurricane, was the house destroyed by wind or by flood? If different policies cover those two losses differently, then the ambiguity becomes economically explosive. The point is exactly the one made earlier in the abstract: definitions of insured loss are part of the institution, not legal clerical work.

Terrorism risk takes us back to the issue of dependence. Traditional insurance policies often exclude war and terrorism because the losses are correlated. This is not an arbitrary exclusion; it is a direct acknowledgment that the ordinary pooling model becomes unreliable when many claims can be triggered by a single common cause. TRIA, beginning in 2002 and later renewed, is the institutional answer Shiller emphasizes: insurers must offer terrorism coverage, but the government stands behind part of the truly large systemic loss.

Finally, catastrophe bonds push insurance back toward finance in a visible way. Shiller's Mexican example is used to make one clean point. A government facing earthquake risk does not want that risk concentrated entirely on its own balance sheet. By issuing catastrophe-linked bonds, it can spread that risk through global capital markets. The exact dollar amount in the transcript is noisy, but the structure is clear: investors receive return in ordinary states of the world and bear losses when the specified catastrophe occurs. Insurance, in that sense, becomes more explicitly financial.

The lecture ends on a sober but forward-looking note. Insurance is painfully slow to improve. Centuries after its invention, millions still remain badly protected, and many important risks are still only partially insured. Some risks change over time—hurricane risk, mold risk, other environmental hazards—and the industry does not yet insure well against changes in probability itself. Many long-term contracts are not properly indexed to inflation. So the institution remains incomplete.

5.8 Summary

The lecture opens by insisting that insurance is finance, and it closes by showing how much work is required to make that claim true in practice. The mathematical core is simple: with many independent risks, the dispersion of the realized claim proportion falls like $1/\sqrt{n}$. That is why risk pooling can protect people against losses that would be unbearable in isolation.

But the lecture's deeper claim is that insurance is not just a theorem. It is a technology and an institutional design problem. Contracts must be written carefully. Moral hazard and selection bias must be contained. Statistics must be gathered and judged. Firms must be organized. Regulators must make long-dated promises credible. And once independence fails, as in AIG or in terrorism risk, the institution may require a different structure altogether.

So the final impression is neither complacent nor cynical. Insurance is one of the great civilizing devices of finance. It has already made life safer and more workable. But it is unfinished, and its future progress will depend on exactly the same combination the lecture has traced from beginning to end: mathematical insight, institutional invention, and careful attention to how real risks actually move together.

Chapter 6

Guest Speaker David Swensen: The Yale Model and Long-Horizon Investing

This chapter follows David Swensen’s guest lecture in Robert J. Shiller’s *Financial Markets* course, curated here by LazyingArt LLC. The lecture opens with a live controversy rather than a formal model. Yale’s endowment had become famous, then controversial, and Swensen uses that controversy to organize the whole hour. We begin with Barron’s criticism of the “Yale model,” move backward to the portfolio world of 1985, then follow the lecture’s central decomposition into asset allocation, market timing, and security selection. Only after that machinery is in place does Swensen return to the questions of diversification, liquidity, alternatives, and institutional performance.

6.1 Opening challenge: the Yale model under criticism

Shiller’s introduction matters because it sets the stakes. He presents Swensen not only as Yale’s endowment manager but as a serious financial innovator, someone associated with the early swap market, and he places Yale’s long-run record in front of the audience before handing over the room. The subtext is clear: we are hearing from someone who has both built instruments and managed institutional capital.

Swensen begins in a different register. He jokes that the introduction would have sounded more flattering before the financial crisis. For many years the publicity around Yale’s portfolio had been good. Then came Lehman, panic, and the inevitable backlash. A Barron’s article from November 2008 had argued that colleges were suffering because the Yale model was too aggressive, too illiquid, and not truly diversified. Swensen’s first move is to turn that criticism into the lecture’s organizing question.

There is also a small but revealing rhetorical point here. Swensen notes the asymmetry in how people speak: when the portfolio succeeds it is the “Yale model,” and when it disappoints it becomes the “Swensen approach.” That joke does more than lighten the room. It marks the lecture’s posture. He is not going to answer the criticism with branding. He is going to ask, from the ground up, what kind of portfolio a perpetual institution ought to hold.

So the lecture begins not with a defense of alternatives as such, and not with a narrow post-crisis

performance argument. It begins with a structural question: what did the old endowment world look like, and what should a long-horizon investor do differently?

6.2 1985 endowments and the two failed tests

To answer that question Swensen goes back to 1985, when he arrived at Yale from Wall Street with no deep portfolio-management track record and did the sensible thing: he looked at what other endowed institutions were doing. The result, as he reports it, was remarkably uniform:

$$(w_{\text{US stocks}}, w_{\text{US bonds+cash}}, w_{\text{alternatives}}) \approx (0.50, 0.40, 0.10). \quad (6.1)$$

The lecture's diagnosis is that this conventional portfolio failed two basic tests. It failed the diversification test, and it failed the equity-orientation test.

On the first point Swensen invokes the finance theory associated with Tobin and Markowitz. Diversification is the famous “free lunch” of portfolio theory. He does not work through mean-variance geometry in the lecture, but the claim can be written in its standard form. If two portfolios have the same expected return, a better-diversified portfolio can carry lower variance,

$$E[R_p] = E[R_q], \quad \sigma^2(R_p) < \sigma^2(R_q), \quad (6.2)$$

and if two portfolios have the same variance, the better-diversified portfolio can carry higher expected return,

$$\sigma^2(R_p) = \sigma^2(R_q), \quad E[R_p] > E[R_q]. \quad (6.3)$$

This is the abstract theory. Swensen's point is that the old endowment portfolios did not actually satisfy its spirit. A portfolio with half its assets in a single asset class is not meaningfully diversified, and a portfolio with ninety percent in domestic marketable securities is not drawing on a wide enough set of underlying return drivers.

His concrete example is interest rates. The lecture's language is intuitive: lower rates are good for bonds, and lower discount rates also raise the present value of future earnings streams, so they can be good for stocks as well. Since no validated board equation survives for this lecture, the clean way to sharpen that intuition is with the standard present-value formulas:

$$P_{\text{bond}} = \sum_{t=1}^T \frac{C_t}{(1+r)^t} + \frac{F}{(1+r)^T}, \quad (6.4)$$

$$P_{\text{equity}} = \sum_{t=1}^{\infty} \frac{CF_t}{(1+r)^t}. \quad (6.5)$$

The lecture does not insist that stocks and bonds always move together. The narrower point is that they can respond in the same direction to an important macro variable, so a portfolio made mostly of domestic stocks and domestic bonds is less diversified than it first appears.

The second failed test is equity orientation. Swensen's intuition is almost embarrassingly simple, and that is part of its force. Endowments have longer horizons than almost any other investors. Their job is to preserve purchasing power in perpetuity. That means they are natural bearers of equity risk. Yet the conventional endowment portfolio held about forty percent in bonds and cash, which Swensen treats as low expected return assets. On this view, the old structure was too timid for the institution it was meant to serve.

These are the two premises from which the Yale model grows: real diversification, and a meaningful commitment to equity-like risk for a perpetual investor.

6.2.1 Question & Answer

Question. If diversification was already standard finance theory, why were endowments still so badly structured?

Answer. Swensen’s answer is historical and institutional. The old endowment world had inherited a convention and learned to call it prudence. Once a portfolio shape becomes standard across peer institutions, it can survive by habit even when it is no longer defensible from first principles. The lecture’s diagnosis is therefore not that diversification theory was unknown, but that portfolio practice had drifted away from the implications of the theory.

6.3 Three tools of return

Having reconstructed the old world, Swensen changes gears and introduces the framework that governs the rest of the lecture. There are, he says, three things we can do to affect portfolio returns. First, we can decide what assets to own and in what long-run proportions. That is asset allocation. Second, we can depart from those long-run weights because of a short-run valuation view. That is market timing.


Third, within an asset class, we can try to beat the market by overweighting some securities and underweighting others. That is security selection.

This is the lecture’s main analytical hinge. Once the three levers are separated, the structure of the argument becomes much clearer. In note form we can summarize the decomposition as

$$R_p = R_{\text{policy}} + \alpha_{\text{timing}} + \alpha_{\text{selection}} - L, \quad (6.6)$$

where R_{policy} is the return from the long-run policy portfolio, the α ’s are gains or losses from active deviations, and L denotes the leakages of active investing: fees, commissions, market impact, loads, taxes, and related costs.

The flow of the lecture can be summarized by the following editorial diagram:



The next step is especially important. Swensen says that asset allocation is “far and away” the most important tool, but he immediately warns us not to misunderstand that as a universal law of nature. If Yale put its entire endowment into Google stock, security selection would dominate. If Yale used the whole endowment to day-trade bond futures, market timing would dominate. These examples are deliberately absurd. Their purpose is to show that the dominance of asset allocation is behavioral rather than metaphysical.

Why, then, does asset allocation dominate in practice? Because real investors usually do something much more boring. They maintain a fairly stable long-run portfolio. They diversify within asset classes. They do not continually transform the entire endowment into one concentrated bet. Under those conditions the main determinant of realized return variation is the set of asset-class exposures themselves.

This is also how Swensen interprets the Ibbotson-style finding that more than ninety percent of the variability of institutional returns is attributable to asset allocation. Indeed, he pushes the point further. If market timing and security selection are zero-sum before costs and negative-sum after costs on average, then the policy portfolio explains more than one hundred percent of net returns:

$$\overline{\alpha_{\text{timing}}} + \overline{\alpha_{\text{selection}}} - \overline{L} < 0. \quad (6.7)$$

In other words, the active deviations do not merely fail to help on average; they subtract from what passive implementation of the policy portfolio would have produced.

6.3.1 Question & Answer

Question. Why can asset allocation be the dominant source of returns even though market timing and security selection also exist?

Answer. Because the lecture's claim is conditional on sensible portfolio behavior. Once we assume that investors hold stable, diversified policy portfolios, the large-scale movements of asset classes dominate the return process. The cases in which timing or selection dominate are logically possible, but they are the pathological cases that Swensen uses precisely in order to set them aside.

6.4 Long horizons, equity risk, and the trap of literalism

Swensen now turns from framework to long-run evidence. Using Roger Ibbotson's historical data, he asks what would have happened to one dollar invested from the end of 1925 to 2009. The lecture is not interested in a tidy annualized abstraction. It wants the cumulative comparison:

$$M_{\text{bills}} = 21, \quad M_{\text{inflation}} = 12, \quad (6.8)$$

$$M_{\text{bonds}} = 86, \quad M_{\text{large stocks}} = 2592, \quad (6.9)$$

$$M_{\text{small stocks}} = 12226. \quad (6.10)$$

The immediate moral is that long-run investors are rewarded for accepting equity risk. Bills do not preserve much after inflation. Bonds are better, but still modest over many decades. Equities are in a different league entirely. So far the lecture's message is straightforward: a perpetual institution should carry meaningful exposure to risky productive assets.

But then Swensen performs one of the lecture's most important reversals. If the historical table is read too literally, it seems to imply something absurdly simple. If small stocks do best by such a large margin, why not put the whole endowment into small stocks? This is where the lecture's rhythm matters. Swensen lets the tempting inference appear before taking it apart.

6.4.1 Question & Answer

Question. If small stocks dominate the long-run data, why not put the whole endowment into small stocks?

Answer. Because long-run arithmetic is not the whole problem. We must also ask whether the strategy is survivable. Swensen's chosen example is the Great Depression, and here the lecture slows down enough for us to work the arithmetic:

$$W_{1929} = W_0(1 - 0.54) = 0.46W_0, \quad (6.11)$$

$$W_{1930} = 0.46W_0(1 - 0.38) = 0.2852W_0, \quad (6.12)$$

$$W_{1931} = 0.2852W_0(1 - 0.50) = 0.1426W_0, \quad (6.13)$$

$$W_{1932} = 0.1426W_0(1 - 0.32) \approx 0.097W_0. \quad (6.14)$$

So the dollar at the peak becomes roughly ten cents at the trough. That is the worked mathematical example at the heart of the section, and it changes the problem completely. The question is no longer whether small stocks were best in the very long run. The question is whether an investor, however strong his stomach, can actually remain invested through such a collapse.

Swensen's answer is emphatically no. At some point dollars turning into dimes will overwhelm the abstract case for long-run expected return. This is why the lecture refuses the crude rule "just own the thing with the highest historical multiple." Heavy equity exposure may be correct for a long-horizon investor, but undiversified concentration in one risky segment is not.

That is also why Swensen pauses to recall the cultural disgust for equities after the crash, the old language about stocks being "insecurities." The point is not nostalgic. It is analytical. A strategy that cannot be held through the path of bad states is not a serious institutional policy. Diversification is what makes equity orientation liveable.

6.5 Market timing and security selection as negative-sum behavior

At this point the lecture pivots from long-run compounding to investor behavior. The second lever is market timing, and Swensen introduces it through Keynes. The relevant passage is memorable because it captures both the behavioral and the arithmetic problem: people who attempt large shifts sell too late, buy too late, do both too often, and pay heavy expenses in the process.

The lecture then supplies evidence. Morningstar compared time-weighted returns and dollar-weighted returns across seventeen categories of domestic-equity mutual funds. The distinction is central. Time-weighted returns describe the fund's own return path. Dollar-weighted returns tell us what investors actually experienced once their inflows and outflows are incorporated. Swensen's summary is strikingly simple:

$$R^{DW} < R^{TW} \quad (6.15)$$

in every one of the seventeen categories. That is a remarkably clean empirical sign of perverse timing. Investors add money after good performance and pull money out after bad performance. They buy high and sell low.

The internet-fund example makes the same point with sharper narrative force. Swensen reports that the time-weighted return across the top internet funds over the six-year window around the bubble

was about 1.5% per year. Yet investors allegedly put in roughly \$13.7 billion and lost roughly \$9.5 billion. The lecture's arithmetic at that moment is rough, but the logic is unmistakable. Investors were not steadily exposed through the whole period. They piled in near the top, then sold in disappointment after the collapse. Once again, the timing of the path mattered more than the headline statistic.

Nor do institutions escape. The October 1987 crash becomes the institutional analogue of the mutual-fund story. Stocks fell sharply, bonds rallied in a flight to safety, and institutions responded by selling stocks and buying bonds. Swensen's conclusion is blunt: both individuals and institutions have a recurring tendency to chase performance and thereby destroy returns.

Security selection is the third lever, and the lecture handles it with the simplest possible relative-market logic. If we buy the market inside an asset class, there is no active selection return:

$$\alpha_{\text{selection}} = 0. \quad (6.16)$$

If instead we overweight one stock and underweight another relative to the market, then someone else must hold the offsetting position. In aggregate, before costs,

$$\sum_i \alpha_i^{\text{selection}} = 0. \quad (6.17)$$

After commissions, market impact, management fees, consultant fees, loads, and taxes, it follows that

$$\sum_i \left(\alpha_i^{\text{selection}} - L_i \right) < 0. \quad (6.18)$$

Swensen applies a similar, though slightly looser, intuition to market timing in the aggregate. Someone's overweight relative to target must be another person's underweight, and once the cost of playing is included, timing too becomes negative-sum on average.

6.5.1 Question & Answer

Question. If active management is zero-sum before costs, why does the average active investor still lose?

Answer. Because the lecture's actual object is not gross alpha but net alpha. Swensen's cited evidence gives only about a 14% chance of beating the market after fees and taxes. He then weakens even that figure. Many funds impose front-end loads, so the game is even more expensive than the published comparisons suggest. And survivorship bias flatters the winners by ignoring dead funds. Swensen's striking database numbers make the point concrete:

$$N_{\text{total}} = 30,361, \quad N_{\text{living}} = 19,129, \quad N_{\text{dead}} = 11,232. \quad (6.19)$$

A large fraction of funds disappear, mostly because they fail. The average investor therefore faces a game that is zero-sum before costs, negative-sum after costs, and statistically hard to win even before we correct the data for survivorship.

6.6 Where active management matters

Once the lecture has shown that active management is structurally costly, a new question arises. If active management is not attractive everywhere, where is it worth attempting? Swensen's answer is one of the lecture's most interesting practical refinements.

He proposes using dispersion as a proxy for opportunity. The exact quartile language in the transcript is a little noisy, so it is safest to speak generally of the spread between stronger and weaker managers in a given asset class. If a market is highly efficient, that spread should be tight. If it is less efficient, the spread should be wider:

$$\Delta_{\text{bonds}} \approx 0.5\% \text{ per year,} \quad (6.20)$$

$$\Delta_{\text{large cap}} \approx 2.0\% \text{ per year,} \quad (6.21)$$

$$\Delta_{\text{foreign}} \approx 4.0\% \text{ per year,} \quad (6.22)$$

$$\Delta_{\text{absolute return}} \approx 7.1\% \text{ per year,} \quad (6.23)$$

$$\Delta_{\text{real estate}} \approx 3.0\% \text{ per year,} \quad (6.24)$$

$$\Delta_{\text{LBO}} \approx 13.7\% \text{ per year,} \quad (6.25)$$

$$\Delta_{\text{VC}} \approx 43.2\% \text{ per year.} \quad (6.26)$$

The lecture's interpretation is not abstractly efficient-market in tone. It is operational. Bonds are “just math”—coupons, principal, probabilities of default. They are easier to analyze, so the scope for persistent skill is smaller. That is why managers in those markets tend to be packed close together. The safest thing for a manager to do in such an environment is often to hug the benchmark.

At the other end are private markets, especially venture capital. There may not even be a benchmark that can be mechanically replicated. The assets are idiosyncratic, the enterprises are private, and manager skill can matter much more. This is why Swensen's argument for alternatives is never simply “private markets are good.” It is instead: if active management is costly everywhere, then we should spend our active effort where the opportunity set is widest.

The audience questions sharpen this point. Swensen is asked whether the growth of hedge funds and private equity has made those markets too efficient to bother with. His answer is nuanced. Fees and leakages have unquestionably grown; that makes the negative-sum problem worse. But he does not think dispersion has collapsed. So long as the spread between the best and worst managers remains large, the activity remains potentially worthwhile for an institution able to identify superior managers.

6.7 Returning to Barron's: diversification, alternatives, and the institutional coda

Only now does Swensen return to Barron's. On diversification, he grants a narrow point: in a panic, diversification can appear to fail. Investors compress the world into two bins, risk and safety. Risky assets are sold together, and U.S. Treasuries become the unique refuge. In that window, the only thing that looks like protection is ownership of the asset everyone is running toward.

But Swensen insists that this is the wrong horizon on which to judge portfolio design. To hold a very large Treasury position all the time is to pay a large opportunity cost year after year. The panic

window may last six months or a year; the cost of defensive underinvestment can last decades. That is why he brings in Japan. If one had wanted an equity-oriented portfolio but held only domestic Japanese equities, then from 1989 to 2009 one would have lived through a roughly seventy-three percent decline:

$$\text{Nikkei}_{1989} \approx 38,000, \quad \text{Nikkei}_{2009} \approx 10,500, \quad (6.27)$$

so that

$$1 - \frac{10,500}{38,000} \approx 0.724 \approx 73\%. \quad (6.28)$$

The lesson is exact and limited. Diversification may disappoint in a panic, but concentration can destroy long-run purchasing power.

On the charge of overemphasizing alternatives, Swensen shifts from principle to record. Over the ten years ending June 30, 2010, Yale's returns in several alternative sectors were stronger than in domestic public equities:

$$R_{\text{domestic eq}} \approx -0.7\% \text{ per year}, \quad R_{\text{bonds}} \approx 5.9\% \text{ per year}, \quad (6.29)$$

$$R_{\text{private eq}} \approx 6.2\% \text{ per year}, \quad R_{\text{real estate}} \approx 6.9\% \text{ per year}, \quad (6.30)$$

$$R_{\text{absolute return}} \approx 11.1\% \text{ per year}, \quad R_{\text{timber}} \approx 12.1\% \text{ per year}, \quad (6.31)$$

$$R_{\text{oil\&gas}} \approx 24.7\% \text{ per year}. \quad (6.32)$$

He then broadens the performance comparison to Yale as a whole. The lecture cites approximately 8.9% per year over ten years versus an average university figure near 4.0%, and approximately 13.1% over twenty years versus a university average near 8.8%. The tone at this point is understandably firm. Barron's, he argues, had looked at too short a window.

The final audience questions turn this defense into practical advice. Swensen says Yale's job differs from a hedge fund manager's in an important way: Yale sits one layer above direct security selection. Its central task is to find and structure relationships with excellent managers. That is why the principal-agent problem becomes so important. Yale must induce agents to act in the university's interest.

He also distinguishes institutions like Yale from ordinary investors on taxes and resources. Yale does not pay taxes; individuals do. Yale can devote a professional staff to manager selection; most individuals cannot. This is why, when asked what ordinary investors should do, Swensen does not tell them to imitate the full Yale model. He tells them to choose a sensible policy allocation and implement it with low-cost index funds.

6.7.1 Question & Answer

Question. If diversification fails during panics, why should a long-term investor still insist on it?

Answer. Because the long-term investor's problem is not to maximize comfort in a panic; it is to preserve purchasing power across decades without being ruined by concentration. Panic behavior and long-run portfolio design are not the same question. A portfolio that looks brilliant only in crisis and costly the rest of the time is not obviously well designed. Swensen's answer is that the right comparison must include both the opportunity cost of too much safety and the catastrophic risk of country or asset-class concentration.

That same logic reappears when Swensen answers a later question about whether stocks are currently expensive. His first response is not a market forecast. It is that the valuation question matters less in a well-diversified portfolio with disciplined rebalancing. Once we have target weights w_i^* , rebalancing pulls realized weights back toward policy:

$$w_i \rightarrow w_i^*. \quad (6.33)$$

The arithmetic is plain. If an asset underperforms and drifts below target, we buy it back up to target; if it outperforms and drifts above target, we trim it back down. The lecture's later numerical guidance is therefore practical and memorable: a diversified portfolio may have something like a 5% to 10% minimum position and a 25% to 30% maximum position in any one asset class. Within such a structure, rebalancing quietly accomplishes much of what investors wrongly try to do with dramatic market timing.

Swensen is also asked about technology exposure. His answer is consistent with the earlier dispersion argument. Yale has a longstanding venture-capital commitment, and it also uses managers who specialize in technology and biotechnology because those sectors can be less efficiently priced than more mature public-equity segments. So even here the lecture's general rule holds: do not speculate because a sector is fashionable, but do recognize that some markets may offer more room for skilled selection than others.

The last important caution concerns risk measurement. Swensen is explicitly skeptical of simple Sharpe-ratio comparisons across institutional portfolios:

$$\text{Sharpe} = \frac{E[R_p - R_f]}{\sigma(R_p)}. \quad (6.34)$$

He does not reject the formula; he rejects the casual use of the denominator. Public markets are marked continuously and exhibit large measured standard deviations. Illiquid assets are appraised infrequently, and appraisal smoothing mechanically lowers reported volatility. So when we compare a marketable-securities portfolio to a portfolio heavy in illiquid assets, we are not necessarily comparing like with like. In his terms, this is an apples-and-oranges problem.

The lecture ends with one final sharp contrast. For those with extraordinary resources and genuine access to superior managers, aggressive active management may make sense. For everyone else, the middle ground is the most dangerous place to stand. One ends up paying high fees for mediocre active results. That is why Swensen's practical advice for most investors is a disciplined passive policy, not a watered-down imitation of institutional complexity.

6.8 Summary

Swensen's lecture begins as a defense of the Yale model, but by the end it has become something more general: a theory of how a long-horizon institution should think. The first principle is that endowments ought to be equity oriented because their horizon is long and their mission is to preserve purchasing power in perpetuity. The second is that real diversification is broader than a conventional mix of domestic stocks and bonds. The third is that once we separate asset allocation from market timing and security selection, we see why active behavior so often subtracts from rather than adds to returns.

The lecture never says that all alternatives are good, or that diversification is magic in every state of the world. Its claim is more precise. Diversification can fail in a panic and still be indispensable

over decades. Active management can be negative-sum in the aggregate and still be worth pursuing in unusually wide-dispersion markets for institutions with exceptional resources. And for most investors, the right response to all this is not complexity for its own sake, but a sensible policy portfolio, disciplined rebalancing, and a refusal to confuse excitement with investment intelligence.

Chapter 7

Efficient Markets

These notes follow Robert J. Shiller's lecture on efficient markets in *Financial Markets*, curated by LazyingArt LLC. The lecture does not begin with a settled doctrine. It begins with a provocation. Efficient markets, Shiller says, is a half-truth. If prices already absorb public information, and if the market therefore ought to beat any one of us, what are we to do with David Swensen's long record at Yale? The lecture answers this by moving in a deliberate sequence: first through the weakness of simple performance statistics, then through the history of the efficient-markets idea, then through technical analysis as a rival way of reading prices, and finally through the random-walk and mean-reverting models that give the theory its mathematical spine.

7.1 Opening paradox: efficient markets and David Swensen

Shiller begins with the strongest classroom form of the hypothesis. Markets incorporate public information. Therefore we should not expect to outsmart the market by trading on what is already known. In the lecture's deliberately blunt slogan, the market knows more than we do, and if we think we can beat it, the market will usually win.

Definition 7.1. In the working sense of this lecture, the efficient-markets hypothesis says that publicly available information is incorporated into market prices so rapidly that one should not expect systematic excess profits from trading on that information after it has become public.

That is already a sweeping proposition. But the lecture does not leave it in the abstract. Shiller brings in David Swensen, whose investment record at Yale had just been discussed in the previous class. Swensen is presented as someone who seems to have beaten the market over a very long period beginning in 1985, and with visible institutional consequences. Yale's financial strength supported policies such as wide need-blind access. Thus the opening problem is not merely statistical. It is institutional, human, and concrete.

7.1.1 Question & Answer

If Swensen appears to beat the market for decades, is that skill, luck, or a problem with the theory?

Shiller refuses to settle the question with a slogan. A defender of efficient markets can always say that if millions of managers try, one of them will look brilliant by luck alone. But the reply comes

quickly: twenty-five years is a long time to wave away with a one-line appeal to luck. At the same time, raw return alone cannot decide the matter. We must ask what risks were taken, how those risks were measured, and in what kinds of markets the manager chose to compete. That is why the lecture turns almost immediately to the Sharpe ratio.

7.2 Sharpe ratios, hidden risk, and the temptation to game the statistics

The bridge into the mathematics comes from a leftover question from Swensen's appearance. Why, if Yale's portfolio had done so well, did nobody talk about the Sharpe ratio? Shiller says he caught himself praising Yale's return without correcting for risk. That self-correction matters, because a high return by itself proves very little. Any expected return can be raised by taking more risk. That was already one of the lessons of the efficient frontier and the tangency line earlier in the course.

In the lecture's loose but serviceable form, we may write the Sharpe ratio as

$$S \approx \frac{E[R_p - R_b]}{\sigma(R_p)}, \quad (7.1)$$

where R_p is portfolio return and R_b is a benchmark return. In Shiller's spoken exposition the benchmark is "the market," so we keep the symbol R_b only as a cautious notation choice.

The ingredients being summarized are the mean and standard deviation of return:

$$\mu_R = E[R], \quad \sigma_R = \text{SD}(R). \quad (7.2)$$

A high excess return accompanied by a high standard deviation is not evidence of exceptional skill; it may merely be the compensation for taking exceptional risk.

Swensen's reply, as Shiller reconstructs it, is that for a broad portfolio full of illiquid assets the denominator may not be measurable in a clean way. Private equity is not continuously traded. Real estate may sit for years with only appraisals rather than repeated market transactions. Appraisal-based series smooth observed returns. In that case $\sigma(R_p)$ can be artificially low before any deliberate manipulation has even begun.

Shiller then makes the argument harsher. Suppose a manager knows that the world will judge him by his Sharpe ratio. Suppose, further, that he has no ethical concern except to look good for a few years, gather assets, and walk away rich. The question becomes: what portfolio design makes the ratio look best, even if the underlying investment is rotten?

The answer, drawn from a paper by William Goetzmann and coauthors, is to reshape the tails. One begins with a bell-shaped return distribution. The mean and standard deviation of that distribution are what the Sharpe ratio sees. But the manager can alter the distribution itself. He can sell off the upper tail, giving away rare large gains in exchange for money now, and at the same time deepen the lower tail, so that rare bad outcomes become much worse. Most of the time nothing dramatic happens. The middle of the distribution looks calm. The portfolio records steady gains. The statistic looks beautiful. What has improved is not the investment. What has improved is the disguise.




Figure 7.1: Bell-curve sketch for the return distribution. The frame supports a generic probability picture over a horizontal return axis, but not a fully labeled statistical graph.



Figure 7.2: Left: a cautious redraw of the bell-shaped return distribution visible on the board. Right: an editorial reconstruction of the lecture's verbal mechanism, with clipped upside and a heavier downside tail. The asymmetry on the right is transcript-based exposition, not a literal board redraw.

7.2.1 Question & Answer

How can a portfolio show a high Sharpe ratio while hiding catastrophic risk?

Because the Sharpe ratio only summarizes the realized sample. If the catastrophic left tail has not yet occurred, the reported mean and standard deviation can look entirely respectable. A manager who has clipped away rare upside and loaded rare downside receives income now, looks smooth for a while, and postpones the disaster to a future date. The true risk has not vanished. It has merely been pushed into rare states of the world that the sample has not yet visited.

Shiller's option language is informal, but the mechanism is clear. One may write away upside on the upper tail and otherwise create extra downside exposure through option positions. The point is not to produce a perfectly specified derivatives strategy on the board. The point is to understand how a statistic based on mean and standard deviation can be gamed by changing the unseen architecture of the tails.

7.3 Integral Investment Management, disclosure, and the moral of the Swensen recap

At once Shiller says that this is not merely a blackboard trick. He points to Integral Investment Management, a hedge fund that used an options strategy of roughly this character. Investors poured money into it. Most notably, the Art Institute of Chicago put about \$43 million into the fund and related vehicles. When markets fell sharply in 2001, the strategy imploded and the Art Institute lost almost all of that investment.

The fund had looked extraordinary. That was the trap. A very high Sharpe ratio was mistaken for evidence of unusually good investing when in fact it was evidence that the measured statistics were concealing the true structure of the risk.

Shiller then sharpens the point further. The explicit derivatives trick is in one sense too obvious. Any experienced professional, once the strategy is laid bare, should recognize it. But the same logic can be implemented more subtly. One can simply buy firms or assets with large left tails: securities that carry a small probability of very large losses. For a while the returns look strong and the volatility looks low. Later the fund blows up. The lecture briefly gestures toward political tail risk of that kind, but the general lesson is sufficient: the market can hide bombs inside apparently ordinary securities.

This prepares the legal and moral turn. As Shiller states it, the law does not permit a manager to rely on boilerplate alone. If there is relevant information that would make the published statistics misleading, the manager must actively disclose it and make the client understand it. A generic sentence saying that past returns do not guarantee future results is not enough if the real issue is a specific embedded tail exposure.

The deeper moral returns us to Swensen. The substance of Swensen's achievement, Shiller suggests, is not mastery of a performance ratio. It has something to do with character, integrity, and the real objectives of the institution he serves. Statistics matter, but they do not tell the whole story. In real financial life, people judge not only numbers but persons and purposes. Only after this detour does Shiller say, in effect, that he wants now to go more directly into the lecture proper on efficient markets.

7.4 Before the phrase: Gibson, Reuters, and Conant

The lecture now resets historically. Shiller wants to reconstruct the intuition of efficient markets before the phrase itself became standard. The first clear statement he can find is George Gibson in 1889. Once shares are known in an open market, Gibson says, the value they acquire there may be taken as the judgment of the best intelligence concerning them.

This is already the core efficient-markets picture. The market is a voting mechanism, but not in the casual sense of polling opinions. If we think a share is undervalued, we buy. If we think it is overvalued, we sell. The better informed traders have stronger incentives to act and, if they succeed repeatedly, more capital with which to act. Price is therefore not the average thought of the public. It is a pressure point where the strongest and best-informed views are forced into action.

Just as important is the speed of information. Gibson writes in an electric age. In 1889 that means telegraphy, ticker machines, and rapid dissemination of quotes. The market is efficient, if it is efficient, not because wisdom falls from the sky, but because there is a relentless incentive to collect




Figure 7.3: Gibson, Reuters, and Conant on the board. The arrangement is historically suggestive, but the frame should be read as board evidence rather than as a formal timeline.

and transmit information faster than one's rivals.

Reuters gives this intuition its institutional form. Before the telegraph became dominant, Reuters used carrier pigeons to move financial information between cities. What sounds quaint is, economically, perfectly modern. If news can be moved even a little faster, money can be made. That is why information transmission becomes an industry. The pigeons are an early version of the same race that later uses telegraph wires, ticker machines, beepers, and electronic terminals.

Shiller makes the same point with a contemporary example. A firm announces a successful new drug. Analysts hear the news immediately. Within seconds and minutes they are on the telephone to specialists, revising estimates and trading before the rest of the public has even begun to process the announcement. Prices jump, overshoot, correct, and settle. By the next morning, when a newspaper reader decides he may want to buy the stock, the market is far ahead of him. This is the part of efficient markets that Shiller clearly wants us to keep. Public information becomes old information with astonishing speed.

This also begins to answer the Swensen puzzle. Shiller connects the point back to the prior lecture. In broad public markets such as bonds and large equities, even top-quartile managers may have little room to outdistance everyone else. But in unusual asset classes such as private equity or certain absolute-return strategies, public information is thinner and the market less crowded. Part of Swensen's achievement, then, may lie in choosing where the game is easier to play.

Conant, writing in 1904, gives the next major statement. He begins by rejecting the notion that stock-market speculation is merely gambling. That cannot be the whole truth, because the stock market is a central institution of the modern economy. If the stock exchanges of the world were closed, people would lose their public signals of value. Capital allocation would become blind. Economic calculation itself would be impaired.

7.4.1 Question & Answer

Why is a stock market a central economic institution rather than just an organized gambling hall?

Because market prices do real social work. They summarize information, permit comparison, and help move capital from one use to another. Conant's point is not that gambling motives disappear in finance. It is that the market's deeper function is price discovery. Without quoted prices, investment decisions lose an indispensable guide. That is why a modern economy cannot simply dismiss speculation as morally noisy clutter. It is entangled with valuation itself.

7.5 From doctrine to revision: Roberts, Fama, CRSP, and the softening of orthodoxy

Only later does the phrase itself become standard. Shiller credits Harry Roberts at the University of Chicago with using the label, and Eugene Fama with making the idea famous. By then the theory had acquired something it did not possess in Gibson's or Conant's time: a modern data infrastructure.

In 1960 the Ford Foundation funded a large effort at Chicago to assemble reliable stock-price histories back to 1926, correcting for stock splits, dividends, and other details that had never been systematically organized. This became the Center for Research in Security Prices, now usually abbreviated CRSP. The famous tape was important not because tape is glamorous, but because finance could suddenly work with an organized universe of data rather than scraps from newspapers.

That changed the rhetoric of the theory. Once the data were assembled, the efficient-markets hypothesis could be tested on a scale that felt authoritative. By the end of the decade there were thousands of articles using the CRSP database. Fama's 1969 review article, *Efficient Capital Markets*, gave the movement its canonical survey. His conclusion was not that every result favored efficiency, but that the evidence pointed toward remarkably efficient markets. The profession heard something stronger still: perhaps the very idea of beating the market was mostly a delusion.

The 1970s were therefore the high point of the doctrine. The theory did not merely describe prices. It seemed to discredit a vast investment industry. Successful managers, on this view, must mostly have been lucky. Yet Shiller's point is historical as well as analytical: the orthodoxy later weakened without becoming useless.

He shows this by turning to textbooks. Earlier editions of major finance texts could still say, with remarkable confidence, that security prices reflect the true underlying value of assets. Later editions retreated. They admitted that prices sometimes become badly misaligned with what appear to be discounted future payoffs. That is not a minor qualification. It is a substantial softening.

Shiller reinforces the point with Jonah Lehrer's phrase "the truth wears off." The suggestion is not that science is fake or that earlier work was worthless. It is that an intellectual movement can begin with real insight, gather prestige, overstate its scope, and then be cut back by later evidence. Shiller even entertains the possibility that early enthusiasm itself can bias a research literature, just as it can in medicine or other sciences. Efficient markets, then, should be treated neither as a revelation nor as a dead theory. It should be treated as a strong idea that was once allowed to harden too far.

7.6 Technical analysis as the foil

After this long historical buildup, Shiller turns to technical analysis. The placement matters. Technical analysis is not introduced as a joke. It is introduced as a serious alternative to the efficient-markets mindset. Instead of reading balance sheets, economic news, or institutional structure, the technician reads price charts and looks for patterns in the path of prices themselves.

Edwards and McGee are the canonical figures here, and McGee appears as a psychologically minded analyst. He is so committed to price-reading that he wants an interior office with no windows, free from the distractions of the outside world. That anecdote is more than colorful. It captures the chartist conviction that the pattern is in the price, and that the price already contains the real drama.

Shiller gives two classic examples. The first is the resistance level. When the Dow first approached 1000, technicians thought it might have difficulty crossing such a psychologically imposing threshold. The second is the head-and-shoulders pattern, in which a higher peak is flanked by two lower peaks and is supposed to herald a collapse. McGee regarded these patterns as expressions of crowd psychology. The market, on this view, has moods that are visible in its shapes.

7.6.1 Question & Answer

Does technical analysis actually work, or does it only find patterns in noise?

Shiller's answer is deliberately balanced. It is too simple to call technical analysis pure bunk. The efficient-markets literature in its strongest phase often dismissed chart reading more completely than the evidence warranted. Shiller even recounts asking Burton Malkiel, after *A Random Walk Down Wall Street* appeared, where the many cited studies disproving technical analysis actually were. He suspected that the dismissal had outrun the bibliography. Later work did find some element of truth in some classic chartist observations.

But the opposite exaggeration is just as false. Technical analysis is not a machine for easy money. At best it is a subtle art that may occasionally augment broader judgment. That is precisely why the lecture now needs a stronger benchmark. If we are to decide whether patterns are meaningful, we must first understand what sheer unstructured price motion looks like. That benchmark is the random walk.

7.7 Random walk, AR(1), and the half-truth conclusion

The logical transition is direct. If prices are informationally efficient, then day-to-day price changes should come only from news. But news, as news, is unforecastable. Therefore price changes should themselves be unforecastable. This is the intuition behind the random walk.

Definition 7.2. A random walk is a process $\{X_t\}$ whose one-step change is an unforecastable mean-zero shock.

In symbols,

$$X_t = X_{t-1} + \epsilon_t, \quad (7.3)$$

$$E[\epsilon_t] = 0, \quad (7.4)$$

where ϵ_t is the innovation term. For classroom convenience Shiller allows us to imagine that the shocks have some standard deviation σ_ϵ and perhaps even a Gaussian distribution:

$$\epsilon_t \sim N(0, \sigma_\epsilon^2). \quad (7.5)$$

He later warns us not to trust the Gaussian approximation too much, but it is the easiest starting point.

Historically, the image is Carl Pearson's drunk at the lamppost. The drunk has no directional plan. Asked for a forecast of his position after some time, we do not infer a trend away from the lamppost. We take the current position as the center of the forecast and let the uncertainty widen with time. That is already the random-walk picture.

Proposition 7.3. *If the shocks ϵ_t are independent with mean zero, then for the random walk*

$$E[X_{t+n} | X_t] = X_t, \quad (7.6)$$

and the forecast dispersion obeys the square-root rule:

$$\text{Var}(X_{t+n} - X_t) = n\sigma_\epsilon^2, \quad \text{SD}(X_{t+n} - X_t) = \sqrt{n}\sigma_\epsilon. \quad (7.7)$$

Proof. Iterating the one-step equation gives

$$X_{t+n} - X_t = \epsilon_{t+1} + \epsilon_{t+2} + \cdots + \epsilon_{t+n}.$$

Conditional on X_t , each shock has mean zero, so

$$E[X_{t+n} - X_t | X_t] = 0,$$

which yields $E[X_{t+n} | X_t] = X_t$. If the shocks are independent, their variances add:

$$\text{Var}(X_{t+n} - X_t) = \sum_{j=1}^n \text{Var}(\epsilon_{t+j}) = n\sigma_\epsilon^2.$$

Taking square roots gives the square-root rule. □

This is the point at which the lecture turns back against technical analysis. A random walk can produce striking local patterns by accident. In fact Shiller's simulation later finds shapes that look head-and-shoulders-like even though the process generating them contains no true chartist signal at all. That is why the random walk is such a powerful foil: it teaches us how easily the eye invents structure.

Shiller then contrasts the random walk with the process written on the board in the lecture's final mathematical section. The screenshot preserves the blackboard comparison itself, and the comparison is part of the pedagogy.

Cleaned up into note-quality notation, the board comparison is

$$X_t = 100 + \rho(X_{t-1} - 100) + \epsilon_t, \quad -1 < \rho < 1, \quad (7.8)$$

$$X_t = X_{t-1} + \epsilon_t. \quad (7.9)$$

The right-hand shock term in the screenshot is not perfectly legible and looks closer to E_t , but the transcript supports the standard reconstruction ϵ_t . The first equation is the AR(1) alternative, centered at 100; the second is the random walk.




Figure 7.4: AR(1) process versus random walk. The screenshot matters here because it preserves the lecture's board layout: Pearson and *Nature* 1905 at left, the mean-reverting alternative in the center, and the random walk at right.

The point of the AR(1) is mean reversion. The anchor level 100 plays the role of Pearson's lamppost. If the process is above 100, it tends to stay above 100 for a while, but it is pulled partway back. If it is below 100, it is pushed partway upward.

Proposition 7.4. *For the AR(1) process*

$$X_t = 100 + \rho(X_{t-1} - 100) + \epsilon_t,$$

if $E[\epsilon_t | X_{t-1}] = 0$, then

$$E[X_t - 100 | X_{t-1}] = \rho(X_{t-1} - 100). \quad (7.10)$$

Hence when $0 < \rho < 1$, deviations from 100 shrink in conditional expectation.

Proof. Subtract 100 from both sides:

$$X_t - 100 = \rho(X_{t-1} - 100) + \epsilon_t.$$

Now take conditional expectation given X_{t-1} :

$$E[X_t - 100 | X_{t-1}] = \rho(X_{t-1} - 100) + E[\epsilon_t | X_{t-1}] = \rho(X_{t-1} - 100).$$

If $0 < \rho < 1$, the deviation keeps its sign but becomes smaller in magnitude on average. □

Worked example. Shiller gives the concrete classroom case $\rho = \frac{1}{2}$. If the previous observation is $X_{t-1} = 120$, then

$$E[X_t | X_{t-1} = 120] = 100 + \frac{1}{2}(120 - 100) = 110. \quad (7.11)$$

The process is still above 100, but only half as far above as before. That is mean reversion in the simplest possible form.

7.7.1 Question & Answer

Why does $\rho = 1$ turn the AR(1) process into a random walk, and what changes when ρ is only slightly below 1?

The algebra is immediate:

$$X_t = 100 + \rho(X_{t-1} - 100) + \epsilon_t, \quad (7.12)$$

$$\rho = 1 \implies X_t = 100 + (X_{t-1} - 100) + \epsilon_t \quad (7.13)$$

$$= X_{t-1} + \epsilon_t. \quad (7.14)$$

So at $\rho = 1$ the anchor disappears from the law of motion. The process is no longer being pulled back. It is exactly the random walk.

If ρ is only slightly below 1, however, mean reversion is present but extremely weak. This is where Shiller introduces the loose-elastic-band metaphor. The drunk is still tied to the lamppost, but by a very loose elastic. He may wander for a long time before being pulled back. That is why an AR(1) with $\rho = 0.99$ or 0.98 can look very much like a random walk over realistic horizons, even though strict long-run mean reversion still exists.

The simulation segment makes this visible, though it does so in a slightly improvised way. Shiller compares actual long stock-market history with artificial series generated from random numbers. The striking point is that the random walk often looks uncannily like real market history. It produces booms, slumps, and plausible-looking patterns, even when no meaningful pattern has been built into the process.

He then compares those simulations with an AR(1) process. If ρ is substantially below 1, the process hugs its trend too tightly. In such a world profitable rules would be too easy: buy when the price is below trend and sell when it is above, because reversion is too reliable. Shiller's judgment is that the actual stock market does not look like that strong-reversion world.

At the same time, he corrects himself during the spreadsheet demonstration. Some of the simulated paths include an upward deterministic term. A stylized way to write that is

$$X_t = X_{t-1} + c + \epsilon_t, \quad (7.15)$$

rather than a pure zero-drift random walk. The exact bookkeeping of the spreadsheet is not the heart of the lecture. The real lesson is comparative. A random walk, even with a simple drift term, often resembles long stock-market history more closely than a low- ρ AR(1) with strong mean reversion.

There is, however, a final caveat. The Gaussian random walk misses something important. In the simulation, Shiller does not get anything like the crash after 1929. The reason is not mysterious. A normal shock distribution has no fat tails:

$$\epsilon_t \sim N(0, \sigma_\epsilon^2) \quad \text{rules out the crash frequency we actually care about.} \quad (7.16)$$

So even the random-walk model that best captures the shape of long market history remains only an approximation.

That is the lecture's last turn. The random walk is illuminating, because it teaches us how hard it is to extract genuine signals from price charts and how quickly public information is absorbed. But it is not the whole story. If the world is more like an AR(1) with ρ very close to 1 than like a

perfect random walk, then there may be long-horizon opportunities, though not easy short-run ones. In that sense efficient markets is neither nonsense nor gospel. It is, just as Shiller said at the start, a half-truth.

7.8 Summary

The lecture begins with Swensen and ends with stochastic processes, but the path between them is the point. We first learn why raw performance statistics are inadequate and why the Sharpe ratio itself can be manipulated when tail risk is hidden. We then step back into the history of the idea, from Gibson's market as aggregated intelligence, through Reuters and the economics of information speed, to Conant's defense of speculation as price discovery rather than mere gambling. Only then do we arrive at the Chicago codification of the doctrine, its later softening, the challenge posed by technical analysis, and the random-walk versus AR(1) comparison that gives the lecture its mathematical payoff.

The conclusion is balanced in exactly the way the opening promised. Public information is absorbed rapidly, so simple money machines are elusive. Yet prices are not so perfect that every deviation from value disappears instantly, nor so Gaussian that crash risk fits neatly inside normal shocks. Efficient markets is useful because it captures something real. It is limited because it captures only part of what real financial markets do.

Chapter 8

Theory of Debt, Its Proper Role, and Leverage Cycles

This chapter follows Robert J. Shiller's eighth lecture in Yale's *Financial Markets* course, curated for the LazyingArt LLC track. The lecture is explicitly split in two. We do the technical things first: an Irving Fisher model of interest, present value, discount bonds, compounding, conventional bonds, the term structure, and forward rates. Only after that do we return to the larger question of what debt markets are really doing in social life. That order matters. The mathematics is introduced as groundwork for judgment, not as a substitute for it. The opening puzzle is therefore a serious one: why are interest rates positive at all, and why do they so often sit at only a few percent per year?

8.1 Why interest rates need an explanation

An interest rate is the percentage one earns on a loan or pays on a loan, depending on which side of the contract one stands. The lecture begins by insisting that we not treat that as a trivial fact. Interest rates are ancient. They have existed for thousands of years. They are usually positive. They are often in the neighborhood of a few percent. Why should that be so? Why not zero? Why not negative? Why not something drastically larger?

Shiller begins, not with a pricing formula, but with the history of thought. In Bohm-Bawerk's late nineteenth-century account, the rate of interest is explained in broad verbal terms by three causes.

- Technical progress: as knowledge accumulates, future production may be better than present production.
- Roundaboutness: production that takes time, including indirect methods and tool-making, may be more productive than direct production.
- Time preference: people may prefer present consumption to future consumption; that is, they may be impatient.

The important point is that this is still literary economics. It is suggestive and powerful, but it is not yet a model that places these causes inside one clear mechanism. Fisher's contribution is to take this loose verbal list and convert it into a disciplined two-period equilibrium picture.




Figure 8.1: Irving Fisher board and cropped two-period diagram. The screenshot fixes the historical framing; the transcript supplies the economic labels and the logic of the geometry.

8.2 From Bohm-Bawerk to Fisher's two-period framework

Shiller briefly acknowledges the textbook summary associated with Fisher: the interest rate may be drawn as the intersection of a supply curve and a demand curve for savings. The supply of saving slopes upward because higher rates induce more people to save. The demand for investment funds slopes downward because lower rates make borrowing more attractive to businesses. This is not wrong. But it is too compressed for the lecture's purpose.

So we move to Fisher's more revealing diagram. Shiller lingers over the historical fact that Fisher was a Yale professor and even remarks that one can imagine Fisher lecturing from the same blackboard. The point of the historical setup is not nostalgia. It is that Fisher's 1930 *Theory of Interest* provides a diagram that brings Bohm-Bawerk's causes together more succinctly than the textbook supply-and-demand cross.

The axes are current consumption and next year's consumption. To animate them, Shiller tells a Robinson Crusoe story. Crusoe has grain. He must decide how much to consume this year and how much to save and plant for next year. If total current grain endowment is e , current consumption is c_0 , and next year's consumption is c_1 , then in a cleaned notation consistent with the lecture's linear technology story we may write

$$c_1 = (1 + r)(e - c_0). \quad (8.1)$$

The slope of the production possibility frontier is therefore

$$\frac{dc_1}{dc_0} = -(1 + r). \quad (8.2)$$

Shiller says on the board that the slope is "minus 1 plus r ." In the notes it is best to normalize that to the standard form $-(1 + r)$.

The first Fisher story is intentionally simple. For every bushel of grain planted, Crusoe may get two bushels next year. In that stripped-down case, $1 + r = 2$. The frontier is a straight line. Crusoe still




Figure 8.2: Cleaned reconstruction of Fisher’s one-person two-period geometry after diminishing returns are introduced. The cropped board confirms the presence of the graph; the transcript supplies the economic structure.

chooses a point on it by comparing present and future consumption, and the difference between his endowment and his current consumption is saving. But if the frontier is linear, the interest rate itself is pinned down by technology. Preferences determine where he lives on the line, not the line’s slope. So impatience has not yet fully entered the rate of interest.

That is why the lecture immediately takes the next step.

8.3 Robinson Crusoe, first alone and then in a loan market

The next step is diminishing returns. If Crusoe plants only a little grain, he may use the best land and do very well. But as he saves more and more, the additional grain planted may be less productive. The production possibility frontier then bends inward. Now the rate of interest is no longer just the slope of a straight technological opportunity set. It is the slope at the point where technology and taste meet.

Crusoe chooses the highest indifference curve tangent to the production possibility frontier. At that tangency,

$$\text{MRS}_{c_0, c_1} = \text{MRT}_{c_0, c_1} = 1 + r. \quad (8.3)$$

Now we can finally say what Fisher adds to Bohm-Bawerk. Roundaboutness and productivity are carried by the frontier. Technical progress can shift the frontier itself. Time preference appears in the slope and shape of the indifference curves. The interest rate is therefore not just impatience and not just productivity. It is the slope at a tangency between the two.

This also yields the lecture’s comparative statics. Suppose Crusoe becomes more impatient. Then he wants more consumption today and less next year. The tangency moves to the right. Because it moves to a steeper part of the frontier, the implied interest rate rises. If, instead, Crusoe is very patient, the tangency lies further up and to the left, and the implied interest rate is lower. So impatience matters, but only through its interaction with the production frontier.

At that point Shiller recaps: we now have all of Bohm-Bawerk’s causes. But he is not satisfied yet. A one-person economy has no trade. So he adds a second Crusoe.




Figure 8.3: Two Crusoes with different patience. Production occurs at the common tangent point P , while lending supports different consumption points A and B on the same intertemporal line.

The second stage is the lending market. There are now two Crusoes on opposite sides of the island. They have the same technology and the same endowment, but not the same patience. One wants more consumption today and saves little. The other is patient and saves heavily for the future. If they remain isolated, they choose different tangencies on the same frontier. But once they discover each other, there are gains from trade.

The reason is not merely that tastes differ. It is that the patient Crusoe is planting so much that he is deep into diminishing returns, while the impatient Crusoe is planting relatively little, where marginal productivity is still high. A loan transfers grain toward the more productive use and simultaneously moves both agents toward preferred consumption plans.

The market interest rate is the slope of the common tangent line. That is the lecture's central theoretical statement. The interest rate is not read off from any one person's impatience. It emerges from heterogeneous preferences interacting with a common technological opportunity set.

8.3.1 Question & Answer

Question. Why is a loan socially useful in the Fisher story?

Answer. Because with diminishing returns and different degrees of patience, isolated production plans are inefficient. The patient Crusoe is planting too much where marginal productivity is low; the impatient Crusoe is planting too little where marginal productivity is high. A loan reallocates grain toward the more productive margin and lets both agents move to higher indifference curves. The common tangent line clears the loan market and makes both better off.

This is precisely where the lecture pauses. Put this way, lending looks obviously good. Shiller does not abandon that conclusion. He says, in effect, let us take this Fisher theory of interest as given. But before we return to criticisms of lending, we must do the arithmetic of finance.

8.4 Debt arithmetic: discount bonds, compounding, and present discounted value

Here the lecture deliberately changes level. We move from the equilibrium logic of interest to the pricing of loan instruments. The first and simplest instrument is the discount bond.

A discount bond promises a fixed payment at a future date and sells today for less than that payment. It has no coupon stream along the way. If a bond pays 100 at maturity T , then under annual compounding its present price is

$$p = \frac{100}{(1+r)^T}, \quad (8.4)$$

$$(1+r)^T = \frac{100}{p}, \quad (8.5)$$

$$r = \left(\frac{100}{p} \right)^{1/T} - 1. \quad (8.6)$$

The rate r inferred from price is the yield to maturity. The point is subtle but important: the bond itself is quoted as a price, but we can back out an annual interest rate from that price once maturity and compounding are specified.

Worked calculation. For a two-year discount bond,

$$p = \frac{100}{(1+r)^2} \quad \implies \quad r = \sqrt{\frac{100}{p}} - 1. \quad (8.7)$$

This is the simplest explicit bridge from bond price to interest rate.

Shiller then stops and carefully distinguishes compounding conventions. Under annual compounding, interest is credited once per year. If one deposits B_0 today, then after T years the balance is

$$B(T) = B_0(1+r)^T. \quad (8.8)$$

But finance often prefers semiannual conventions, because conventional bonds pay coupons every six months. The transcript's notation becomes unstable here, so it is best to normalize it once:

$$z = \frac{r}{2}, \quad n = 2T, \quad p = \frac{100}{(1+z)^n}, \quad (8.9)$$

where z is the half-year rate and n counts half-year periods.

The lecture then takes the limiting case: compound more and more frequently until compounding becomes continuous. The full sentence on the board is partly obscured, but the standard completion consistent with the lecture is

$$B(T) = B_0 e^{rT}. \quad (8.10)$$

So annual compounding and continuous compounding are not two unrelated formulas. They are two versions of the same intertemporal calculation under different conventions.




Figure 8.4: Annual versus continuous compounding on the board. The screenshot matters because it preserves the comparative board layout, not just the final exponential term.

Present discounted value. This is the central technical concept of the middle of the lecture. If a single payment x arrives at date T , then under annual compounding

$$\text{PDV} = \frac{x}{(1+r)^T}. \quad (8.11)$$

If annual payments $\{x_t\}$ arrive over time, then

$$\text{PDV} = \sum_{t=1}^{\infty} \frac{x_t}{(1+r)^t}. \quad (8.12)$$

If payments arrive every six months, the corresponding semiannual expression is

$$\text{PDV} = \sum_{t=1}^{\infty} \frac{x_t}{(1+r/2)^t}. \quad (8.13)$$

And if payments arrive continuously, then the sum becomes an integral:

$$\text{PDV} = \int_0^{\infty} x(t)e^{-rt} dt. \quad (8.14)$$

Shiller emphasizes that this is what finance repeatedly does. It takes many future payments and turns them into one present number. Once that number has been computed, very different instruments can be compared on a common basis.

8.5 Perpetuities, annuities, and conventional bonds

Now the board becomes formula-dense, but the lecture still proceeds in a definite order. We do not start with the full conventional bond. We begin with the simplest infinite payment stream, then truncate it, and only then add principal repayment.




Figure 8.5: Annuity setup beside the consol benchmark. The screenshot preserves the moment when the lecture moves from the perpetuity formula to the finite annuity case.

A consol, or perpetuity, pays the same coupon forever. The board writes the special case of a one-pound coupon:

$$PDV_{\text{consol}} = \frac{\pounds 1}{r}. \quad (8.15)$$

In normalized form we write

$$PDV_{\text{consol}} = \frac{C}{r}, \quad r = \frac{C}{PDV_{\text{consol}}}. \quad (8.16)$$

This is the benchmark against which the annuity is introduced.

An annuity is a finite payment stream. It pays x each period from $t = 1$ to $t = T$, where T is the last payment date. The lecture writes the final formula on the board, but its logic is especially clear if we derive it from the consol. An annuity is just a perpetuity with its tail removed:

$$PDV_{\text{annuity}} = \frac{x}{r} - \frac{x}{r(1+r)^T} \quad (8.17)$$

$$= \frac{x}{r} \left(1 - \frac{1}{(1+r)^T} \right). \quad (8.18)$$

That is the formula Shiller arrives at in the annuity panel. It is important not only as algebra but as a building block. The lecture's concrete example is the traditional home mortgage: fixed payments for a finite horizon and then no more.

The final step in this pricing sequence is the conventional bond. A conventional bond pays coupon C every six months and then pays principal at maturity together with the final coupon. The board states the decomposition more clearly than the full algebra, so we write the normalized formula cautiously:

$$PDV_{\text{bond}} = \frac{C}{z} \left(1 - \frac{1}{(1+z)^n} \right) + \frac{F}{(1+z)^n}, \quad z = \frac{r}{2}, \quad n = 2T. \quad (8.19)$$

The first term is the coupon annuity. The second is the discounted principal. This is exactly the lecture's conceptual point: a conventional bond is an annuity plus a discount bond.




Figure 8.6: Finite-annuity formula applied to a conventional bond. The board makes visible both the formula and its immediate reuse for bond coupons.




Figure 8.7: Coupon timeline for a conventional bond. The diagram is subordinate to the board evidence, but it makes the annuity-plus-principal structure explicit.

8.6 Forward rates from the term structure

The next pivot is one of the lecture's best. We move from pricing a single stream to relating prices across maturities. Shiller motivates this through Sir John Hicks and a story about trying to identify who first formulated forward rates. Hicks's reply points, not to a heroic isolated discovery, but to coffee-hour conversations at the London School of Economics in the 1920s.

That anecdote matters because it shows what the problem was. Open the paper in 1925 and one sees yields for one year, two years, three years, and longer. Those are all rates from today to some future date. But what about the rate between two future dates, say from 1926 to 1927? Hicks's insight is that this rate need not be quoted separately. It is already implicit in the term structure observed today.

8.6.1 Question & Answer

Question. How can we lock in a one-year interest rate for a future year using only bonds traded today?




Figure 8.8: Forward-rate identity and expectations-theory trading setup. The screenshot preserves the board's mixture of algebra, dates, and trading instructions.

Answer. Let r_1 be today's one-year yield and r_2 today's two-year yield. Suppose that in 1925 we know we will have £100 available to invest in 1926, and we want to lock in now the one-year return from 1926 to 1927. The lecture's construction is to buy, in 1925,

$$\frac{(1+r_2)^2}{1+r_1}$$

two-year discount bonds and simultaneously short one one-period bond.

The short one-period bond implies a payment of £100 in 1926. That is exactly the amount we are prepared to invest at that date. The long two-year position, scaled as above, pays

$$£100 \cdot \frac{(1+r_2)^2}{1+r_1}$$

in 1927. Therefore the gross return from 1926 to 1927 is

$$1+f = \frac{(1+r_2)^2}{1+r_1}, \quad (8.20)$$

$$f = \frac{(1+r_2)^2}{1+r_1} - 1. \quad (8.21)$$

That is the one-step forward-rate identity written on the board.

Shiller then gives the identity its economic interpretation. Expectations theory says, in words, that the forward rate equals the expected future spot rate. In compact notation, one may write

$$f \approx \mathbb{E}_t[r_{t+1}^{\text{spot}}]. \quad (8.22)$$

The key phrase in the lecture is that the whole future seems to be laid out in this morning's paper. Once the term structure is observed, forward-rate formulas extract implied future one-period rates all along the horizon.




Figure 8.9: Timeline for the one-step forward-rate construction. The trade uses 1925 positions to lock in the gross return between 1926 and 1927.

But Hicks immediately qualifies the pure expectations theory. Forward rates are not simply forecasts. They also contain risk compensation. In a cleaned analytic normalization,

$$f = \mathbb{E}_t[r_{t+1}^{\text{spot}}] + \pi, \quad \pi > 0. \quad (8.23)$$

So forward rates tend to lie above optimally forecast future spot rates. The gap is the risk premium. This is the lecture's closing technical point.

8.7 From the theory of interest to the problem of usury

Only now does Shiller return to what he postponed near the start. The Fisher story makes lending look unambiguously welfare-improving. The loan market allows gains from trade. The term structure lays out rates for many horizons. The whole apparatus looks elegant and socially useful. Why, then, have religious and social traditions so often treated interest as dangerous or immoral?

The lecture answers by stepping back into the language of usury. The old Latin word *usura* can mean use, and it can also mean interest. More troublingly, it can shade into excessive or abusive interest. That ambiguity matters. Shiller does not try to settle biblical or Islamic exegesis. Instead he uses the ambiguity to make a conceptual point: suspicion of lending persists because in real life a difference in present and future consumption plans is not always an innocent difference in taste.

8.7.1 Question & Answer

Question. If lending raises welfare in the Fisher story, why has interest so often been condemned?

Answer. Because in actual debt markets the borrower may be impatient in a harmful rather than a merely descriptive sense. The Fisher model treats a desire for present consumption as a preference. Real life may force a different judgment. The borrower may be making a mistake, acting on temptation, ignoring future distress, or responding to manipulation. In that case the right response may be advice or protection, not simply more credit. Usury, in the lecture's closing definition, is abusive lending undertaken without genuine concern for the borrower's welfare.

That is why Shiller brings behavioral economics back into the story. If one of the Crusoes wants to consume far too much today, perhaps the real issue is not that a loan has failed to appear. Perhaps the issue is that he should be told not to do it. Once we say that, the path from Fisher's clean geometry to the real world becomes less smooth.

The examples at the end of the lecture are chosen for exactly this reason. Vacation loans sound dubious. Yet Modigliani's remark about honeymoons complicates the matter: some expenditures that look like consumption may actually be investments in memory, relationship, and life trajectory. So the lecture refuses a simple moralism. It does not say that such loans are always wrong. It says that the real question is whether lenders are helping borrowers make good intertemporal choices or exploiting their vulnerabilities.

Elizabeth Warren enters at the close as the institutional form of this concern. Her work on bankruptcy, household distress, and abusive lending practices is used to show that modern debt markets can systematically harm people while appearing, on the surface, to offer them choice. The Consumer Financial Protection Bureau appears in the lecture as a modern answer to a very old problem. We do not reject lending. We regulate it.

This is the balance Shiller wants to keep. The original Fisher and Bohm-Bawerk story about interest was basically right. Even apparently soft loans, including honeymoon loans, are not automatically illegitimate. But the old suspicion of usury is not empty rhetoric. It survives because lending can become predatory, and because a useful debt market still needs government rules against abuse.

8.8 Summary

The lecture begins with a puzzle and ends with a distinction. The puzzle is why interest exists at all and why it is positive and modest in scale. The distinction is between useful debt and abusive debt. In between, the lecture unfolds step by step. Bohm-Bawerk offers the old verbal causes. Fisher turns them into a two-period geometry in which technology and preferences jointly determine the interest rate. The arithmetic of finance then carries that rate into discount-bond pricing, compounding conventions, present discounted value, perpetuities, annuities, and conventional bonds. Forward rates show that the term structure already contains prices for future intervals, though not without risk premia. And only then does the lecture return to the moral question. Lending is not repudiated. It is defended, but defended with the important qualification that real borrowers are behaviorally fragile and that debt markets need regulation against usury.

Chapter 9

Corporate Stocks

These companion notes follow Robert J. Shiller's Yale lecture on corporate stocks, with curation by LazyingArt LLC. We begin where the lecture begins, not with a formula but with a small company and a problem of fairness: who gets how many shares when one founder contributes money, another contributes time, and no one can value these contributions exactly in advance? From there the argument widens to the world scale of listed equity, then narrows again to boards, dividends, financing, and finally to Xerox and Microsoft, where the arithmetic of market capitalization and shareholder equity becomes concrete.

9.1 Corporate Stocks as an Organizational Form

We begin with the familiar picture of a corporation: people set up a company, divide it into shares, later issue more shares as the company grows, and sometimes sell those shares to someone else. The lecture reminds us that this form is old. The Dutch East India Company appears again as the first celebrated corporation with traded shares. The point is not antiquarian. The corporation is a durable scheme for organizing business.

Shiller then grounds that abstraction in his own experience. He helped found Case-Shiller Weiss, Incorporated, with Alan Weiss, Chip Case, and a businessman, Chuck Longfield. Only one of the four put money into the company; the others put in time and expertise. The immediate question was the one that always lurks behind the word *equity*: how should the firm be divided? In the story, they did not pretend to solve that valuation problem exactly. They divided it four ways, partly out of fairness, partly because they wanted to begin.

That anecdote is not ornamental. It tells us what the corporate form is doing. A corporation lets heterogeneous contributions be bundled into a common project. It lets people stay involved at different intensities. It also gives the firm a way to manage resentment. Shiller says that when someone seemed to be doing more work than someone else, the natural salve was more shares, that is, a bonus. A share is therefore not just a claim on money at the end. It is also an instrument for motivating labor and settling conflict inside the organization.

This is why the lawyers matter so much in the story. Shiller says that founding a company gave him a new appreciation for lawyers because the charter anticipated conflicts before the founders fully understood what those conflicts would be. The charter and bylaws do not create the energy of the enterprise, but they keep the enterprise from dissolving as soon as the first resentments appear.

That is the lecture's first deep point: the corporation is a social technology for combining money, labor, incentives, and conflict management.

He adds that the company was eventually sold to Fiserv in 2002, when it had only 12 employees, and that later the indexes themselves were acquired by S&P and became widely known as the S&P/Case-Shiller indexes. The company never became a giant firm, but that is not the lesson. The lesson is that even a small company already exhibits the essential corporate logic that we later see again, at much larger scale, in Microsoft or Xerox.

9.2 The Scale of Listed Equity in the World Economy

From the small company the lecture suddenly turns to the scale of listed equity around the world. The World Bank slide is important because it shows two different magnitudes at once: the dollar value of listed stocks and the value of those stocks relative to GDP.




Figure 9.1: World Bank stock-market capitalization table.

The cleaned mathematical object behind the slide is

$$\frac{\text{market capitalization}}{\text{GDP}}. \quad (9.1)$$

The lecture immediately warns us not to read this too literally. We are comparing a stock with a flow: a stock of wealth in the stock market divided by one year's GDP. The ratio is still informative, but only if we remember what is being divided by what.

A few visible rows from the slide are enough to carry the argument:

Region	Fraction of GDP	US\$ trillions
United States	0.817	11.737
United Kingdom	0.696	1.851
Canada	0.669	1.002
East Asia and Pacific	0.559	3.244
Euro Area	0.380	5.153

The lecture's emphasis is that the United States is not merely large in absolute terms. It is also large relative to GDP. In 2008 its listed market capitalization was worth almost a full year's GDP. The United Kingdom is also large relative to GDP, while the Euro Area is large in dollar magnitude but much smaller on this stock-to-flow comparison.

Shiller connects that scale to history. The United States became a major capitalist center early, and he points to the New York corporate law of 1811 as part of the legal background that helped set the form in motion. He also expects the chart eventually to look different. Stock markets are spreading and growing elsewhere because, in his view, the corporate form works. Even countries that once defined themselves against capitalism have adopted stock markets in some form. The lecture uses that historical spread as evidence that the institutional device has been learned and retained.

But almost immediately Shiller pulls back. He does not want to exaggerate the size of the stock market. If we round U.S. listed market value to \$12 trillion and divide by roughly 300 million people in 2008, then

$$\frac{12 \text{ trillion dollars}}{300 \text{ million people}} \approx \$40,000 \text{ per person.} \quad (9.2)$$

For a family of four this becomes about

$$4 \times \$40,000 = \$160,000. \quad (9.3)$$

That is not nothing, but it is not the whole economy. Shiller presses the point further by comparing stocks with housing wealth. He says that U.S. houses were worth about \$18 trillion in 2010, more than the stock market, and that total household assets were about \$69 trillion. Stocks are important, but they are only one claim on national wealth.

The lecture also insists on why the stock market is smaller than many people vaguely imagine. Corporate stocks are only claims on profits after many other payments have already been made. Corporations pay wages, interest, taxes, and many other expenses before common shareholders receive what remains. So even in a highly capitalist economy, the stock market is not the whole social product. It is the capitalized value of a residual claim.

9.3 Corporation, Equity, and Governance

At this point the lecture narrows again and asks what sort of thing a corporation actually is. Before that, though, it makes a useful distinction. The World Bank slide is about traded stocks. Not every corporation is publicly traded. Shiller's own company was never traded; that is private equity. Public equity means shares traded on an exchange. He gives the German terminology as a reminder that this distinction is not peculiar to the United States: *AG*, *Aktiengesellschaft*, indicates a stock-market company, while *GmbH* indicates the nontraded form.

He then pauses over the word *equity*. In ordinary English the word has many meanings, but here it means shares in a corporation. He speculates that the term evokes equality or fairness, and institutionally that is a useful guide. Common shares are, in the ordinary case, treated equally. That is the financial meaning that matters in the lecture.

The next term is *corporation* itself. From the Latin *corpus*, body, the corporation is a legal person. Not a natural person, but an artificial one. It can own property, incur liabilities, make contracts, and continue over time even as the human beings around it enter or exit. Shiller even notes that corporations in some form go back to Rome, where *publicani* had something like a stock market in the Roman Forum, though corporate proliferation on a mass scale is mainly a nineteenth-century phenomenon.

Once we have legal personhood, governance follows. The standard picture is simple:

- the shareholders own the corporation through shares;
- each share normally carries one vote;
- shareholders elect a board of directors;
- the board hires and supervises the chief executive officer.

This is what the lecture calls shareholder democracy. It is modeled on political democracy, but it is not one-person-one-vote. It is one-share-one-vote. That difference matters, because control is tied to ownership.

Shiller makes the institutional details concrete. In the United States one chooses a state in which to incorporate. Delaware is popular because of its corporate law and tax environment. State law typically requires an annual shareholder meeting. Yet the shareholders do not run the company day to day. The board does. The board is small, perhaps six to ten people, and is expected to know the company well enough to judge management. Directors are not generally the employees who run the firm. They are the people brought in to supply judgment, perspective, and oversight. The CEO, by contrast, is the energetic full-time executive, working long hours and reporting to the board.

The board is not meant to be decorative. It can replace the CEO. That is why the lecture brings in Carl Icahn. The point is not to celebrate every activist investor, but to show that shareholder control can become real when a large shareholder pushes a complacent board to act. Boards owe duties of loyalty to shareholders, not merely friendliness toward management.

The lecture then turns to the contrast between for-profit and nonprofit corporations. Yale itself is the nearby example. Yale is a nonprofit corporation: it has a board, legal personhood, and an institutional mission, but it has no shareholders entitled to profits. Any surplus stays inside the institution and is used for education and research. That sounds very different from a for-profit corporation, yet Shiller emphasizes that the organizational sociology can be surprisingly similar. The same kinds of people can sit on both sorts of boards, and they may behave in similar ways. The distinction is not that one has structure and the other does not. The distinction is that in the for-profit case there are owners with claims on distributed profits.

9.4 Shares, Market Capitalization, and the Logic of Dividends

Only after all this structure is in place does the lecture ask the direct financial question: what is a share? The arithmetic is elementary, but Shiller deliberately repeats it, because students and

investors alike often lose sight of the denominator.

Let P_t be the price per share at time t , N_t the number of shares outstanding, and MC_t the market capitalization. Then

$$MC_t = P_t N_t. \quad (9.4)$$

Equivalently,

$$P_t = \frac{MC_t}{N_t}. \quad (9.5)$$

If we own n_t shares, our ownership fraction is

$$\frac{n_t}{N_t}. \quad (9.6)$$

This is why the raw number of shares owned is meaningless without the total number outstanding. If a company has 1,000,000 shares and we own 1,000, then

$$\frac{1000}{1,000,000} = \frac{1}{1000}.$$

If instead the company has 10,000,000 shares, then the same 1,000 shares amount only to

$$\frac{1000}{10,000,000} = \frac{1}{10,000}.$$

That corrected arithmetic is essential, because the lecture keeps returning to ownership fractions, not raw share counts.

The lecture also insists that common shares are treated equally. If the firm pays anything to one common shareholder, it must pay the other common shareholders in proportion to their holdings. That is the practical meaning of equal shares.

This leads naturally to dividends. A dividend is paid only when the board decides to pay it. Shiller remembers exactly such a conversation in his own small company. Employees had been given tiny amounts of stock as motivation, and the board wondered whether it should pay a dividend simply so that those employees would receive checks in the mail and feel that their shares were worth something. That anecdote matters because it places us exactly where the lecture wants us: between the legal discretion of the board and the economic point of owning stock.

9.4.1 Question & Answer

Question. If boards are not legally required to pay dividends, why do common shares have value at all?

Answer. The lecture's answer is that the whole point of a for-profit corporation is eventually to distribute profits to its owners. A young company may postpone that day. Microsoft is the example Shiller keeps in reserve: for a long time it paid no dividend and simply kept the money, because shareholders believed that the real opportunities still lay ahead. But a company that would never distribute anything would cease, in economic substance, to look like a for-profit company at all. Shareholders tolerate delay, not permanent denial.

That is why the lecture corrects a common misunderstanding. People say they buy stocks because the price will go up. Shiller says that in theory this is backward. The price should go up only because investors expect more dividends later. The benchmark statement is

$$P_t = \sum_{j \geq 1} \frac{\mathbb{E}_t[\text{Div}_{t+j}]}{(1+r)^j}. \quad (9.7)$$

The lecture does not derive this formula formally. It uses it as the disciplinary idea behind the rest of the discussion. A share is valuable because it is a claim on expected future payouts.

The remaining worry is board discretion. What if the board simply hoards the cash? Shiller's answer is governance. Shareholders elect the board. If the board accumulates money and refuses to distribute it despite strong profits, shareholders can replace the board. In the lecture, Carl Icahn reappears here in thought if not in person: put someone like that on the board and the money starts to come out. So the present-value logic is not detached from institutional control. It depends on it.

9.5 Dividend Pricing and Financing Mechanics

Once that puzzle has been resolved, the next piece is almost embarrassingly simple. If the company pays cash out, the company is worth less by that amount. This is the logic of the ex-dividend price.

Let the aggregate cash dividend at time t be D_t , with N_t shares outstanding. Then the dividend per share is

$$d_t = \frac{D_t}{N_t}. \quad (9.8)$$

The lecture's pricing logic can be written as a short derivation:

1. before the dividend, the firm owns the cash that is about to be distributed;
2. the board pays out aggregate cash D_t ;
3. assets fall by D_t , while liabilities do not automatically change;
4. therefore equity value falls by roughly D_t ;
5. per share, the price should fall by roughly d_t .

So the ex-dividend adjustment is

$$P_t^{\text{ex}} \approx P_t^{\text{cum}} - d_t. \quad (9.9)$$

Shiller treats this as everyday arithmetic, not a deep theorem. In the old newspapers, and now on websites, the ex-dividend adjustment is marked because the observed price decline is not bad news. It is merely cash leaving the company.

This is also why "selling dividends" is misleading and, in the lecture's telling, ethically dubious. A broker should not urge a client to rush into a stock merely because a dividend will be paid in three days. If one buys before the dividend, one gets the dividend but pays the cum-dividend price. If one buys after the dividend, one misses the dividend but buys the stock more cheaply. Apart from taxes and second-order frictions, it comes to about the same thing.

From there the lecture moves back to the small-company picture. Shares are valuable partly because they are saleable. If one founder gets tired and wants out, the firm or the other founders can buy

back the shares. If a new investor wants in, the firm can sell shares at the current price, which may be very different from the price at which the company was first organized. Saleability lets the organization persist even when particular people come and go.

The same logic explains financing. If the company needs a new factory, it can issue more shares and raise money. In a private company there may be no exchange quote, but a broker can still search for wealthy buyers or other investors willing to purchase the new issue. This is equity financing. The alternative is debt financing, by bank loan or bond issue. Debt carries fixed repayment obligations; equity carries a residual claim on profits. The lecture plainly intends that contrast, despite a garbled line in the transcript.

Once the firm borrows, common shareholders become more exposed. Their claim is the residual

$$SE_t = A_t - L_t, \quad (9.10)$$

where A_t denotes assets and L_t liabilities. Equity is what remains after creditors have been paid. This is the lecture's leverage point: debt does not merely add financing, it also amplifies the risk borne by common shareholders because their slice is now thinner and comes later in the priority structure.

9.5.1 Question & Answer

Question. Why is dilution not automatically bad, and why is a buyback not automatically good?

Answer. Because neither event can be judged from the share count alone. Suppose a firm has 1,000,000 shares outstanding and we own 1,000 of them. Then we own one-thousandth of the firm. Now suppose the board issues another 1,000,000 shares at \$10 per share. The lecture's arithmetic is

$$N = 1,000,000, \quad n = 1000, \quad \frac{n}{N} = \frac{1}{1000}, \quad (9.11)$$

$$N' = 2,000,000, \quad n = 1000, \quad \frac{n}{N'} = \frac{1}{2000}. \quad (9.12)$$

Our ownership fraction has been cut in half. That is dilution. But the firm has also received

$$1,000,000 \times \$10 = \$10,000,000 \quad (9.13)$$

in cash. So the right question is not simply whether our fraction fell. The right question is whether the firm sold the new shares at a good enough price and whether the cash will be used well.

The reverse side is repurchase. If the firm buys back shares, the share count falls. If it buys back 5% of all shares and we tender 5% of our own holdings, then we end up with cash in hand and essentially the same ownership fraction in what remains. That is why Shiller says that a buyback can feel much like a dividend. One historical reason firms preferred repurchases was tax treatment: dividends were taxed as income while capital gains could be taxed more lightly. By the time of the lecture, those rates were closer together, but tax considerations still remained complicated enough to matter.

The lecture then uses this financing discussion to introduce common and preferred stock. Common stock is the real ownership claim, the ultimate residual claimant. Preferred stock sits above common. It usually promises a fixed dividend, often has no ordinary vote, and typically must be paid up

before common dividends resume. It is therefore safer than common stock but lacks the full upside of common. The government purchase of preferred shares in General Motors during the bailout illustrates the point. Preferred shares let the government provide capital without looking as though it had fully nationalized the company by taking ordinary voting control.

9.6 Do Stock Markets Matter for Big Firms, and How Do Firms Choose Dividends?

At this point the lecture tests its own story. It is easy to see how a small company of founders, employees, and early investors uses shares to organize itself. But what about a giant company? Bill Gates and Paul Allen started Microsoft in 1975 as a little company. Once a firm grows very large, does the stock market still play the role we have been describing, or does it become mostly a market in old shares?

Shiller raises the challenge in several forms. He briefly invokes the old Marxian suspicion that stock markets may become detached from real production. He then turns to the modern finance version of the same complaint, Stuart Meyers's pecking-order theory. On Meyers's aggregate U.S. data for 1973–1982, the financing mix looked roughly like

$$\text{financing mix} \approx 62\% \text{ retained earnings} + 6\% \text{ net equity} + \text{remainder debt.} \quad (9.14)$$

That picture suggests a hierarchy. Firms prefer to use their own retained earnings first, then debt, and only as a last resort issue new shares. If that were the whole story, one might conclude that large firms hardly use the stock market for financing at all.

9.6.1 Question & Answer

Question. If big firms often finance internally, what is the stock market still doing for capitalism?

Answer. The lecture's answer is that the aggregate net number is too crude. What Meyers calls 6% equity financing is *net* equity issuance:

$$\text{net equity issuance} = \text{new shares sold} - \text{share repurchases.} \quad (9.15)$$

A low net number can arise even when many firms are issuing shares, provided many others are buying them back.

That is exactly the Fama–French reply. They argue that Meyers was looking at aggregate data during a period of relatively low net issuance. But when one asks how many firms actually issue equity in a given year, the picture changes. Shiller quotes their result that 67% of U.S. firms issued new shares each year during 1973–1982, and 86% did so during 1993–2002. So the stock market still matters. It matters for financing, and it matters for price discovery even when mature firms are not constantly flooding the market with new stock.

Once that question is settled, the lecture turns back from market-wide challenge to boardroom behavior. Even if firms retain earnings, borrow, issue, and repurchase in different mixtures, someone still has to decide what dividend to pay. That is where John Lintner enters. The route is important: the mathematics comes only after the institutional behavior has been described.

Lintner, a Harvard Business School professor and one of the early figures behind the capital asset pricing model, interviewed managers and directors about dividend decisions. He found a great deal of verbal complexity. Boards want the company to be well regarded. Mature firms take pride in paying dividends. Above all, boards dislike cutting dividends once they have started. That single behavioral regularity does a great deal of work in the lecture.




Figure 9.2: Lintner dividend adjustment on the board.

The board image is valuable precisely because it catches the equation in mid-emergence. On the left the earlier labels “common,” “preferred,” and “John Lintner” are still visible from the surrounding discussion. In the center we can clearly read $\text{Div}_t - \text{Div}_{t-1}$ and the beginning of the right-hand side, $= \rho$. The full equation is not completely legible on the board, so the displayed formula below is a cautious transcript-assisted reconstruction:

$$\text{Div}_t - \text{Div}_{t-1} = \rho(\tau \text{EPS}_t - \text{Div}_{t-1}) + \varepsilon_t, \quad (9.16)$$

with

$$0 < \tau < 1, \quad 0 < \rho < 1.$$

The interpretation is simple and powerful:

1. the board has a target payout ratio, so the desired dividend level is roughly τEPS_t ;
2. the board dislikes cuts, so it does not jump all the way to target when earnings rise;
3. instead, dividends move only part of the way toward target, at speed ρ ;
4. the disturbance term ε_t absorbs the boardroom complexity that the model does not explicitly represent.

This is the lecture’s clearest mathematical core. The equation does not drop from the sky. It comes from an institutional fact: firms fear the reputational cost of cutting dividends. If earnings surge, the board raises dividends cautiously, because an aggressive increase today may force an embarrassing cut tomorrow. Shiller also notes that payout ratios have trended downward over long stretches of time. Earlier shareholders wanted a larger immediate fraction of earnings. Modern shareholders are often more willing to let firms retain cash for internal growth.

9.7 Xerox and Microsoft: Book Equity and Market Value

At this point the lecture does what it promised earlier and turns to concrete firms. Xerox comes first. The story is almost mythic: founded in 1906 as the Halloyd Photographic Company, transformed by Chester Carlson's invention of the plain-paper dry copier, and elevated into one of the glamour growth firms of the 1960s and 1970s. Then came digital copying, which Xerox did not own, and the firm nearly collapsed. Ann Mulcahy, hired in 2001, is presented as the executive who kept the company from extinction.

Because Xerox is a public company, its balance sheet is visible in SEC filings. Shiller uses that institutional detail to make a larger point. Going public gives a firm access to public markets, but it also subjects it to continuous disclosure. His own private company never had to publish a quarterly balance sheet on `sec.gov`. Xerox did.

The balance-sheet lesson is basic but essential. Assets are what the firm owns. Liabilities are what it owes. Receivables, inventory, buildings, short-term debt, long-term debt, deferred taxes, and even a small amount of preferred stock all appear in the lecture's summary. What finally matters for common shareholders is the residual, the shareholder's equity, also called stockholder's equity. In 2010 Xerox's shareholder equity was about \$11.9 billion, up from roughly \$4.9 billion. That is the book residual after liabilities are subtracted from assets.

But Shiller is careful. Shareholder equity is an estimate, not a market price. Buildings and equipment are carried at accounting values, and no one knows exactly what they would fetch if sold off in a hurry. So book equity is a liquidation-style benchmark, not the same thing as market value.

For Xerox the market-cap arithmetic is

$$\text{Xerox market cap} \approx (1.4 \text{ billion shares})(\$11.80) \quad (9.17)$$

$$\approx \$16.4\text{--}16.5 \text{ billion.} \quad (9.18)$$

So the market valued Xerox above its book shareholder equity of about \$11.9 billion.

Microsoft sharpens the contrast. Founded in 1975 by two young founders, it had by 2010 assets of about \$92 billion, including a great deal of cash and short-term investment, some real estate and equipment in Redmond, and even about \$10 billion of stock holdings. On the liability side Shiller mentions ordinary items such as accounts payable, taxes, and unearned revenue, the latter being money already received on products for which future service is still owed. After subtracting liabilities, Microsoft had shareholder equity of about \$48 billion.

Yet the market capitalization was far larger:

$$\text{Microsoft market cap} \approx (8.497 \text{ billion shares})(\$26) \quad (9.19)$$

$$\approx \$221 \text{ billion.} \quad (9.20)$$

That is more than five times book shareholder equity. Shiller even suggests that Microsoft lends some support to the pecking-order intuition: for a long time it retained cash and paid no dividend, and investors were willing to accept that because they believed in the firm's opportunities and management.

9.7.1 Question & Answer

Question. Why can market capitalization exceed shareholder equity on the books?

Answer. Because a public company is not purchased merely as scrap. If we buy the company in order to shut it down, sell the buildings, collect the receivables, unload the inventory, and pay off the debts, then book shareholder equity is the rough benchmark for what might be left. But if we buy the company as an ongoing concern, then we are buying more than booked assets. We are buying future profits, organizational capital, reputation, installed customer relationships, and growth opportunities.

That is why Xerox can sell above its book equity even after severe competitive trouble. The market is saying the company is worth more alive than dead. And that is why Microsoft can sell vastly above its book equity. Its value lies overwhelmingly in future operating profits, not in liquidation of the Redmond campus and short-term securities alone.

The lecture also points out the converse. If market capitalization fell far below shareholder equity, the firm could become vulnerable to takeover and liquidation. Buy the company cheaply, shut it down, sell the assets, pay the debts, and keep the difference. That is why CEOs and boards watch market capitalization so closely against book equity. Still, the comparison is never exact, because book shareholder equity is itself an accounting estimate. It is informative, but it is not a literal auction value observed every day.

9.8 Summary

The lecture unfolds in a very deliberate order. We begin with the corporation as an organizational invention, one that lets people combine money, labor, incentives, and conflict resolution inside a durable legal form. We then widen to the world scale of listed equity, but only long enough to see both its importance and its limits. After that we rebuild the institution from the inside: legal personhood, public versus private equity, one-share-one-vote, boards, CEOs, and the contrast between for-profit and nonprofit corporations.

Only then do we turn fully to valuation. Market capitalization is price times shares outstanding. Ownership is a fraction, n_t/N_t , not a raw share count. Common stock is ultimately a claim on dividends, even if the board delays those dividends for years. The ex-dividend price fall, dilution arithmetic, buybacks, leverage, and preferred stock all become simple once we keep the residual nature of common equity in view. The stock market still matters even for big firms, because net issuance is not the same thing as inactivity, and price discovery remains central. Lintner's model then distills messy boardroom judgment into a partial-adjustment rule for dividends. Xerox and Microsoft close the chapter by showing why book shareholder equity and market capitalization can differ so sharply. The corporation is not just a legal shell and not just a balance sheet. It is a going concern whose claims are priced through expectations about future distributions.

Chapter 10

Real Estate

These notes follow Robert J. Shiller’s lecture as it unfolds in real time. We begin where he begins: with a roadmap. First comes the long history of mortgage lending; then commercial real estate finance; then residential real estate finance; and, if time permits, the larger architecture of securitization. That order matters. The lecture does not treat real estate finance as a list of institutions. It treats it as a chain of problems: first the existence of collateralized lending, then the legal preconditions for scaling it, then the ownership forms used in practice, then the redesign of the home mortgage, and finally the system-wide consequences of turning mortgages into securities.

10.1 Mortgage Lending in Long Historical Perspective

The lecture opens by widening the frame. Real estate finance, we are told, is about financial contracts that use real estate as collateral, and that subject is easier to understand if we refuse to start in the twentieth century. Shiller therefore begins with the word *mortgage*. He cautiously glosses an old Latin ancestor as a “death pledge,” and then notes that medieval French usage fed into the English word. The exact philology is not the main point here; the main point is that the institution is old.

The lecture then pushes the history much farther back than medieval England. Drawing on the work of Yale historian Valerie Hansen, Shiller points to trade documents from the Tang period and to Sogdian materials from the Silk Road world. Those records suggest that loans secured by collateral are more than a thousand years old. In some of the cases discussed in the lecture, the collateral could include property and even persons held in slavery, together with contractual obligations to maintain that collateral. The moral content is disturbing, but analytically it serves a purpose: the mortgage is not a modern invention.

That is why the lecture immediately introduces a contrast. Mortgage lending is ancient, but large-scale modern mortgage finance is not. The idea of collateral has been around for a very long time. What arrived much later was the institutional environment that could support a broad and trusted mortgage market. That distinction carries us into the next step.

10.2 Property Rights as the Precondition for Mortgage Markets

If mortgages are so old, why did modern mortgage markets develop so late and so unevenly? Shiller's answer is that the bottleneck was not conceptual novelty. It was legal and administrative clarity.

The lecture's most vivid example is a Hartford Courant advertisement from 1778. Elisha Cornwell sold his farm on mortgage terms for 800 pounds. The buyer did not fully pay him, yet proceeded to mortgage the same farm again, first for 880 pounds and later for 1,000 pounds. Cornwell had to warn later victims in the newspaper. That detail is important. It shows just how underdeveloped mortgage institutions were. The problem was not merely that one man behaved badly. The deeper problem was that there was no systematic public way to determine who actually owned the farm.

So we should not summarize the point too quickly as "property rights matter." The lecture is more concrete than that. Mortgage lending requires a title system that lets strangers verify ownership, and it requires foreclosure rules that let lenders recover the collateral if the borrower does not pay. Without those two things, the contract cannot be scaled.

That is why Shiller dwells on the Prussian *Grundbuch*. In 1872, Prussia created a central book of ownership, and in 1897 the institution became national within Germany. The lecture uses that example to mark a turning point: once the state provides a reliable registry, mortgage lending can move beyond local trust networks and become a national financial business.

The argument then broadens again. Hernando de Soto's *The Mystery of Capital* is brought in not as a side citation but as evidence that the same problem persists in much of the developing world. One may be able to go to a village and ask around who owns a parcel of land, but a serious financial transaction cannot rest on village hearsay. Nor can mortgage lending flourish if foreclosure takes years or if courts do not function well enough to make collateral recovery credible.

10.2.1 Question & Answer

Question. Why can mortgage lending not rest on local knowledge or hearsay about ownership?

Answer. Because a mortgage is not merely a social understanding. It is a priced contract. The lender must know that the borrower really owns the property and that this fact can survive dispute, resale, and enforcement. If ownership is known only informally, the collateral is unreliable. If foreclosure is legally possible on paper but practically impossible in court, the collateral is again unreliable. In that case the lender will either refuse to lend or charge terms severe enough to make the market thin. The lecture's claim is therefore stronger than a slogan: mortgage markets scale only when ownership and repossession become legible and enforceable.

There is a hard institutional symmetry here. It may seem cruel to throw someone out of a house after default, and the lecture acknowledges that moral discomfort. But Shiller immediately adds the lender's side: if the house cannot be recovered, no one will make the mortgage in the first place. That is part of the same development story. The growth of mortgage finance in the twentieth and twenty-first centuries depends on improvements in both title and enforcement.

10.3 Commercial Real Estate Before Democratization

With that long history in place, the lecture narrows to the United States. Shiller tells us explicitly that finance tends to standardize within countries: laws, regulators, and public expectations push people toward a small number of accepted contract forms. Innovation is therefore slow, and this is already one of the lecture's recurring themes.

He then asks a practical question: who owns the commercial buildings around us? The answer is often not a public corporation. In the lecture, the typical answer is a partnership.

The reason is tax treatment. Corporations are subject to corporate tax, and their owners are taxed again on the income they receive. Partnerships allow income to pass through more directly. If one owns a building through a partnership rather than through a corporation, one can avoid this double taxation. That is why so much stand-alone commercial real estate has been organized in partnership form.

The lecture gives that partnership form a more specific name: the direct participation program, or DPP. A DPP is direct in the sense that the investor participates directly in the profits of the real-estate venture as a partner rather than as a shareholder. It is typically not offered to the general public. Instead, it is restricted to accredited investors, whom the lecture describes roughly as wealthy investors meeting SEC wealth or income thresholds.

Three structural features matter in the lecture's account. First, the DPP is usually project-specific. Second, it has a limited life rather than perpetual duration. Third, it often distinguishes between general partners and limited partners. The general partner manages the business and lacks limited liability in the strongest sense; the limited partners are passive investors and enjoy limited liability.

That structure fits a building rather well. A group can acquire or develop a property, manage it for a period, depreciate it over the life of the venture, and then sell it and wind the vehicle down. The lecture's local example is 360 State Street in New Haven: Shiller does not claim to know its actual financing structure, but uses it as a plausible illustration of how a real-estate partnership could be organized around one building, with a plan for operation and an eventual exit.

So this part of the lecture does more than describe a niche legal form. It reveals a tension that will matter later. The tax-efficient structure for commercial real estate is also a structure that excludes most ordinary investors.

10.4 REITs and the Democratization of Commercial Real Estate

That tension is the entry point for REITs. Shiller explicitly frames them as part of the democratization of finance.

The political complaint was easy to state. If partnerships and DPPs are tax-efficient, why should access be limited to accredited investors? Why should most households be confined to ordinary corporate forms while a narrower wealthy group receives the better tax treatment available in commercial real estate? Congress's answer, in 1960, was to create the real estate investment trust, or REIT.

The lecture presents the REIT as a compromise. It opens real-estate investing to the broader public, but only by imposing a narrow legal definition of what the vehicle may do. In the lecture's compressed summary, a REIT must satisfy a set of threshold conditions:

- at least 75% of assets must be in real estate or cash;
- at least 75% of income must come from real estate;
- at least 90% of income must come from real-estate dividends, interest, and capital gains;
- at least 95% of income must be paid out;
- no more than 35% of income may come from sales of properties held for less than four years.

The point of these thresholds is not just definitional neatness. Once Congress creates a vehicle that avoids ordinary corporate taxation, every company would like to qualify. So the state confines the privilege to a carefully delimited real-estate form. One should read these figures here as lecture-era statutory conditions, not as a polished current-law summary.

10.4.1 Question & Answer

Question. If partnerships avoid double taxation, why not simply let everyone invest through partnerships directly?

Answer. Because the government does not want the partnership form to become a universal substitute for the corporation. The partnership form is therefore hemmed in by restrictions: limited life, project-specific structure, access limits, and distinctions between active and passive partners. The REIT is a different compromise. It extends access to small investors, but only by binding the vehicle to real estate, payout rules, and other statutory constraints. In the lecture's terms, it is democratization, but democratization through a regulated legal form.

Shiller then slows us down again with a historical theme he clearly cares about: innovation in finance is slow. REITs are created in 1960, but they do not immediately dominate the market. The lecture describes a succession of booms. The first comes in the late 1960s, when regulated deposit-rate ceilings make REIT payouts comparatively attractive. Another comes after the tax reform of 1986, which removes some of the tax attractions of DPPs and partnerships. A third expansion comes with the real-estate boom of the 1990s, when specialized REITs become more visible and the form finally becomes broadly familiar.

That long gestation matters because the lecture is making two parallel claims. One is that innovation is slow and path-dependent. The other is that finance tends, over long stretches of time, toward wider participation.

10.5 The Great Depression and the Redesign of the Home Mortgage

Only after commercial real estate has been organized in these terms does the lecture pivot to residential real estate. The pivot is motivated by scale. Housing is socially larger and financially more pervasive than commercial real estate. In the United States, the lecture gives a rough homeownership rate of about two-thirds, and treats that fact as partly the product of public policy that encourages mortgage lending.

To explain why the American home mortgage took its modern form, the lecture turns to the Great Depression. This is not just historical background. It is the crisis that exposes a flaw in the earlier contract design.

Before the Depression, mortgages in the United States were often short, typically two to five years, and they were balloon-payment mortgages. The lecture's arithmetic example is:

$$H_0 = \$10,000, \quad B_0 = \$5,000, \quad D_0 = \$5,000, \quad (10.1)$$

where H_0 is the house value, B_0 the mortgage balance, and D_0 the down payment. The borrower pays monthly interest, but the principal is largely untouched until maturity:

$$x_t = \frac{r}{12} B_0 \quad \text{for } t = 1, \dots, T-1, \quad B_T = B_0. \quad (10.2)$$

The whole arrangement depends on refinancing. When the balloon date arrives, the borrower expects to return to the bank and roll the loan.

The Depression breaks that expectation. Unemployment rises to about 25%. House prices fall sharply, often by more than half. In the lecture's example, a house bought for \$10,000 with a \$5,000 mortgage may fall to about \$4,000, so that

$$H_1 \approx \$4,000 < B_0 = \$5,000. \quad (10.3)$$

Now the borrower arrives at the refinancing date with exactly the wrong balance sheet: weak income, weak collateral, and a full principal balance still due. The bank refuses renewal. The family is then pushed toward bankruptcy or distress sale, and the original down payment is lost. The lecture is careful here. The problem is not merely that borrowers stopped paying. The deeper problem is that the contract hid its fragility in the refinancing date.

This is why the Roosevelt administration treats the old mortgage form itself as defective. The Federal Housing Administration, created in 1934, begins insuring mortgages. Mortgage insurance means that if the borrower defaults and the house value is insufficient, the government makes up the lender's shortfall. At the same time, FHA-insured mortgages are required to be long-term and non-balloon. The lecture first emphasizes 15-year amortizing mortgages, and then notes that by the early 1950s the FHA is pressing toward 30-year mortgages.

Shiller pauses over the word *amortize*: it means to pay down the balance. An amortizing mortgage therefore has no giant payment at the end. The family makes a fixed monthly payment, and when the term is over, the balance is zero. That is the conceptual payoff of the crisis story. The redesign is not cosmetic. It moves the risk out of a catastrophic lump-sum date and spreads it over a long, predictable payment stream.

10.6 The Mathematics of the Fixed-Rate Amortizing Mortgage

At this point the lecture deliberately slows down. Shiller says that the arithmetic of the amortizing mortgage can be confusing, and that in the 1930s it took real explanation. This is the mathematical center of the chapter.

We take the notation directly from the lecture. Let m be the maturity in months, let r be the annual mortgage rate, let B_0 be the initial mortgage balance, and let x be the fixed monthly payment. The contractual condition is that the present value of all monthly payments equals the amount initially borrowed. In clean notation, the lecture's formula is

$$B_0 = \sum_{t=1}^m \frac{x}{(1 + r/12)^t} \quad (10.4)$$

$$= x \sum_{t=1}^m (1 + r/12)^{-t} \quad (10.5)$$

$$= x \frac{1 - (1 + r/12)^{-m}}{r/12}. \quad (10.6)$$

This is simply the annuity formula, written with the monthly discount rate $r/12$. Shiller briefly self-corrects in the lecture after an oral slip; the correct rate here is the monthly rate $r/12$, and that is the form we keep.

Solving for the payment gives

$$x = B_0 \frac{r/12}{1 - (1 + r/12)^{-m}}. \quad (10.7)$$

That tells us the constant monthly payment. But the lecture is equally interested in something subtler: the internal composition of that payment. Each month's payment splits into interest and principal:

$$x = I_t + P_t, \quad (10.8)$$

$$I_t = \frac{r}{12} B_{t-1}, \quad (10.9)$$

$$P_t = x - I_t, \quad (10.10)$$

$$B_t = B_{t-1} - P_t. \quad (10.11)$$

This recursion is the essential mechanism. Interest is computed on the remaining balance B_{t-1} . As the balance falls, interest falls. Since the total payment x is fixed, the principal component must rise. The household sees a stable monthly bill, but internally the contract is steadily shifting from interest to amortization.

The lecture then produces visual evidence from an old mortgage table. Before computers, lenders relied on printed tables to carry out exactly this arithmetic.

The projected page is partly blurred, but its structure is clear. The headers YRS, MOS, INTEREST, PRINCIPAL, and BALANCE repeat across the page so that a long schedule can be printed compactly. The left side shows early months, when interest is relatively high and the balance is still close to the original loan amount. The right side shows later months, where interest is smaller, principal is larger, and the balance has fallen substantially. Shiller's pointing gesture is part of the evidence here: he is drawing attention to a later row precisely to show the declining interest share.

The lecture's worked example uses a 10-year mortgage at 5%, with the table scaled per \$1,000 of initial balance. Thus

YRS	MOS	INTEREST	PRINCIPAL	BALANCE
0	1	4.17	6.44	993.56
0	2	4.14	6.47	987.09
0	3	4.11	6.50	980.59
0	4	4.09	6.52	974.07
6	6	≈ 1.74	≈ 8.87	≈ 407.61

Table 10.1: A compact reconstruction of the lecture’s 10-year, 5% amortization example per \$1,000 borrowed. The early rows are visible in the projected table; the later checkpoint follows the lecture’s spoken arithmetic and should be read cautiously because the projected page is partly obscured.

10.6.1 Question & Answer

Question. If the monthly payment is fixed, why does the interest share fall while the principal share rises?

Answer. Because the contract computes interest on the remaining balance, not on the original balance. Early in the mortgage, the remaining balance is large, so the interest charge is large. Later in the mortgage, the balance is smaller, so the interest charge is smaller. Since the total payment x is fixed, every fall in I_t automatically shows up as a rise in P_t . This is the core time structure of the amortizing mortgage: a constant total payment with a changing internal composition.

That changing composition is what makes the fixed-rate amortizing mortgage both safer and easier to live with. The family no longer faces a giant balloon date. Instead, it faces a predictable monthly obligation, while the contract quietly pays itself down in the background.

The lecture now gives the instrument its standard name: the conventional fixed-rate mortgage. But Shiller does not stop there. He immediately turns success into a puzzle. Why did such a sensible contract become dominant in so few countries?

10.7 From Fixed-Rate Mortgages to Securitization, Crisis, and Reform

The lecture’s first answer to the international puzzle is that the fixed-rate mortgage is costly to supply. If a lender guarantees a rate for 20 or 30 years, that promise has value and must be priced. Borrowers may resist the higher rate. Regulators may also resist the maturity mismatch it creates: banks fund themselves with deposits that can leave quickly, but the mortgage asset stays on the books for decades. In a run, the bank cannot liquidate that asset cheaply. So the long fixed-rate mortgage is not just a household contract. It is a balance-sheet problem and a regulatory problem.

Shiller uses other countries to sharpen that point. Canada, for example, did not make the long-term fixed-rate mortgage standard in the same way as the United States. He mentions a Canadian episode around 1980, when rates rose sharply and refinancing became painful for many households. He also mentions Alastair Darling’s call for longer-term mortgages in the United Kingdom after the crisis. The underlying theme is that the American mortgage form, however useful, is institutionally contingent.

The lecture then turns from the mortgage contract itself to the state-supported market that grew

around it. In 1938, the federal government created what became known as Fannie Mae, initially as a government agency that bought mortgages in order to support lending and stimulate the depressed housing sector. The economic logic is simple: when a bank makes a mortgage and then sells it, it gets cash back and can lend again. A secondary market therefore supports the primary market.

Later Freddie Mac is created. Shiller presents Freddie as initially associated more explicitly with mortgage securitization: buying mortgages, packaging them into securities, and selling them with a guarantee. Over time both Fannie and Freddie operate in that space. So the mortgage moves one step farther away from the household. It becomes the raw material of tradable securities sold to investors.

10.7.1 Question & Answer

Question. Why were Fannie Mae and Freddie Mac rescued after repeated claims that they were not government-guaranteed?

Answer. Because the lecture's distinction between explicit and implicit guarantees collapses in crisis. Formally, officials could say that the government did not stand behind these private corporations. But investors around the world treated them as government-linked, and the United States benefited from that belief. Once the crisis of 2008 arrived, letting them fail would have damaged not only mortgage markets but also the perceived safety of U.S. financial commitments more broadly. The legal disclaimer did not survive the reputational and systemic cost of enforcing it.

This is one of the lecture's most memorable institutional lessons. One can deny a guarantee repeatedly. But if the state has created the institution, allowed the market to rely on it, and let the exposure become enormous, the denial may not be credible when the crisis finally arrives.

The lecture briefly compares Canada again at this point. Canada has a government-owned housing-finance institution, the Canada Housing and Mortgage Corporation, which does some work analogous to the FHA and the U.S. secondary-market agencies. But Shiller's point is that Canada did not let that support structure become as large or as aggressively expansionary as the United States did. That difference helps explain why the Canadian housing bubble was less extreme.

The final movement of the lecture broadens into the full securitization chain. The textbook, Shiller says, discusses securities such as collateralized mortgage obligations (CMOs) and collateralized debt obligations (CDOs). A CMO is a mortgage-backed structure that divides cash flows into tranches, especially by prepayment exposure. A CDO is a further layer, often holding mortgage-linked debt instruments as assets. By the time we reach this stage, the original home mortgage sits inside a much more complicated financial machine.

The lecture insists on the sequence of intermediaries. Someone originates the mortgage by dealing directly with the homeowner. That mortgage is then sold to an investor or securitizer. A servicer may collect the payments and deal with delinquency. The mortgage pool can then be reorganized into CMOs, and those securities can in turn become inputs into CDOs. What appears at the household level as one mortgage has become, at the system level, a chain of layered claims held all over the world.

This is where the lecture places the moral-hazard problem. If the originator knows the mortgage will be sold immediately, the originator may become cynical about repayment quality. Shiller gives

a blunt example: a mortgage broker may encourage a family to overstate income because the broker expects to pass the mortgage on and avoid the loss. That is why the post-crisis reform he highlights is so simple in form and so revealing in purpose. If M denotes a mortgage pool, then the rule is

$$\text{originator retains } 0.05M, \quad \text{originator sells } 0.95M. \quad (10.19)$$

The idea is to keep some risk with the party that screened the borrower. Europe moved in that direction, and the United States incorporated the same 5% principle into Dodd–Frank. The lecture also notes related professionalization, such as stricter licensing of mortgage brokers.

The ending is deliberately guarded. Shiller does not say that we should abolish securitization and go back to 1778. He says instead that the basic system has been retained, while reforms try to reduce the worst incentive failures. That is a more modest conclusion, but it is also a more faithful one.

10.8 Summary

The lecture moves in a deliberate arc. Mortgage lending is ancient, but broad mortgage markets are modern because they depend on clear title, enforceable collateral, and administratively legible ownership. Commercial real estate then shows us how tax law and regulation shape the ownership form, first through partnerships and DPPs and later through REITs as a democratizing compromise. Residential finance raises the stakes: the Great Depression reveals that the old balloon mortgage hid a dangerous refinancing risk, and the FHA-backed shift to long-term amortizing mortgages redesigns the household contract. The mathematical core of that redesign is an annuity-pricing equation together with a simple monthly recursion for interest, principal, and balance. Finally, the lecture zooms back out. The fixed-rate mortgage becomes an international puzzle, then the raw material for securitization, implicit guarantees, global investment, crisis, and reform. The result is a chapter in which a household payment stream and a world financial system turn out to be two views of the same contract.

Chapter 11

Behavioral Finance and the Role of Psychology

In this lecture Robert J. Shiller asks us to widen finance until it includes the actual people who inhabit it. Rationality remains useful; the lecture never denies that. But rationality is only a partial language for institutions whose success depends on minds, moods, habits, framing, reputation, and character. Behavioral finance therefore appears here not as a rejection of finance theory, but as a correction to an overnarrow picture of human conduct.

These notes follow Lecture 11 of the Yale Financial Markets series, curated by LazyingArt LLC. The lecture unfolds in a deliberate order. It begins with a moral and institutional preface, then slows down for the best-known formal core in behavioral economics, and only then accelerates into a wider catalog of recurring psychological mechanisms that matter for markets.

11.1 Behavioral Finance and the Limits of Rationality

Shiller begins by saying quite plainly that behavioral finance has become an important element of modern finance, even if it remains controversial. The controversy is itself revealing. Economists have long relied on rationality as a simplifying principle, and that principle has indeed been useful. But the lecturer's point is that its usefulness is limited. People are often rational, not completely rational. They can be intelligent in many dimensions and still make serious mistakes.

That point matters in finance because financial institutions are designed for real people, and their functioning depends on the behavior of real people. The lecture places this alongside a broader scientific development: a revolution in neuroscience. The human brain is not simple, so one should not expect the economy built out of human decisions to be simple either.

Already in the opening minutes the lecture introduces the darker implication. We are aware of one another's weaknesses, and that awareness can create an impulse to exploit. If we see somebody behaving foolishly, we may be tempted to turn that error to our own advantage. That is the danger that motivates the entire lecture. But before the lecturer catalogs any particular bias, he insists on placing that danger in perspective.

11.2 Morality, Exploitation, and Character

The lecture's first real pivot is moral. We are told to expect a discussion of human failings, but we are also told not to begin from cynicism. The business world, Shiller says, does not exploit human weakness as relentlessly as one might expect. Successful businesses have long-run interests. They care about reputation. A strategy that is blatantly exploitative may generate short-run gain, but it works against durable advantage.

He adds a second restraint: morality itself is part of human behavior. On the lecturer's telling, we do not merely maximize against one another. We also have an evolved impulse to be moral, and that impulse limits how much satisfaction we get from exploiting the mistakes of others. The lecture never says manipulation is absent. It says the market system contains more moral and institutional restraint than a shallow cynicism would predict.

11.2.1 Question & Answer

Question. If people are imperfect and manipulable, why are financial institutions not simply machines for exploiting them?

Answer. The lecture's answer is cumulative rather than doctrinaire. Exploitation is checked by long-run reputation, by repeat business, by trust, by regulation, and by the ordinary desire not to be the sort of person who preys on others' mistakes. Opportunities for exploitation are real, but they do not exhaust the structure of finance. This is why the lecture can honestly turn to human error without concluding that markets are nothing but organized predation.

To put that optimistic counterweight on firmer ground, the lecture turns to Adam Smith.

11.2.2 Adam Smith and Praiseworthiness

Shiller returns not only to *The Wealth of Nations* of 1776, but to the earlier *Theory of Moral Sentiments* of 1759. That move is essential to the lecture's rhythm. Before we see any value function or weighting function, we are asked how an economy can work if people are selfish, vain, and sometimes manipulative.

Smith's answer, as presented here, begins with praise. Human beings want approval. We crave the good opinion of others. But Smith's deeper observation is that mature people do not stop there. They move from wanting praise to wanting praiseworthiness. We do not merely want to be admired; we want to deserve admiration. That is the step that makes an economy workable.

The lecture lingers on Smith's example of mathematicians. Mathematicians are usually obscure. The public does not understand what they do. Yet they continue doing it. Why? Because the motive is not public applause as such. The motive is to be the kind of person whose work is genuinely worthy. Shiller uses this example to argue that finance, like other mature professions, depends heavily on people who have internalized that desire for deserved approval.

The claim is not naively idealistic. The lecturer says explicitly that not everyone reaches this mature state. That is why he next introduces personality psychology and the problem of character.

11.2.3 Personality Differences and Manipulability

The discussion of the DSM-IV appears in the lecture for a very specific reason: to complicate the Smithian picture without abandoning it. People are not all the same. Some are more manipulative, more deceptive, and less constrained by remorse than others. The lecture's rough prevalence claims should be quoted cautiously, because the transcript is noisy at this point, but the magnitudes are clear enough to preserve:

$$\Pr(\text{APD} \mid \text{male}) \approx 0.03, \quad (11.1)$$

$$\Pr(\text{APD} \mid \text{female}) \approx 0.01, \quad (11.2)$$

$$\Pr(\text{APD} \mid \text{prison inmate}) \approx 0.40. \quad (11.3)$$

The lecture's descriptive list is more important than the exact percentages: lack of remorse, frequent lying, lack of empathy, superficial charm, shallow emotion, and a manipulative relation to others. These observations are not the mathematical spine of the chapter, but they do explain why the lecturer refuses to rest content with general moral optimism. Finance is populated by real mixtures of character.

From there the lecture broadens again. We are all manipulable, not only by extreme personalities but by ordinary commercial tactics. The example of pricing points, such as \$9.99 instead of \$10.00, is intentionally simple. It is not cited as a moral catastrophe. It is cited to show how competition itself rewards psychological leverage. The lecturer's conclusion is subtle: individual acts may look somewhat manipulative, yet what matters in the long run is the total character of the person or institution through time.

With that introduction behind us, the lecture finally turns to the famous formal core.

11.3 Prospect Theory: Setup and Board Architecture

The lecturer says he wants to start with what is probably the most famous element of behavioral economics: prospect theory. It was introduced by Daniel Kahneman and Amos Tversky as a theory of how people make decisions about prospects, that is, gambles, under uncertainty. The lecture also warns us not to overread what follows. There is a huge literature here, and what we are about to see is only a quick sketch.

Prospect. A prospect is a gamble, or more generally a choice problem under uncertainty.

Prospect theory. In the lecture's own decomposition, prospect theory has two parts: a value function, which describes how gains and losses are experienced, and a weighting function, which describes how probabilities are psychologically processed.

The retained board image is important because it preserves the lecture's pacing. On the left the board reads "Value Function" and "Weighting Function?"; on the right it reads "Prospect theory" and "Kahneman & Tversky." Nothing has yet been turned into a polished formula. The lecture first names the two pieces, then says, "I'll draw a picture," and only then begins the mathematical interpretation.




Figure 11.1: Blackboard headings introducing prospect theory, the value function, and the weighting function.

11.4 The Value Function

Now the lecture slows down visibly. The lecturer draws a vertical line and a horizontal line through the middle of the board. The horizontal axis is wealth, or money, or something like that, with zero in the middle. The vertical axis is value, something like utility. The left side is losses; the right side is gains.

For note-writing purposes it is convenient to introduce notation that the board itself does not write explicitly. Let W denote wealth and let r denote the psychological reference point. We then write

$$x = W - r. \quad (11.4)$$

This is a cautious reconstruction for clarity. The board itself labels regions as losses and gains rather than introducing a symbol x .

The lecture's qualitative claims can now be summarized cleanly. The curve passes through the origin, the gain side is concave, the loss side is convex, and there is a kink at the origin:

$$v(0) = 0, \quad x < 0 \Rightarrow \text{losses}, \quad x > 0 \Rightarrow \text{gains}. \quad (11.5)$$

$$v''(x) > 0 \quad \text{for } x < 0, \quad v''(x) < 0 \quad \text{for } x > 0. \quad (11.6)$$

We should be careful about what is and is not being claimed. No closed-form formula is written on the board. What is written, and then verbally interpreted, is a qualitative graph. On the gain side the curve is like diminishing marginal utility in ordinary economic theory. On the loss side it bends the other way. The value of gains and losses is therefore not linear.




Figure 11.2: Prospect-theory value function with losses on the left, gains on the right, and value on the vertical axis.

11.4.1 Question & Answer

Question. What does the kink at the origin mean, and why does the reference point matter?

Answer. The reference point is the psychological zero from which gains and losses are measured. Usually it is today's wealth, but the lecture insists that it need not be fixed. It can be shifted by framing, that is, by the way a prospect is presented. Once the reference point is set, the kink says that small movements around that point matter disproportionately. We react strongly to little losses near zero. That is why framing changes behavior: the same objective prospect, described relative to a different reference point, can produce a different decision.

A cautious note-writing summary of the lecture's verbal asymmetry is

$$-v(-\varepsilon) > v(\varepsilon) \quad \text{for small } \varepsilon > 0. \quad (11.7)$$

This inequality is not written on the board, but it captures the lecturer's explicit claim that we are "spooked" by small losses and "less encouraged" by small gains.

The lecture's own example is a loss of \$5 on the way to class. Relative to lifetime wealth, \$5 is negligible. Relative to the reference point of "my money this morning," it is vivid and painful. The psychological measurement is local, not actuarial.

From the graph to an insurance market. The lecture immediately turns the graph into an institutional argument. The logic runs as follows.

1. Place the origin at the current reference point, usually today's wealth.
2. Read the kink as heightened sensitivity to small losses near that point.




Figure 11.3: Clean qualitative redraw of the prospect-theory value function based on the board sketch.

3. Notice that firms can identify small, salient, bounded losses that consumers are actively imagining.
4. Sell insurance on those narrow losses rather than on the household's largest lifetime risks.
5. The lecture's main examples are funeral insurance, airline flight insurance, and diamond-ring insurance.

This is one of the lecture's sharpest claims. The insurance people *ought* to buy is concentrated on truly large losses, such as the death of a wage earner. But the insurance that often sells most easily is insurance attached to a narrow, emotionally salient event. Funeral insurance has been sold since ancient Rome precisely because it isolates a visible expense that people can picture. Diamond-ring insurance works similarly. It is not the household's central financial vulnerability, but it is vivid.

Framing belongs here as well. Kahneman and Tversky's point, as Shiller states it, is that the same prospect can be presented in different language so as to suggest a different reference point. The prospect itself has not changed. The behavior does.

The lecturer also adds an important qualification before moving on. These distortions are real, but not total. People eventually learn. They come to distrust crude manipulation. So even here, at the most formal point in the lecture, the narrative is still double-sided: the bias is real, and yet the institution adapts around it.

11.5 Weighting Probabilities and Rare Events

The second half of prospect theory concerns not value but probability. The lecture's point is that objective probabilities and psychological decision weights are not the same thing. A probability is a number between 0 and 1, but human beings do not naturally think in a smooth continuum of such numbers.

For note-writing purposes we introduce a weighting function $w(p)$, again with the warning that the

lecture gives only a qualitative sketch:

$$p \in [0, 1], \quad w(0) = 0, \quad w(1) = 1. \quad (11.8)$$

$$w(p) \approx 0 \quad \text{for very small } p, \quad w(p) \approx 1 \quad \text{for very large } p. \quad (11.9)$$

The crucial part is the middle. Very small probabilities may be rounded down to zero. Very large probabilities may be rounded up to one. But if a small probability is not rounded away, it may be exaggerated sharply. The lecture describes the simplest version of the weighting function as almost a broken-line psychological sketch, not as a smooth calibrated formula. That matters. We should not import a polished textbook parameterization where the lecture gives only a rough intuition.




Figure 11.4: Transcript-based schematic of the lecture’s qualitative weighting function. The retained screenshots do not preserve this graph, so the diagram is an interpretive reconstruction rather than frame-backed evidence.

The lecture’s analogy is memorable: our minds handle probabilities as if there were only a few categories. Roughly speaking, the psychological categories are “cannot happen,” “maybe,” and “will happen.” That is why tiny probabilities can be either ignored or inflated.

The institutional example is airline flight insurance. The crash probability of a commercial flight is extraordinarily small. Most travelers mentally round it to zero and move on. But some do not. For them the event takes on a disproportionately large psychological weight, especially at the moment of boarding. That is why airport vending machines once sold flight insurance at highly unfavorable prices: the machines did not need every traveler, only the subset who could not bring themselves to round the probability down.

The lecture immediately complicates its own example. Those machines have largely disappeared. Why? Because the error is not immutable. People learn, experience accumulates, and trusted intermediaries matter. If my ordinary insurance agent does not recommend a special airport policy, I begin to ignore the vending machine. That is why the lecture briefly invokes Gerhard Gigerenzer and then Nick Barberis’s experiments with highly quantitative students: the tendencies are real, but they are not destiny.

The lecturer’s own transition is telling: prospect theory explains a great deal, but it does not explain everything. At that point he explicitly says, in effect, let me move on.

11.6 Other Behavioral Mechanisms

The rest of the lecture accelerates. It does not try to build one grand system out of the remaining material. Instead it moves through a sequence of mechanisms, each connected to some financial surface.

11.6.1 Regret and Gambling

Regret theory is introduced first. It is related to prospect theory, but it isolates something more specific: the painful experience of realizing that one has made a mistake. The lecture's phrase is the ordinary one, "I was kicking myself." Regret is therefore not just a bad outcome. It is a bad outcome with ego attached to it. That can distort decision-making, because we may avoid actions whose later failure would feel humiliating.

The next topic is gambling. The lecture treats gambling as a human universal rather than as a narrow pathology. Anthropologists find it in every society, and the lecture quotes rough prevalence figures from one U.S. study:

$$\Pr(\text{gambled at least once in the year}) \approx 0.61, \quad (11.10)$$

$$\Pr(\text{compulsive gambler} \mid \text{male}) \approx 0.011, \quad (11.11)$$

$$\Pr(\text{compulsive gambler} \mid \text{female}) \approx 0.005. \quad (11.12)$$

Again the numbers are supportive rather than central. The real point is that gambling is psychologically connected to competence, self-image, stimulation, and even relief from depression. In that sense it can resemble an addiction. Yet Shiller does not reduce the stock market to pathology. The same appetite for stimulation can be channeled into productive enterprise. That is one of the lecture's recurring moves: a dangerous psychological impulse can sometimes be organized into something socially useful.

11.6.2 Overconfidence: a Worked Example

The lecturer then slows down again for a classroom experiment. Overconfidence is not presented as mere swagger. It is operationalized through the notion of a 90% confidence interval. If a subject reports an interval $[L, U]$ for an unknown quantity X , then the intended meaning is

$$\mathbb{P}(L \leq X \leq U) = 0.9. \quad (11.13)$$

In ordinary language, if we keep giving honest 90% confidence intervals, we should be right nine times out of ten.

The lecture asks for three such intervals: one for world population, one for the mass of the Earth, and one for the number of languages in the world. The stable numerical answers preserved in the transcript are

$$X_{\text{world population}} = 6,901,330,581, \quad (11.14)$$

$$M_{\text{Earth}} = 5.974 \times 10^{24} \text{ kg}, \quad (11.15)$$

$$N_{\text{languages}} = 6909. \quad (11.16)$$

For the Earth-mass question the lecture briefly changes units. The stable conversion is

$$1 \text{ metric ton} = 10^3 \text{ kg.} \quad (11.17)$$

The transcript around the verbal conversion is noisy, so the kilogram figure is the safer quantity to preserve.

Now comes the diagnostic step. The class raises hands if the true value falls inside the interval. The reported frequencies are only rough show-of-hands counts, but they are sufficient:

$$\text{hit rate}_{\text{world population}} \approx 0.8, \quad (11.18)$$

$$\text{hit rate}_{\text{Earth mass}} \approx 0.1, \quad (11.19)$$

$$\text{hit rate}_{\text{languages}} \approx 0.1. \quad (11.20)$$

The inference is straightforward and worth stating explicitly.

1. A nominal 90% interval should contain the truth about 90% of the time.
2. The realized inclusion rates are far below 0.9.
3. Therefore the intervals were too narrow.
4. Therefore subjects knew less than they thought they knew.

That is overconfidence in an operational sense. The lecture's explanation is also important. We tend to think from the first perspective that comes to mind. We do not canvass enough ways in which the world might surprise us. Since imagination is narrower than reality, the interval is too narrow as well.

Shiller then widens the application. Overconfidence does not stop with the self. It extends to our friends, to leaders, to central bankers, and to charismatic CEOs. This is where the lecture invokes Rakesh Khurana's *Search for the Charismatic CEO* and Taleb's *Fooled by Randomness*. Chance success is quickly rewritten as genius, and institutions can then be reorganized around that illusion.

11.6.3 Cognitive Dissonance

Cognitive dissonance is presented as a bias in belief maintenance. It is painful to admit that a belief to which we are attached is wrong, so we selectively seek confirming evidence and evade disconfirming evidence. The lecture's early example is the car experiment: after buying a car, people attend to advertisements for the car they chose and avoid advertisements for the alternatives they rejected.

The same mechanism appears in finance. Shiller cites work by Will Goetzmann and coauthors on mutual-fund investors. When a fund performs badly, many investors do sell. But some hang on, and the striking point is that they often do not even have an accurate impression of how badly the fund has done. Unpleasant evidence is not neutrally processed. It is blocked out.

The lecture then turns to a more institutionally awkward example, drawn from Mullainathan's work on financial advice. Actors presented advisers with extremely concentrated portfolios: all money in money-market funds, or all money in tech stocks, or all wealth in one's employer's stock. Professionally, an adviser should challenge the concentration. In practice many advisers did not, especially when the client's position seemed to express a deep personal conviction.

11.6.4 Question & Answer

Question. If advisers know that clients are biased, when should they challenge those beliefs and when should they work around them?

Answer. The lecture does not offer a formula, and that is precisely the point. If an adviser immediately attacks a new client's deepest convictions, the client may simply leave. The adviser may therefore soften the challenge, delay it, or fail to make it at all. That may look like cowardice, but the lecture treats it as a real professional dilemma. A bluntly correct intervention can be institutionally ineffective. The moral problem is therefore not only what is true, but how truth can be introduced without losing the relationship through which correction might eventually occur.

11.6.5 Anchoring, Representativeness, and Social Contagion

Anchoring comes next. The lecture retells Kahneman and Tversky's wheel-of-fortune experiment. Subjects are first exposed to a random number and then asked a numerical question. The random number has no logical bearing on the answer, yet answers drift toward it. When ambiguity is genuine, judgment is pulled by whatever salient number happens to be nearby.

Representativeness is a different distortion. Here the mind overreads remembered patterns as if they were reliable guides. The lecture's example is the head-and-shoulders pattern in technical analysis. A rare pattern, once noticed, is then seen too often. This creates room for manipulation. If traders believe a head-and-shoulders formation predicts decline, a manipulator in a thin market can engineer price trades that mimic the pattern and then short into the resulting sell signal.

The lecture's final behavioral mechanism is social contagion. This is really social psychology rather than individual psychology. Our beliefs are interdependent. Herd behavior does not usually feel like conscious imitation. It feels like ordinary common sense inside a shared field of opinion. Shiller invokes Durkheim's phrase "collective consciousness" to capture this. The practical conclusion is plain enough: shared narratives and shared memory can generate large market swings.

11.7 Institutions, Shared Value, and Integrity

The lecture ends where it began, but with more detail behind the optimism. There is manipulation. There is exploitation. There are antisocial personalities. There are recurrent errors in value perception, probability weighting, confidence, updating, and social judgment. But these are not the only facts.

Institutions evolve. Some manipulations become crimes. The lecture's own example is the deliberate manufacture of a head-and-shoulders pattern in order to profit from the public's response. That is not merely clever. It is market manipulation, and the lecturer points out that it is prosecuted.

Beyond formal law, there is still the older Smithian theme. People who rise to important positions in finance often do so partly because they have an internal compass. The lecturer explicitly returns to praiseworthiness here. The desire to be the kind of person one ought to praise is not a decorative moral. It is part of the explanation of why financial markets do not become wholly dominated by exploitation.

The lecture's last two references reinforce this. Michael Porter and Mark Kramer are cited for the idea of shared value: mature business leadership does not think only in terms of narrow short-run shareholder maximization. It understands that long-run value is created with society, not against it. Anna Bernasek's *The Economics of Integrity* makes the same point through a different example. We are safe drinking milk not only because of regulation, but because the people producing and handling it usually think it is obvious that one should keep it clean. Integrity is not an afterthought. It is one of the hidden supports of economic life.

That is the lecture's final balance. Behavioral finance matters because people are vulnerable to error. It also matters because finance is an institutional achievement that learns how to live with that vulnerability.

11.8 Summary

We began where the lecture begins: with the claim that finance must be built for real people rather than idealized agents. From there the lecture deliberately refused a purely cynical interpretation. Human weakness creates opportunities for exploitation, but reputation, morality, regulation, and character all limit how far exploitation can organize the system.

Only then did the lecture turn to prospect theory. The value function gives the chapter its first mathematical spine: gains and losses are measured relative to a reference point, the gain side is concave, the loss side is convex, and the origin is kinked. The weighting function gives the second: objective probabilities are psychologically rounded, simplified, and sometimes exaggerated. Around that core the lecture assembles regret, gambling, overconfidence, cognitive dissonance, anchoring, representativeness, and social contagion.

The concluding claim is not that markets are clean, and not that they are rational in any simple sense. It is that financial institutions evolve, people learn, regulation matters, and integrity remains economically important. Behavioral finance therefore enlarges finance theory. It does not replace the institutional world with a list of biases; it explains why that world must be designed for the complicated creatures who live inside it.

Chapter 12

Misbehavior, Crises, Regulation and Self Regulation

This lecture picks up exactly where the discussion of behavioral finance leaves us. If markets are shaped by human foibles, then we have to ask what institutions do about them. Regulation is the answer, but not in only one sense. It is directed, first, at manipulation, abuse, and the exploitation of human weakness; it is directed, second, at the more technical problem of keeping the financial system working at all. The lecture unfolds from that starting point: first *too big to fail*, then the split between microprudential and macroprudential regulation, then the need for referees, and only after that the climb from regulation within the firm to regulation at the international level.

12.1 Regulation After Behavioral Finance

We begin, then, with the lecturer's explicit transition from the previous lecture. Regulation is not an afterthought to behavioral finance. It is what we do once we admit that finance is inhabited by imperfect people, by temptation, by persuasion, by manipulation, and by the possibility of plain bad character. At the same time, regulation has a more technical side: it deals with monopolies, externalities, and the need to prevent the whole financial system from malfunctioning.

The lecture chooses one system-level example almost immediately, because it gives the discussion urgency before any taxonomy appears. That example is *too big to fail*. When a large financial firm collapses, the government often rescues it, not because anyone intended to create such a privilege, but because the failure of a sufficiently large firm can drag the whole system down with it. In a compact reconstruction,

$$P(\text{bailout} \mid \text{large-firm failure}) > P(\text{bailout} \mid \text{small-firm failure}). \quad (12.1)$$

That inequality is not a formal theorem from the board; it is a way of writing the lecture's causal claim. Once it is true in practice, large firms acquire an implicit government guarantee. The mechanism the lecture stresses is then

$$\text{large-firm failure} \Rightarrow \text{threat to the system} \Rightarrow \text{likely bailout}, \quad (12.2)$$

$$\text{likely bailout} \Rightarrow \text{implicit guarantee} \Rightarrow \text{greater risk-taking} \Rightarrow \text{greater systemic risk}. \quad (12.3)$$

That is the key analytical move. A rescue may stabilize one crisis in the short run, but the expectation of rescue changes behavior in advance. The guarantee is effectively free insurance provided by taxpayers, and it pushes risk upward. In that sense regulation is not only about stopping swindlers; it is also about stopping the system from training its largest institutions to gamble more aggressively.

12.2 Too Big to Fail, Microprudential, and Macroprudential Regulation

From there the lecture introduces a pair of terms that became especially prominent after the financial crisis of the 2000s: *microprudential* and *macroprudential*. “Prudential” here means regulation concerned with prudent standards of business conduct. The difference lies in the object of protection:

Microprudential regulation $\longrightarrow$ protect firms, households, and individual investors, (12.4)

Macroprudential regulation $\longrightarrow$ protect the system as a whole. (12.5)

The lecture is careful about this distinction. Microprudential regulation is the kind that protects small people, small firms, ordinary borrowers, and ordinary investors from being fooled or exploited. Macroprudential regulation is the kind that worries about crises that spread through the system. On this classification, *too big to fail* is the paradigm macroprudential problem.

The lecturer also gives us a historical claim. Regulation has long involved both kinds, but until recently the mix leaned more heavily toward the microprudential side. Only after world financial crisis did the macroprudential side move to the center of the discussion. That historical change matters for the whole rest of the chapter, because many of the agencies at the end of the lecture are explicitly designed to think at the system level.

12.2.1 Question & Answer

Question. Why can rescuing a large firm create more risk rather than less?

Answer. Because the rescue does not only affect the moment of collapse; it affects incentives before collapse. If the market learns that sufficiently large firms will be rescued to prevent system-wide failure, then those firms behave as though they possess a contingent public guarantee. They can therefore take risks that smaller firms cannot take. The immediate rescue may avert panic, but the expectation of rescue breeds the next round of excessive risk-taking.

12.3 Why Markets Ask for Referees

Before the lecturer turns to institutions, he pauses to justify regulation itself. He does not assume that the case for rules is obvious. Instead he says that regulators are like referees at a sports event. They enforce rules, decide when the rules have been broken, and punish violations. Players may complain about referees, but they still want them, because without them dangerous play would become nearly necessary in order to win.

That same idea is carried over to finance. In a competitive system, if everyone else is doing something shady, then one may feel compelled to do it too in order to survive. The lecture calls this a race to the bottom. In schematic form,

$$\text{if others cheat} \Rightarrow \text{pressure to imitate} \Rightarrow \text{demand for enforceable rules.} \quad (12.6)$$

This is one of the lecture's most important normative points. Regulation need not be viewed as something imposed from outside against the wishes of all market participants. Competitors themselves may ask for it, because a well-policed game is preferable to a game in which success requires fouling.

The lecturer pushes the analogy further. Children on an empty field also create rules, and in effect referee one another. That is already a kind of self-regulation. He also insists that we should respect regulators as professionals. In sports we admire players more than referees; in finance we admire dealmakers more than regulators. But refereeing is a real craft, and regulation is a serious career. Some people move from Wall Street into regulation not because they failed, but because they would rather enforce decent rules than spend their whole lives in a morally compromised environment.

12.3.1 Question & Answer

Question. If competition is supposed to discipline behavior, why do competitors themselves ask for regulation?

Answer. Because competition disciplines some margins and corrupts others. It may pressure firms to perform, but it can also pressure them to copy profitable misconduct. When rule-breaking becomes a competitive advantage, honest behavior is no longer individually sustainable. A referee solves that coordination problem by making restraint universal rather than optional.

12.4 Five Levels of Regulation

Only after those motivational steps does the lecture reset itself with a blackboard outline. On the left are the prudential terms; on the right is a hierarchy of regulatory scale. The board is not a formal diagram. It is an outline of the lecture's architecture.

The lecturer partially blocks the two prudential words on the left board, but the transcript makes the intended pair clear. A clean typographic reconstruction is

$$\text{Too big to fail} \quad \text{Microprudential} \quad \text{Macroprudential.} \quad (12.7)$$

Beside that vocabulary the lecturer writes the five levels through which he will now climb:

1. Within firm
2. Trade groups
3. Local government
4. National
5. International




Figure 12.1: Blackboard outline of prudential regulation and the five levels of regulatory scope.

It is important not to overread the board. There is no drawn arrow linking the prudential terms on the left to the five-part ladder on the right. The relation is conceptual rather than graphical. The prudential distinction tells us what regulation may be trying to protect; the five-level ladder tells us where regulation is located and how broad its scope has become.

The lecturer immediately adds a historical observation: the older forms of regulation are the narrower ones, and over time regulation has become broader in geographical scope. That is why international regulation is saved for the end. The lecture is not only a taxonomy; it is also a historical ascent.

12.5 Within-Firm Regulation: Boards, Character, and Tunneling

The first and lowest level is within the firm. The lecturer insists on counting this as regulation even though the government is not yet involved. Firms set up their own rules, and in the modern corporation the board of directors can be understood as a kind of internal regulatory body. He even says that boards could be called boards of regulators, though no one actually speaks that way.

The key element is the outside director. Inside directors work for the firm; outside directors do not, except through their board service. They are therefore the ones who make the board look most like a regulator. The lecture then ties this to the previous discussion of character. Some people have more character than others, and society tries, in its imperfect way, to keep people of weak character out of positions where they can do the most damage. A board of directors is part of that filtering mechanism. We place on it people with reputations for being serious, high-minded, and publicly accountable.

That is why the lecture pauses over the Yale Corporation. The point is not celebrity for its own sake. It is reputational merger. Richard Levin, Dan Malloy, Fareed Zakaria, Indra Nooyi, Mimi Gardner Gates: names like these signal that one is joining one's reputation to the reputation of the institution. Indra Nooyi also introduces the idea of interlocking directorship, because a director is

judged partly by the other boards on which she is willing to serve. The lecture's point is that one would hesitate to propose a shady deal in front of such people.

The deeper reason for this internal regulation is that the corporation is extremely vulnerable to insider abuse. Once we move from a nonprofit example to a for-profit firm, the lecturer says, there are simply too many opportunities for sleazy behavior. That is why he introduces what he treats as perhaps the central internal pathology: *tunneling*. Tunneling means taking value that belongs to the corporation and sneaking it into private pockets instead of leaving it with the shareholders who supplied the capital.

The internal logic can be summarized as

outside capital $\Rightarrow$ delegated control $\Rightarrow$ opportunities for tunneling $\Rightarrow$ need for internal regulation.
(12.8)

The lecture then accumulates examples, and the accumulation matters:

- **Asset sales.** The firm sells an asset below market value to a relative or associate, with compensation to the insider coming later and privately.
- **Contracts.** The firm overpays a favored supplier, again with some return favor arranged off the books.
- **Executive compensation.** The chief executive pays friends or cronies far too much and defends the arrangement as a legitimate judgment about merit.
- **Expropriation of corporate opportunities.** A business opportunity that properly belongs to the firm is diverted into a new vehicle privately controlled by insiders.
- **Insider trading.** Management uses nonpublic information, acquired because of its position in the firm, to trade for private gain.

The lecture wants us to feel how fragile the corporation would be if these channels were left uncontrolled. Shareholders do not monitor every asset sale, every contract, every compensation package, every deal opportunity, every piece of private information. A corporation is therefore, as the lecturer puts it, a delicate thing. It works only if there are institutions inside it that restrain tunneling.

12.5.1 Question & Answer

Question. Why would anyone invest in a corporation if insiders have so many ways to tunnel value out of it?

Answer. Because the corporation is not just a productive unit; it is a governance arrangement. Boards of directors, especially outside directors, serve as internal regulators. They bring reputation, judgment, and community standards into the firm. The whole point is to make outside investors believe that managers cannot quietly convert corporate value into private value at every margin.

The lecture then sharpens the board's role by naming two legal standards:

Duty of care : act as a reasonable, prudent, rational person would act, (12.9)

Duty of loyalty : place the shareholders' interest ahead of private insider gain. (12.10)

The *duty of care* means that the board member must get information, watch management, and exercise judgment. A board member is not the manager of the company, but neither is he a ceremonial guest who attends four meetings a year and nods at the chief executive's presentation. He must know enough to act as a regulator. The *duty of loyalty* is usually interpreted, in the lecture's baseline formulation, as loyalty to shareholders. The lecturer notes that corporate social responsibility is widening this discussion, but he returns to the basic point: because the temptation to tunnel is so persistent, the board's first loyalty remains to those who put up the money.

That is why the lecturer treats within-firm regulation as perhaps the most essential form of regulation. Even before we move outward to exchanges or governments, the corporation itself must build mechanisms that make outside investment possible. Firms also create compliance departments, statements of purpose, and internal rulebooks. But the board remains the centerpiece, because it combines oversight, reputation, and fiduciary obligation in a single institution.

12.6 Trade Groups, Local Regulation, and National Disclosure

The lecturer then says explicitly that he wants to move up another step. We leave the single firm and turn to trade groups: organizations formed by firms themselves, sometimes as self-regulatory bodies. The historical example he chooses is the New York Stock Exchange, founded in 1792 after the first American stock market crash.

The Buttonwood Agreement matters here because the lecturer actually reads it and forces us to confront its blunt economic content. It is not a long ethical manifesto. It is a very short agreement that, in essence, establishes a minimum commission and gives preference to negotiation among the members themselves. The commission term can be summarized as

$$c_{\text{Buttonwood}} = \frac{1}{4}\% = 0.25\% = 0.0025. \quad (12.11)$$

Taken narrowly, that looks like price-fixing combined with exclusion of outsiders. In other words, it looks like a cartel. But the lecture does not leave the matter there. The more charitable interpretation is that the stockbrokers were reacting to a speculative boom and crash associated with a dubious promoter, and they were trying to exclude bad actors while maintaining incomes high enough to support what they regarded as responsible conduct. The lecturer is explicit that this second interpretation is partly reconstructed between the lines. The tension is the point: a trade group can stabilize conduct and create monopoly power at the same time.

12.6.1 Question & Answer

Question. How can self-regulation help markets without simply turning into cartelization?

Answer. Because the two functions are genuinely separable even when they arrive in the same institution. A trade group can screen participants, police conduct, and build reputational norms. But it can also fix prices and exclude rivals. Public regulation then has to do two things at once: preserve the useful parts of self-regulation, such as surveillance and ethical rule enforcement, while dismantling monopoly privileges.

That later dismantling comes with *May Day*, May 1, 1975, when fixed commissions on the New York Stock Exchange were made illegal. The lecturer's point is not that the exchange disappeared. It is

that the old cartel structure was broken. Members could no longer maintain mandatory commissions and closed dealing relations among themselves. The national market system widened the broker's obligation: not loyalty to a closed club, but duty to seek the best price for the customer on any exchange. The British *Big Bang* of 1986 plays the same role in the lecture's comparative story. In both cases, deregulation of commissions is itself a form of regulation against cartelization.

What remains even after that change is surveillance. The New York Stock Exchange still monitors trading records for manipulation. So the lecturer does not treat the post-1975 world as a world without referees. It is a world in which one particular kind of self-regulation—fixed commissions and monopoly dealing—has been rejected, while surveillance and conduct policing remain.

From trade groups the lecture moves to local regulation. In the United States, before the 1930s, regulation was overwhelmingly local. The federal government did very little. States developed *Blue Sky Laws* during the Progressive Era, beginning in Kansas in 1911, to block fraudulent or wildly overpromoted securities. The colorful name is tied, at least in one common explanation, to sales pitches promising that prices could rise with nothing above them but the blue sky.

Local regulation, however, ran into a structural problem:

$$\text{state-by-state regulation} \Rightarrow \text{cross-state evasion} \Rightarrow \text{need for national regulation.} \quad (12.12)$$

That is the mechanism the lecture wants us to see. Once telephones and boiler rooms make it easy to sell worthless securities across state lines, fragmented state enforcement becomes inadequate. A boiler room in one state can pressure victims in another state, and the geographical fragmentation of law becomes an invitation to fraud. So the Blue Sky system remains historically important, but it is no longer sufficient.

The next step is national regulation under the New Deal. In 1934 the United States created the Securities and Exchange Commission. The lecturer does not present it merely as another agency on a list. He makes it vivid through figures such as William O. Douglas and Arthur Levitt, each of whom, in his own way, confronted Wall Street's hostility to disclosure-based regulation. The guiding principle comes from Louis Brandeis: "sunshine is the best disinfectant." In finance that means the following:

$$\text{secrecy} \Rightarrow \text{scope for abuse,} \quad \text{disclosure} \Rightarrow \text{public scrutiny} \Rightarrow \text{cleaner markets.} \quad (12.13)$$

That is why the SEC is built around disclosure. Public companies must file documents; those documents are made available through EDGAR; and, in principle, anyone can inspect them. The lecture returns to this point because it wants disclosure to be understood not as bureaucratic clutter but as a regulatory technology.

The SEC also draws the line between public and private securities. A firm that wants to go public must pass through the regulated procedure of an IPO. A private fund is given more freedom, but only on the condition that access is restricted. The lecture describes two broad hedge-fund categories in approximate quantitative form:

$$N_{\text{class 1}} \leq 99, \quad y_{\text{single}} \geq 200,000, \quad y_{\text{married}} \geq 300,000, \quad W_{\text{investable}} \geq 1,000,000, \quad (12.14)$$

$$N_{\text{class 2}} \leq 500, \quad W_{\text{individual}} \geq 5,000,000, \quad W_{\text{institution}} \geq 25,000,000. \quad (12.15)$$

The statutory labels are noisy in the transcript, so the important point is conceptual: private funds are permitted greater latitude because participation is restricted to wealthy investors presumed able to protect themselves.

The lecture then uses the fee formula “2 and 20” to make that world feel concrete. If A denotes assets under management and Π denotes trading profits, a cautious reconstruction is

$$f(A, \Pi) = 0.02A + 0.20 \max\{\Pi, 0\}. \quad (12.16)$$

Worked example. If $A = 10^9$, then the management fee alone is

$$0.02 \times 10^9 = 2 \times 10^7 = 20,000,000. \quad (12.17)$$

That is the lecturer’s arithmetic check. Even before performance fees, the manager’s compensation can be enormous. Losses do not come symmetrically out of the manager’s pocket. The point is not to build a theory of hedge-fund contracting. The point is to dramatize the two-tier system of investor protection: one heavily regulated public tier for ordinary investors, and one much looser private tier for wealthy investors.

The national section ends, appropriately, with surveillance. The SEC does not only require disclosure; it also watches the market, often together with self-regulatory organizations. The lecture’s Lotus case is a compact illustration:

$$N_{\text{Lotus tippees}} \approx 25, \quad V_{\text{Lotus trades}} \approx \$500,000. \quad (12.18)$$

A private comment about IBM’s plan to acquire Lotus spreads through a network of acquaintances, trading surges, phone records are subpoenaed, and the chain is reconstructed. The Emulex case makes the complementary point: false information can also be used to manipulate markets, and surveillance can trace the trade, the short cover, the computer, and the person who sent the fake press release. In both stories the lecture returns us to its earlier analogy. Markets do not function well without referees.

Two final institutions complete the national level. FASB, recognized by the SEC, defines GAAP, providing accounting standards for disclosure. SIPC, created in 1970 after brokerage failure exposed the fragility of accounts held in street name, protects brokerage customers in a role parallel to that played by deposit insurance for bank depositors. Both are examples of how national regulation is partly delegated and partly specialized.

12.7 From Dodd–Frank to Europe and the International Layer

The final movement of the lecture shifts from the older national framework to the post-crisis expansion of macroprudential and international regulation. In the United States the key event is the Dodd–Frank Act of 2010. The lecturer reads it as a push away from a predominantly microprudential regime and toward a more macroprudential one. Earlier legislation had already addressed some fraud against individuals. Dodd–Frank adds institutions whose job is explicitly to worry about systemic risk.

The clearest example is the Financial Stability Oversight Council, FSOC. Its purpose is to think about *too big to fail* and about risks that threaten the whole system rather than merely one individual contracting party. Alongside it stands the new consumer-protection bureau created by Dodd–Frank, whose purpose is more microprudential: to focus specifically on ordinary borrowers, mortgage customers, and credit users who were not well served by the older regulatory structure. The pair reproduces, at the institutional level, the distinction made near the start of the lecture.

The same broadening then appears in Europe. The European Supervisory Framework, also associated with 2010, includes several new bodies:

- the European Systemic Risk Board in Frankfurt;
- the European Banking Authority in London;
- the European Securities Market Authority in Paris;
- the European Insurance and Occupational Pension Authority in Frankfurt.

The lecturer is careful here too. These institutions are new, and at the time of the lecture their eventual functioning is not yet fully known. Still, the direction is unmistakable: regulation is being lifted from the level of separate national jurisdictions toward a wider economic region.

From there the lecture goes fully international. The problem is easy to state. If regulation is costly, then financial firms can move offshore. If they dislike U.S. rules, there is always the Bahamas or the Cayman Islands. So we are led to another causal chain:

tighter national regulation $\Rightarrow$ offshore migration $\Rightarrow$ need for international coordination. (12.19)

This is why the international layer matters. The Bank for International Settlements, in Basel, does not legislate for the world, but it convenes and coordinates. The lecture says it brings together roughly 57 central banks. The Basel Committee operates similarly through recommended banking standards rather than direct command. Its major sequence is

Basel I (1988), Basel II (2004), Basel III (2010). (12.20)

The lecture's tone here is sober. International bodies do not usually possess the direct coercive force of national regulators. Their authority comes through recommendation, coordination, and repeated agreement. But that does not mean they are unimportant. After a global banking crisis, recommendation itself can have real force if the major countries agree to take it seriously.

A chronology helps keep the lecture's historical scale in view:

1792 : Buttonwood Agreement and the New York Stock Exchange, (12.21)

1911 : Kansas Blue Sky law, (12.22)

1934 : Securities and Exchange Commission, (12.23)

1970 : SIPC, (12.24)

1975 : May Day and the end of fixed commissions, (12.25)

1986 : Big Bang in the United Kingdom, (12.26)

2008 : G20 expansion, (12.27)

2010 : Dodd–Frank, European supervisory reforms, Basel III approval. (12.28)

The path from the G6 to the G7 and then to the G20 marks the same broadening in political form. Finance ministers first coordinate among a narrow group of large countries, then among a broader one. The G20, in the lecture, represents a more global forum, one that now includes countries such as China and India rather than overwhelmingly reflecting Europe and North America. The Financial Stability Board, also in Basel, is presented as the G20-linked body making recommendations for global regulatory procedures.

The lecture closes with guarded optimism. We are seeing financial regulation become larger in scale because finance itself has become larger in scale. A worldwide crisis has forced regulation to think not only about fraud against individuals, but about the stability of the world economy. That, for the lecturer, is an inspiring development, provided the international cooperation continues.

12.8 Summary

The lecture begins with human misbehavior and ends with world-scale institutions, and the movement between those two endpoints is carefully staged. We start with the thought that regulation must answer both to psychology and to system design. We then see why *too big to fail* creates moral hazard and why that problem is macroprudential. We pause to ask why markets want referees at all, and the answer is that competition without enforceable rules can become a race to the bottom.

Only then do we climb the five levels of regulation. Within the firm, boards of directors, reputation, and fiduciary duties hold back tunneling. At the level of trade groups, self-regulation can improve conduct but also harden into cartelization. At the local level, Blue Sky laws matter historically but fail once finance easily crosses state lines. At the national level, the SEC builds regulation around disclosure, surveillance, and investor protection. And at the international level, the lesson of crisis is that finance outgrows any one jurisdiction, so regulation must become more coordinated, more macroprudential, and more global.

That is the real shape of the lecture. It is not a miscellaneous list of agencies. It is a guided ascent from the moral fragility of the individual firm to the systemic fragility of the world financial order.

Chapter 13

Banks

These notes follow Robert J. Shiller’s Lecture 13 on banks in *Financial Markets*, curated by LazyingArt LLC. After a brief administrative aside, the lecture lays out its path very clearly: origins, theory, deposit insurance, regulation after the world financial crisis, shadow banking, and comparisons with earlier crises. We will keep that order. The subject throughout is the traditional deposit-taking bank. It is not the investment bank, whose characteristic activity is the underwriting of securities, and it is not the central bank, whose business is monetary management.

13.1 What a Bank Is: Spread, Notes, and Liquidity

We begin exactly where the lecture begins, with an attempt to define a bank by what it does. The first candidate is spread income. A bank borrows at one rate and lends at another:

$$\text{spread income} = r_L - r_D, \quad (13.1)$$

with

$$r_L > r_D. \quad (13.2)$$

Here r_L is the lending rate and r_D is the deposit or borrowing rate of the bank.

That is a good beginning, but the lecturer immediately pulls back from treating it as a finished definition. Historically, another aspect of banking was note issue. Two centuries ago many people would have answered the question “What is a bank?” by saying that a bank prints money. In the modern world that function has largely migrated to the central bank, though the lecture still notes limited private note issue in places such as the United Kingdom and Hong Kong. So note issue belongs to the story, but it no longer gives us the central modern mechanism.

The lecture then moves to the mechanism that matters most for theory: liquidity creation. Banks borrow short and lend long. That is a maturity discrepancy, not merely an interest-rate spread:

$$T_L > T_D, \quad (13.3)$$

where T_L denotes the maturity of loans and T_D the maturity of deposits or other short liabilities.

The lecturer’s example is the household mortgage. The homeowner wants funds for perhaps thirty years. The depositor does not want to lock up funds for thirty years. The depositor wants access on

short notice, even on demand. The bank stands between them, and by standing between them it creates liquidity. We can therefore summarize the basic mechanism as

$$\text{borrow short, lend long} \Rightarrow \text{liquidity for customers} \quad \text{and} \quad \text{run vulnerability.} \quad (13.4)$$

This is the first deep hinge of the lecture. The bank is useful for exactly the same reason that it is fragile. If everyone asks for cash at once, the bank cannot call in thirty-year loans tomorrow morning. The liquidity that looks effortless in normal times reveals a structural vulnerability in bad times.

13.1.1 Question & Answer

Question. How can banks create liquidity and yet be structurally vulnerable to crisis?

Answer. Because the two facts are the same fact viewed from opposite sides. The bank promises depositors short access to funds while investing in longer assets. Under ordinary conditions, not everyone withdraws at once, so the arrangement works beautifully. But if withdrawals become simultaneous, the bank cannot instantly turn long assets into cash at par. The lecture's point is that fragility is not something added onto banking after the fact. It is built into the same maturity transformation that makes banking socially useful.

13.2 Origins of Interest and Early Banking

Having arrived at fragility, the lecture steps backward. It asks where the idea of interest itself came from. The transcript is noisy in places, but the movement is clear. Economic historians are invoked to trace the word for interest back to ancient Sumer, where the same term also meant lamb. The proposed intuition is that lending may have grown out of older arrangements in which one rents out productive resources and claims part of the resulting produce. If land supports sheep, lambs appear; if capital is lent, some return should appear as well.

The lecture does not turn this into a formal theorem, and we should not force it into one. Its point is historical and institutional. Long before modern banks, people had already learned to think about productive resources, temporary transfer, and compensation for use. It also emphasizes that one does not need coined money in order to do banking. In the ancient world one could lend barley or wheat and take interest in barley or wheat. Banking begins when someone intermediates between those who have a stock of something and those who need temporary use of it.

The lecture then glances toward China, noting Song-dynasty records of interest rates and paper money, and then turns to Renaissance Italy, where recognizably modern banking institutions appear. The transcript garbles the Siena bank's name and even wobbles slightly on the founding year, so it is better to keep the claim cautious. The lecture points to the old Siena bank as the oldest surviving bank, philanthropic in origin, founded in the fifteenth century, and later associated with an early government guarantee of deposits. The important point is not antiquarian precision. The point is that banking and deposit insurance are not recent improvisations. They are old institutional responses to recurrent problems.

The English goldsmith bankers then provide the lecture's cleanest origin story. Goldsmiths had vaults. People left gold with them for safekeeping. The goldsmith issued a note promising payment

of the gold in storage. Those notes could then be endorsed and transferred. Eventually they circulated to bearer. At that stage the note itself became a medium of exchange. Then a second discovery followed: once the notes circulated, few people came back to withdraw the underlying gold. The goldsmith could lend against what sat in the vault. So safekeeping, note issue, transferability, and lending all grew together. Banking, in this telling, is not a single invention. It is a cluster of practices that coalesced around storage, transfer, and intermediation.

13.3 Modern Bank Types and Scale

Only after the historical detour does the lecture return to the modern, regulated landscape. That return matters. Because banking generated repeated trouble, governments gradually defined legal bank types and regulated them. Once we come back to the present, we do not encounter a generic bank. We encounter a structured and heterogeneous system.

The lecture begins with the commercial bank, the most important type in the United States. Commercial banks take deposits, pay interest, and make loans, historically most characteristically business loans. In 2010, the lecture reports

14.6 trillion dollars total U.S.-located commercial-bank assets, 10.1 trillion U.S.-chartered.
(13.5)

The gap matters because a substantial share of commercial banking activity in the United States is conducted by foreign institutions operating there.

The lecture then turns to smaller categories. Savings banks held

1.2 trillion dollars, (13.6)

and credit unions

0.9 trillion dollars. (13.7)

Both are tied to social movements rather than merely to profit maximization in the narrow sense. Savings banks are linked to nineteenth-century philanthropic efforts to promote thrift among households whom commercial banks had often ignored. Credit unions are club-like institutions organized around a workplace or a group affiliation. Both, in the lecture, are especially associated with mortgage lending.

Institution type	Characteristic function in the lecture	Magnitude in the lecture
Commercial bank	Deposits and business lending	14.6 trillion U.S.-located
Savings bank	Savings and mortgage lending	1.2 trillion
Credit union	Member-based deposit and loan services	0.9 trillion
Building society (U.K.)	Mortgage-oriented savings institution	U.K. analogue, no figure given

Table 13.1: Major bank types as organized in the lecture.

The lecture also makes a balance-sheet point that is easy to lose if we move too quickly. These are asset totals, not market capitalizations. A bank's assets are offset by liabilities, above all deposits. So even before formal theory begins, the lecture is already asking us to think institutionally and balance-sheet-wise.

13.4 Diamond–Dybvig, Liquidity, and Multiple Equilibria

Only now does the lecturer announce the theory of banks explicitly. The formal center is Diamond–Dybvig, published in the *Journal of Political Economy* in 1988. The lecture does not derive the model in its complete technical form. It does something more limited and more faithful to the course: it extracts the model’s conceptual structure and uses it to organize the rest of the discussion.




Figure 13.1: Blackboard outline of Diamond–Dybvig, multiple equilibria, and bank runs.

The blackboard stack is worth preserving almost verbatim:

$$\begin{array}{l}
 \text{Diamond–Dybvig} \\
 \text{JPE 1988} \\
 \text{Multiple equilibria} \\
 \text{Bank run}
 \end{array}
 \tag{13.8}$$

The visual order on the board matters. We start with the model and its publication marker, move to the phrase *multiple equilibria*, and end on the concrete phenomenon *bank run*. The lecturer’s gesture toward the last line makes the progression unmistakable.

What, then, is the model doing in the lecture? It treats liquidity as a genuine economic good. The lecturer’s analogy is portfolio diversification: diversification is not a physical resource dug out of the ground, but a benefit produced by arranging a portfolio well. Likewise, liquidity appears when a bank is organized in the right way. Depositors gain access, borrowers gain long financing, and no additional physical resource has been created. That is why the lecture calls liquidity an economic good that one somehow gets “for nothing.”

But the same institution admits more than one equilibrium. In a cleaned reconstruction of the lecture’s logic,

$$\text{confidence} \Rightarrow \text{good equilibrium}, \tag{13.9}$$

$$\text{panic expectations} \Rightarrow \text{bad equilibrium} \Rightarrow \text{bank run}. \tag{13.10}$$

Or, more compactly,

$$\text{multiple equilibria} \Rightarrow \text{scope for a bank run}. \tag{13.11}$$

If people believe the banking system is sound, it functions well. If they suddenly expect others to withdraw, the system can collapse precisely because the assets are long and the liabilities are short. The lecture is careful, though, not to overclaim. It immediately notes that the recent global crisis also involved a real-estate bubble and declining asset prices, elements not built into the clean Diamond–Dybvig setup. The model does not explain everything banks do and everything that can go wrong. But it does capture the essential link between liquidity provision and run-prone equilibrium structure.

13.4.1 Question & Answer

Question. Why can deposit insurance keep the system in the good equilibrium?

Answer. Because the bad equilibrium relies on the incentive to rush for the exit before others do. If depositors believe that even a failing bank will repay them, the incentive to run first is weakened. So deposit insurance is not merely ex post compensation. It is an ex ante device for altering expectations:

$$\text{insured deposits} \Rightarrow \text{less incentive to withdraw early} \Rightarrow \text{run prevention.} \quad (13.12)$$

That is why the lecture presents Diamond–Dybvig not only as a theory of bank runs, but also as an economic rationale for deposit insurance.

13.5 Banks as Information and Monitoring Institutions

The lecture then broadens the theory. Diamond–Dybvig explains why banks create liquidity and why runs can occur. But banks do more than maturity transformation. They also solve information problems that would otherwise cripple direct finance.

Suppose a firm needs money for a factory. It could go to an investment bank, issue bonds or commercial paper, and try to borrow directly from the public. Why, then, do we still need banks? The lecture’s answer begins with adverse selection. The public cannot easily judge the quality of a company. Informed investors may take the good paper first. Uninformed investors then start to suspect that what reaches them is precisely the bad paper. Once that suspicion takes hold, direct lending to the public becomes difficult.

The lecture is concrete here rather than abstract. Banks solve this problem by locating themselves inside communities. They employ loan officers who know local businesses, hear local gossip, track local reputations, and gather information that cannot be written down neatly in a securities prospectus. The point is not that banks eliminate uncertainty. The point is that they specialize in screening and can therefore intermediate where atomized public investors cannot.

13.5.1 Question & Answer

Question. Why borrow from a bank instead of issuing securities directly to the public?

Answer. Because the bank is not merely passing funds through. It is producing information. A bank can spend resources to know the borrower, while scattered public investors cannot do so

cheaply. The bank therefore stands between savers and borrowers as a delegated expert. In the lecture's logic, banks are useful not only because they create liquidity, but because they reduce the adverse-selection problem that would otherwise poison direct debt issuance.

The lecture then adds moral hazard. A borrower who has received funds may take actions that are far riskier than the lender intended. The racetrack example is deliberately theatrical: a distressed firm borrows money, takes a wild gamble, keeps the upside if it wins, and shifts much of the downside onto others if it loses. The lecture does not formalize expected utility, and we should not pretend that it does. But it gives a clear limited-liability intuition: once debt is in place, borrowers may want to swing for the fences.

Banks respond not only by screening up front, but by monitoring after the fact. Commercial loans are often formally short-term and repeatedly renewed. That keeps the borrower under observation:

$$\text{short-term loan} \Rightarrow \text{renewal} \Rightarrow \text{monitoring leverage.} \quad (13.13)$$

If the bank sees behavior that looks like moral hazard, it can tighten terms or refuse renewal. So the lecture's second theoretical spine is this: banks are not only liquidity providers, but also information processors and monitors.

13.6 Deposit Insurance in Practice: FDIC, FSLIC, Northern Rock, and IKB

Having moved from liquidity to information, the lecture returns to institutions and asks what governments actually do about runs. Deposit insurance first appeared in the lecture as a theoretical answer inside Diamond–Dybvig. It now reappears as a historically uneven institutional device.

The basic idea is simple. People hear rumors, panic, and try to withdraw deposits. Governments therefore guarantee deposits in order to prevent runs. But the history is “checkered,” as the lecture puts it, because guarantees can be limited, insurers can fail, and what really stops the run is often not the formal scheme alone but the state's willingness to stand behind it.

The lecture first turns to the United States and the FDIC, created in 1933. Its striking claim is not merely that the FDIC has not failed. It is that the FDIC changed public psychology. Since 1933 the United States did not experience the old-style depositor panic on insured commercial banks that had characterized earlier crises. The lecture explicitly connects this to multiple equilibria: if people believe the system is sound, it is much more likely to remain sound.

The contrast case is the FSLIC, which insured savings and loan institutions and did fail. In the S&L crisis of the 1980s, the insurer ran through its reserves and became insolvent. The U.S. government then stepped in and paid the losses. The lecture gives the bill as about \$150 billion. The institutional moral is sharp: the visible insurer may not be the ultimate guarantor. The ultimate guarantor is the state, which cannot afford a collapse in confidence.

To keep the board evidence close to the institutional narrative, we can pair the screenshot with the minimal cleaned reconstruction it supports:

$$\text{S\&L crisis,} \quad \text{1980s,} \quad \$150 \text{ billion,} \quad \text{Northern Rock.} \quad (13.14)$$

Northern Rock provides the lecture's modern run example. In 2007, a rumor spread that the bank held large amounts of subprime exposure. Depositors lined up outside the branches. Newspaper



Figure 13.2: Crisis-comparison board notes with S&L and Northern Rock.

photographs amplified the panic. Yet the United Kingdom already had deposit insurance. The lecture states the scheme cautiously as follows:

$$\text{full insurance up to } \pounds 3,000, \quad 90\% \text{ insurance up to } \pounds 75,000. \quad (13.15)$$

That did not stop the run. According to the lecture, the relevant fact is simple: many depositors had more than the fully insured amount, and partial insurance was not enough to eliminate the incentive to run. Confidence returned only when the authorities effectively guaranteed everyone.

The same pattern appears in Germany with IKB. There, the government intervened before a full-scale run:

$$\text{IKB bailout} = 1.5 \text{ billion euros}. \quad (13.16)$$

Again the logic is not subtle. Governments move because once confidence is lost, the cost of restoring it later can be much higher.

13.6.1 Question & Answer

Question. Why can a bank run still happen even when deposit insurance exists?

Answer. Because a formal scheme may be too narrow to change incentives for the relevant depositors. If large balances are only partly protected, or if the public doubts the scope of the guarantee, depositors still want to be first in line. The lecture's lesson is that what matters is not the legal wording alone, but the credibility of the backstop. A limited guarantee may not stop a run. A credible blanket guarantee often does.

13.7 Basel Regulation and the Arithmetic of Capital Requirements

At this point the next transition is almost forced. If we insure banks, then we must regulate them. Otherwise the insurance itself creates bank moral hazard. The lecture says this very directly. A

bank can do its own version of the racetrack gamble: it can take on very risky assets, enjoy the upside if things go well, and shift losses onto the insurer if things go badly. Insurance without regulation is therefore unstable.

This leads into Basel. The lecture gives the chronology in sequence:

$$\text{Basel I (1988),} \quad \text{Basel II (2004),} \quad \text{Basel III (2010),} \quad (13.17)$$

with Basel III phased in gradually until 2019. Basel is not world government. The Basel committee cannot command states. It issues recommendations. But the recommendations were widely adopted because banking is international and countries do not want radically inconsistent regulatory environments.

The lecture then reduces the Basel framework to one essential device: capital requirements based on risk-weighted assets. The bank must hold enough capital for the risks it takes. That requires a regulatory measure of risk. In the lecture's simplified version, the key weights are

$$0\%, \quad 20\%, \quad 50\%, \quad 100\%. \quad (13.18)$$

The lecture's classification can be summarized as

$$w_{\text{gov}} = 0, \quad (13.19)$$

$$w_{\text{muni/Fannie/Freddie}} = 0.2, \quad (13.20)$$

$$w_{\text{mort}} = 0.5, \quad (13.21)$$

$$w_{\text{biz}} = 1. \quad (13.22)$$

The transcript is loose on the expansion of the OECD acronym, so it is better to keep the label simple and follow the lecture's substance rather than its exact wording: advanced-country sovereign bonds receive the zero weight in the simplified discussion.

Now comes the most extended arithmetic of the lecture. Suppose a bank has four blocks of assets, each equal to 100 million dollars: government bonds, Fannie/Freddie exposure, mortgages, and commercial loans. Then

$$A = 100 + 100 + 100 + 100 = 400 \text{ million.} \quad (13.23)$$

But total assets are not the same thing as risk-weighted assets. Once the Basel weights are applied, we obtain

$$\text{RWA} = 100(0) + 100(0.2) + 100(0.5) + 100(1) = 170 \text{ million.} \quad (13.24)$$

This is the first key worked step. The bank owns 400 million in assets, but from the regulatory point of view it carries only 170 million of risk-weighted assets.

The lecture then jumps to Basel III common-equity requirements. In its simplified form,

$$\text{CE}_{\min} = 0.045 \text{ RWA}, \quad (13.25)$$

$$\text{effective CE target} = (0.045 + 0.025) \text{ RWA} = 0.07 \text{ RWA}, \quad (13.26)$$

$$\text{countercyclical target} = (0.045 + 0.025 + 0.025) \text{ RWA} = 0.095 \text{ RWA}. \quad (13.27)$$

The lecture explains the 2.5% capital conservation buffer carefully. It is not quite the same as the hard minimum, but banks that fall short of it face payout restrictions. So in practice the lecturer treats 7% as the operative target.

Applying that target to the worked bank gives

$$0.07 \times 170 \text{ million} = 11.9 \text{ million.} \quad (13.28)$$

So the bank's common equity should be about 11.9 million dollars in the practical Basel III sense of the lecture.

To connect the regulation back to balance sheets, the lecture reminds us that shareholders' equity is what remains after liabilities are subtracted from assets:

$$E = A - L. \quad (13.29)$$

Common equity is then obtained from total shareholders' equity by subtracting preferred equity:

$$CE = E - \text{preferred equity.} \quad (13.30)$$

The most illuminating step comes next. Suppose the bank actually has 12.9 million of common equity. It is therefore 1 million above the practical requirement. What can it do with that extra million? The answer depends on the asset weight. If it buys additional Fannie/Freddie paper at a 20% weight, then

$$\frac{1,000,000}{0.2 \times 0.07} \approx 71.4 \text{ million,} \quad (13.31)$$

which the lecture rounds to about 70 million. But if the bank uses the same slack for 100%-weighted business loans, it gets only

$$\frac{1,000,000}{0.07} \approx 14.3 \text{ million,} \quad (13.32)$$

which the lecture rounds to about 14 million.

This is where the lecture becomes most interesting. The arithmetic does not merely measure safety. It changes incentives. If one dollar of excess common equity supports far more low-weighted mortgage-related exposure than business lending, banks are nudged toward the former and away from the latter. The lecture then makes that criticism explicit by linking low weights on Fannie and Freddie to the distorted buildup before the financial crisis.

13.7.1 Question & Answer

Question. Why can risk-weighted capital rules push banks toward assets that are socially or economically inferior?

Answer. Because scarce equity is allocated where regulation lets it go furthest. If the same 1 million of surplus capital supports about 70 million of low-weighted mortgage-related exposure but only about 14 million of business loans, the bank has a clear balance-sheet incentive to prefer the former. The lecture's criticism of Basel is therefore not that capital requirements are useless. It is that the chosen weights can favor assets that look safer to the regulator while starving other forms of lending that may be more productive for the economy. Basel can reply that its job is bank soundness, not industrial policy. But that reply explains the distortion; it does not remove it.

13.8 Global Banking Crises, Contagion, and Shadow Banking

From Basel the lecture widens again. The recent global crisis is vivid, but it is not unique. Banking crises recur, and they spread.

Mexico is the first example. The lecture describes the privatization of banks under Salinas, a lending boom, and the absence of effective regulation. Its key quantitative marker is very sharp:

$$\frac{\text{bank lending}}{\text{GDP}} \approx 10\% \text{ in 1988, } \quad 40\% \text{ in 1994.} \quad (13.33)$$

That explosive increase is not presented as a curiosity. It is evidence of what happens when liberalization outruns supervision. The lecture's interpretation is that moral hazard grew in the background because lenders and borrowers came to believe that the government would somehow make them whole if things went badly. It turned out that it could not.

The Asian crisis is then described as something like an international bank run. International banks had lent heavily into several Asian countries. Then international investors wanted their money back. Funds were withdrawn, dependence on foreign lending was exposed, and the crisis spread from Asia into Russia and then Brazil. The lecture is using these examples to make a single point: the world financial system is interconnected, and that interconnection is one of the reasons countries now try to coordinate bank regulation through groups such as the G20.

Argentina is treated more briefly. The lecture mentions it as another case in which banking disruption became severe enough that the government shut down parts of the system. The details are not developed here, and so the note should not pretend that they were. The example functions as one more reminder that banking crises are recurrent, severe, and politically destabilizing.

The final horizon is shadow banking. Here the lecture becomes deliberately gestural. Shadow banking means institutions that are not commercial banks but are nevertheless doing bank-like things. The lecture names major investment banks associated with the recent crisis and emphasizes that they were outside the Basel III framework as commercial banks. Innovation keeps producing new structures, and regulators struggle to keep up. The lecturer does not conclude that innovation should be stopped. He concludes that regulation will necessarily become more complex.

13.8.1 Summary

We can now see the lecture unfold as a single argument. We begin by trying to define a bank through spread income, but that is not enough. We then arrive at liquidity creation, and with it at fragility. History shows that these functions are old. Modern institutional types show that the system is heterogeneous and already regulated before theory begins. Diamond–Dybvig explains why liquidity is valuable and why multiple equilibria make runs possible. Information and monitoring explain why banks still matter even when securities markets exist. Deposit insurance stabilizes expectations, but because it creates moral hazard it immediately leads us to regulation. Basel supplies the arithmetic of that regulation, yet its risk weights can distort bank behavior. The international crises at the end of the lecture are therefore not an appendix. They are evidence that banking failures are recurrent, contagious, and central to macroeconomic life. That is why the lecture closes with a defense of regulators: they are trying to manage a system that is both indispensable and structurally delicate.

Chapter 14

Guest Speaker Maurice “Hank” Greenberg

Robert J. Shiller introduces Maurice “Hank” Greenberg in exactly the right way: not as a celebrity interlude, but as evidence that finance is larger than formal pricing formulas. It is also about underwriting, incentives, organizational design, political access, and the governance of risk. These notes, curated for the Lazyearn track by LazyingArt LLC, follow that cue. We do not have a blackboard derivation lecture here. We have something more institutional and, in its own way, just as analytical: a narrated account of how a financial organization is built, how it scales, and how a change in contractual structure can turn a large balance sheet into a liquidity trap.

14.1 Opening Frame: Finance Beyond Formulas

Shiller’s opening tells us how to read the lecture. We are interested in finance, but not only in finance as a collection of models. We are also interested in how a firm is financed, how people are chosen and disciplined, how risk is judged, and how character enters decision-making. Greenberg’s story therefore begins with military service, education, and contingency, but it does not remain mere biography. It is setting up a style of reasoning.

Greenberg enlists before finishing high school, serves in Europe in World War II, returns to school, goes on to law school, serves again in Korea, and comes back with a degree, a family, and the need for a job. That background matters in the lecture because it explains the cast of mind from which he later speaks. He is not introducing himself as a theorist who discovered insurance from inside a model. He is introducing himself as someone trained by responsibility and pressure, and then forced to learn a business from the ground up.

The first real hinge comes only when he stumbles into insurance. Up to that point the lecture is telling us who is speaking. Then it tells us what the business is.

14.2 From Soldier to Underwriter: Learning to Price Risk

Greenberg returns from Korea, decides almost immediately that he does not want to practice law, walks into an insurance company, and begins as a junior underwriter. At that moment the lecture

stops moving by anecdote and starts moving by mechanism. Greenberg asks his own question: what is a junior underwriter?

Definition 14.1. A junior underwriter analyzes risks and decides whether the company wants to write those risks or not.

If we want a cautious mathematical restatement of that task, we can write it this way. Let i denote a proposed risk, let π_i denote the underwriting profit from accepting it, and let $\mathcal{I}$ denote the information available at the decision date. Then the underwriting decision is an accept-reject rule:

$$\text{write risk } i \quad \text{only if} \quad \mathbb{E}[\pi_i \mid \mathcal{I}] > 0. \quad (14.1)$$

This equation is not spoken symbolically in the lecture, but it is a faithful reconstruction of Greenberg’s verbal definition.

The special-risk division in which he works makes the point concrete. These are not standardized risks arriving in neat homogeneous piles. They are sports risks, accidental-death coverage for executives, and other unusual exposures. The job is therefore not clerical. One has to inspect the structure of the risk, decide whether the premium is adequate, and decide whether the firm even wants such an exposure on its books.

14.2.1 Question & Answer

What exactly does an underwriter do?

Greenberg’s answer is short and very precise: an underwriter studies a risk and decides whether the company should write it. In note form we may separate that into two acts. First comes analysis: what is the risk, how unusual is it, and what could it cost? Then comes selection: even if the risk can be priced, does it belong in this firm’s portfolio? That is why underwriting is the lecture’s first analytical core. From this point on, the rest of the lecture becomes a scaled-up version of the same problem.

14.3 Turning Around American Home and Building AIG

Greenberg next meets C. V. Starr and enters a larger institutional setting. Starr’s organization has an international side, especially in Asia, and one troubled domestic insurer, American Home. Greenberg is sent to American Home, and he immediately asks the next natural question: what was wrong with it?

The diagnosis is operational rather than mystical. American Home wrote business through agents, but its agency organization was poor, and therefore the quality of its business was poor. In compact form, the logic is

$$\text{weak agency organization} \Rightarrow \text{poor business quality} \Rightarrow \text{persistent losses}. \quad (14.2)$$

The problem is not merely that the company is losing money. The deeper problem is that its distribution structure is selecting the wrong risks.

Greenberg’s remedy is correspondingly structural. He gets rid of the agents, shifts toward the brokerage business, and starts writing large commercial risks rather than the older, weaker book:

$$\text{brokerage model} \Rightarrow \text{larger commercial risks} \Rightarrow \text{need for reinsurance}. \quad (14.3)$$

That last step is decisive. Once the firm begins to write large property, casualty, marine, and aviation risks, it needs reinsurance on scale and on credible terms. So Greenberg goes to London and gets backing from Lloyd’s. The lecture treats this not as ornament, but as the next link in a sequence: underwriting skill leads to larger risks, larger risks lead to reinsurance dependence, and successful expansion quickly runs into a capital constraint.

That capital constraint produces the next move. American Home grows rapidly, and Greenberg looks for other insurers with the same sort of fixable weakness. National Union is acquired and repaired in much the same way. New Hampshire follows. Then, in 1967, a holding company is placed over these firms and AIG is born. The order matters. AIG is not introduced as a brand in search of a theory. It is the organizational consequence of a problem-solving sequence:

repair business model $\Rightarrow$ improve underwriting $\Rightarrow$ obtain reinsurance $\Rightarrow$ solve capital bottleneck $\Rightarrow$ create holding
(14.4)

That is the lecture’s first full institutional arc. We begin with the same judgment problem that defined the junior underwriter. We end with a firm architecture built around it.

14.4 Innovation, Diversification, and Opening Markets

Once AIG exists, Greenberg turns to the next question: how does such a firm grow without becoming more fragile? His answer has several parts, but in the lecture they belong together. Product innovation, diversification, expense discipline, and the opening of closed foreign markets are not separate stories. They are parts of one operating philosophy.

The product side is concrete and important. AIG introduces directors-and-officers liability insurance, political-risk insurance, and kidnap-and-ransom insurance. These examples matter because Greenberg is showing us how insurance expands: one identifies a need created by business activity, then creates a contract that makes that need insurable. In this sense, innovation is not decoration at the margin. It is part of the firm’s core growth engine.

The next piece is diversification. Property-casualty insurance is volatile. Earthquakes, hurricanes, and changing economic environments do not strike evenly across time, product, or geography. Greenberg says directly that greater diversification brings greater stability. In careful mathematical form, if

$$X = \sum_{j=1}^n X_j, \quad (14.5)$$

where each X_j is the profit or loss from a line of business or country exposure, then

$$\text{Var}(X) = \sum_{j=1}^n \text{Var}(X_j) + 2 \sum_{j < k} \text{Cov}(X_j, X_k). \quad (14.6)$$

If the covariance terms are not all perfectly positive, then aggregation stabilizes outcomes:

$$\text{Var}\left(\sum_{j=1}^n X_j\right) < \sum_{j=1}^n \text{Var}(X_j). \quad (14.7)$$

This variance formula is a standard reconstruction, not a lecture-board formula, but it says exactly what Greenberg means when he tells us that more diversification gives more stability.

14.4.1 Question & Answer

Why does international diversification make an insurance company more stable?

Because the insurer’s bad outcomes do not all come from the same place at the same time. If earthquakes, hurricanes, political disruptions, and business cycles are distributed across countries and lines of business, then losses do not pile up in perfect synchrony. The mathematics is just the covariance expansion written above. The lecture’s phrase is simpler and better: more diversification provides greater stability.

Greenberg then gives one of the lecture’s clearest numerical comparisons. Most insurance companies, he says, operate at an expense ratio of about 30%, whereas AIG ran at about 19%:

$$\text{industry expense ratio} \approx 30\%, \quad \text{AIG expense ratio} \approx 19\%. \quad (14.8)$$

He immediately answers the question he has raised himself: how did they do that? By being efficient and by using reinsurance properly. The lecture is careful here. The number matters, but the mechanism is organizational, not mysterious.

The growth story then moves outward into world politics. AIG ultimately operates in 130 countries. But markets do not open themselves. Japan is described as a closed market, and China becomes the long strategic example. Greenberg says it takes from 1975, his first visit to China, to 1992, to obtain the first life-insurance license granted to a foreign company:

$$1975 \longrightarrow 1992. \quad (14.9)$$

He also emphasizes the ownership asymmetry:

$$\text{usual foreign ownership ceiling} = 49\%, \quad \text{AIG ownership} = 100\%. \quad (14.10)$$

This part of the lecture is easy to flatten, but it should not be flattened. Greenberg is not merely saying that AIG expanded abroad. He is saying that diversification, product design, trade diplomacy, and foreign-policy engagement were all parts of one business model. Financial services were not initially part of major trade negotiations; he says AIG worked to bring them into the negotiating framework. He also insists that there was a line the firm would not cross: bribery. That matters analytically because it shows us that discipline, reputation, and market access are linked. In the lecture, international expansion is never just geography. It is structure.

14.5 Organization, Incentives, and the Internal Logic of AIG

At this point Greenberg turns inward and, in effect, asks the lecture’s next question: what kind of structure makes a company like this possible? He has already hinted at the answer when he says that senior management at AIG was like a band of brothers. Now he spells out the underlying machinery.

The first element is mobility. AIG’s overseas cadre includes what Greenberg calls the MOP system, “Mobile Overseas.” A manager may be in Nigeria now and in Singapore six months later. The point is not just willingness to travel. The point is that an international insurer cannot be run as a collection of immobile local outposts. It needs people who understand that the firm is one organization moving across many jurisdictions. Greenberg even says the MOP group was like AIG’s own State Department.

The second element is compensation. Here the lecture becomes especially precise. He says that at its peak AIG was the largest and most profitable insurance company in history, and he asks how such an organization can remain coherent. His answer begins with compensation rules:

$$g \approx 15\% \text{ per year,} \quad \text{salary} \leq 1 \text{ million dollars.} \quad (14.11)$$

The salary cap is important because it keeps fixed compensation from becoming the whole story. Performance matters more than salary. Growth targets matter. And the real retention device lies elsewhere.

Greenberg then gives the ownership arc. The private Starr entities initially own all of AIG, but rapid growth requires capital and therefore dilution:

$$\text{initial private ownership of AIG} = 100\%, \quad \text{later control} \approx 15\%. \quad (14.12)$$

At the same time, the enterprise becomes enormously larger:

$$\text{market capitalization: } 300 \text{ million dollars} \rightarrow 200 \text{ billion dollars.} \quad (14.13)$$

This is one of the lecture’s cleanest finance lessons. Control falls in percentage terms, but the object controlled becomes vastly larger. Dilution is not the negation of growth; it is one of growth’s financing conditions.

The most effective incentive device in the lecture is the deferred stock allocation. Greenberg says that if performance targets are hit over a two-year period, AIG shares are set aside for retirement:

$$\text{performance targets hit over a two-year window} \Rightarrow \text{AIG shares reserved for retirement.} \quad (14.14)$$

If an employee leaves, those shares are left behind. This is the lecture’s version of golden handcuffs. The mechanism is worth stating in slow order:

1. Growth and performance goals are set.
2. Shares are allocated to selected individuals if those goals are met.
3. The shares are deferred until retirement.
4. Early departure means forfeiture.

The result is retention without contractual lock-in.

14.5.1 Question & Answer

If top managers had no long-term contracts, what kept them aligned with the firm?

Greenberg answers this directly in the student questions. He refused a contract for himself because, as he puts it, if he were not doing the job, the company should not be forced to keep him. The anti-contract principle therefore has one side of discipline: poor performance should be removable. But it also has one side of attraction: strong performance is rewarded by bonuses, status, and deferred stock. In the lecture, alignment comes from culture and incentives, not from legal tenure. One stays because one wants to build, and because leaving means giving up a great deal.

That is why the lecture keeps returning to culture. Greenberg says later, in one of the final questions, that culture starts at the top, and that the CEO is the top risk manager in the firm. We should hear that already here. The compensation structure is not an isolated payroll policy. It is part of a theory of organizational control.

14.6 AIG Financial Products, Risk Monitoring, and the Crisis Mechanism

Now the lecture breaks. Greenberg has spent a long time describing a successful organization, and then he asks, almost abruptly: so what happened? That hard turn should remain hard on the page.

The first part of the answer, in Greenberg’s telling, is not yet about CDOs or CDS. It is about organizational weakening. He says he was forced out in 2005, in what he presents as the result of Elliot Spitzer’s politicized regulatory campaign. That claim belongs to Greenberg’s account, and it should be read that way. But within the narrative it has a clear function: it explains how a firm that had been built around close monitoring and internal discipline came to lose that discipline.

Greenberg says his succession plan had been gradual and internal. The successor team was supposed to come from within because the firm had a culture built over many years and a company of 92,000 employees could not be handed over lightly. He then gives a very specific operational detail. Senior management met every Monday morning to receive real-time reports across the organization, and there was a second weekly meeting devoted to hedging and risk management. Those meetings, he says, were later discontinued.

That is the bridge to AIG Financial Products. Before the crisis, AIG had already expanded further into life insurance, consumer finance, thrift operations, and related financial services, supported by a huge asset base:

$$\text{assets to invest} \approx 1 \text{ trillion dollars.} \quad (14.15)$$

AIG Financial Products, in Greenberg’s telling, began as a successful derivatives business. It was hedged. It was monitored. The firm even maintained mirrored computer systems off-site so that transactions could be tracked independently. He summarizes the control principle in a phrase worth preserving:

$$\text{enterprise risk management} = \text{monitoring of market risk and credit risk across the firm.} \quad (14.16)$$

Only after that setup does the lecture turn to the main mechanism. The transcript is garbled in part of this section, so we have to proceed carefully, but the core sequence is clear. AIG Financial Products became exposed to instruments that Greenberg describes as CDOs largely built out of real-estate-related assets. Originally, he says, a credit default swap responded only when the underlying instrument defaulted. Later that logic changed. The contract could now generate collateral demands simply because the instrument had lost market value.

The cleanest schematic comparison is:

$$\text{old rule: } \mathbf{1}_{\{\text{default}\}} \Rightarrow \text{payment,} \quad (14.17)$$

$$\text{new rule: } \Delta P_{\text{CDO}} < 0 \Rightarrow \text{collateral posting.} \quad (14.18)$$

That is the decisive contractual shift. Under the old regime, one waits for default. Under the new regime, one responds to a mark. In Greenberg’s language, the instrument did not have to default. It only had to lose value, and then the protection seller had to put up collateral.

This lets us state the crisis mechanism in the lecture’s own order:

1. The value of the CDO falls.
2. Counterparties demand additional collateral.

3. A downgrade from AAA activates collateral clauses that had previously remained dormant.
4. Cash needs rise by billions.
5. The firm can fail on liquidity even if its regulated insurance subsidiaries remain solvent.

In compact mathematical form,

$$P_{\text{CDO}} \downarrow \Rightarrow C \uparrow \quad (14.19)$$

$$\Rightarrow \text{cash demand } \uparrow \quad (14.20)$$

$$\Rightarrow \text{liquidity stress.} \quad (14.21)$$

A small numerical example makes the lecture’s point precise. Suppose a CDO is marked from par to 60:

$$P_0 = 100, \quad (14.22)$$

$$P_1 = 60, \quad (14.23)$$

$$\Delta P_{\text{CDO}} = -40. \quad (14.24)$$

Under the original default-based logic, this by itself does not trigger payment if no default has occurred. Under the later collateral logic, the drop in value can trigger a collateral call of roughly the same order of magnitude even though the loss has not yet been realized through default. This is not a literal contract term from the lecture; it is a compact reconstruction of the mechanism Greenberg describes.

The next problem is price discovery. Greenberg says there was no exchange on which CDOs traded, so there was no common market price:

$$P_{\text{CDO}}^{(1)} \neq P_{\text{CDO}}^{(2)} \neq \cdots \neq P_{\text{CDO}}^{(m)}. \quad (14.25)$$

Different broker-dealers quoted different values. Then the very collateral number that was supposed to measure risk became contestable. In Greenberg’s account, Goldman Sachs marked the relevant CDOs at the lowest values, which drove especially severe collateral demands.

The AAA issue then enters in a sharp and simple form:

$$\text{AAA rating} \Rightarrow \text{no collateral posting}, \quad \text{below AAA} \Rightarrow \text{collateral posting.} \quad (14.26)$$

Greenberg says AIG was triple-A while he was there, and that it lost that rating the day he left. Once the rating was gone, the collateral machinery came alive.

Definition 14.2. Liquidity is the ability to produce cash fast enough to meet immediate obligations such as collateral calls. Solvency is the ability of a firm, or of a regulated insurance subsidiary, to remain economically able to meet its obligations over time.

14.6.1 Question & Answer

Why could AIG fail on collateral and cash even if its insurance subsidiaries were still solvent?

Because the collateral mechanism moved on a faster clock than solvency. Greenberg stresses that the insurance companies were state-regulated, their capital was protected, and policyholders were

protected. The parent company’s problem was different. Once collateral calls rose into the billions, it needed cash immediately. A firm can therefore be solvent in its underlying insurance entities and still fail at the parent level because it cannot fund a liquidity shock.

The rescue terms are stated numerically and should be preserved as stated:

$$85 \text{ billion dollars at } 14.5\% \text{ interest,} \quad 79.9\% \text{ equity taken by the Fed.} \quad (14.27)$$

Greenberg later says that government ownership rose to about

$$\text{government ownership of } \text{AIG} \approx 92\%. \quad (14.28)$$

He also contrasts AIG with Citigroup:

$$\text{Citigroup support: } \approx 30\% \text{ equity,} \quad \text{AIG support: } \approx 92\% \text{ equity.} \quad (14.29)$$

Here the notes have to remain careful. Greenberg also says that many of the relevant CDO values were really around 40 to 60 cents on the dollar, but that counterparties were paid at 100 cents on the dollar. That is part of his account of the rescue and its unfairness, and it should remain attributed as such. The analytical point is that once price discovery fails, valuation and settlement become political as well as financial.

14.7 From AIG to the Wider Crisis: Diagnosis and Policy Lessons

Having explained the local mechanism, Greenberg widens the frame. He now asks why the country got into this trouble at all. The lecture changes level here. We move from one firm’s breakdown to a causal chain for the crisis as a whole.

The chain begins with housing policy. Home ownership is widened without sufficient distinction between those who can afford mortgages and those who cannot. Historically, local banks knew borrowers and held the mortgages. Under the later system, Fannie Mae and Freddie Mac bought mortgages from local banks, which then serviced the loans without retaining the same economic stake. That weakens local discipline at origination.

The next step is leverage. Greenberg says investment banks moved from leverage ratios on the order of five to seven times capital to much more extreme numbers:

$$L = \frac{A}{E}, \quad \text{investment-bank leverage} \approx 30 \times \text{ to } 50 \times \text{ capital,} \quad \text{older benchmark} \approx 5 \times \text{ to } 7 \times . \quad (14.30)$$

The lecture treats this as an obvious increase in fragility.

Then comes securitization. Mortgages from different regions are pooled and sold as if geographic spread were enough to make them safe. The rhetoric of diversification reappears here, but now in a distorted form. In the insurance discussion, diversification stabilized a sound underwriting book. In the structured-finance discussion, diversification is used as a sales pitch to justify inflated AAA ratings on bad underlying assets.

The sequence Greenberg gives can be written schematically as

$$\text{weak mortgage discipline} \Rightarrow \text{mortgage sale into secondary market} \quad (14.31)$$

$$\Rightarrow \text{CDO packaging} \quad (14.32)$$

$$\Rightarrow \text{inflated AAA ratings} \quad (14.33)$$

$$\Rightarrow \text{CDS exposure} \quad (14.34)$$

$$\Rightarrow \text{collateral stress} \quad (14.35)$$

$$\Rightarrow \text{system-wide breakdown.} \quad (14.36)$$

He then adds mark-to-market accounting as an amplifier:

$$V_t^{\text{book/held-to-maturity}} \longrightarrow V_t^{\text{market}}. \quad (14.37)$$

The lecture’s claim is not that mark-to-market accounting invented bad assets. It is that, introduced at exactly the wrong moment, it forced severe writedowns and destroyed capital in ways Greenberg sees as artificial. Again, that is his diagnosis, but it fits the overall structure of his argument: several individually dangerous elements came together at once.

At the end of the lecture the student questions sharpen these claims rather than digressing from them. On government intervention, Greenberg says the least amount necessary, and only rarely for truly systemic institutions. On boards, he says boards failed when they did not grasp what was happening after his departure. On banks, he distinguishes commercial from investment banks, supports restrictions like the Volcker Rule on putting the whole institution at risk, and says an investment bank should not be allowed to sell a client a product and then take a position against it. On entrepreneurship, he says insurance remains full of opportunity because new products and new economic environments constantly generate new insurable needs. On China, he says the insurance industry there has grown rapidly but still has much to learn about risk, and he uses the shift from salaried sellers to commission-based agents as an example of institutional modernization.

14.7.1 Question & Answer

What regulation does Greenberg think would have prevented the collateral spiral?

His answer is explicit. First, return the CDS contract to a default-based trigger:

$$\text{CDS should respond only to default, not merely to a temporary fall in market value.} \quad (14.38)$$

Second, create exchange trading so that there is genuine price discovery:

$$\text{exchange trading} \Rightarrow \text{price discovery} \Rightarrow \text{more disciplined collateral and regulation.} \quad (14.39)$$

Third, if the instrument functions as insurance, regulate it as insurance and require reserves. The logic is that price discovery, reserve requirements, and a default trigger all work in the same direction: they break the link between temporary marks and system-wide liquidity panic.

The last Q&A of the lecture returns, fittingly, to the point that has been latent all along. Greenberg says the culture starts at the top, and so does risk management. The CEO is the top risk manager in the company. That closing remark is not an afterthought. It gathers the whole lecture into a single principle. Underwriting, incentives, governance, product design, and crisis response all come back to who defines the risk discipline of the institution.

14.8 Summary

Greenberg’s lecture belongs in a finance course precisely because it moves beyond narrow formula manipulation without ever ceasing to be analytical. It begins with underwriting, where a firm studies a risk and decides whether to accept it. It scales that logic into reinsurance, acquisitions, diversification, and global market opening. It then turns inward to compensation, mobility, and culture as the internal technologies that let a large organization act coherently.

The crisis section gives the chapter its sharpest mathematical content. Once credit protection is tied to changes in marked value rather than realized default, and once collateral has to be posted into a market without real price discovery, the distinction between solvency and liquidity becomes decisive. AIG, in Greenberg’s account, did not first fail because its insurance subsidiaries were worthless. It failed because a contractual rule transformed falling marks into immediate cash demands. The lecture’s broader lesson is therefore plain. Financial institutions do not break only because assets are bad. They also break because incentives, governance, valuation rules, and funding structures interact in the wrong order and at the wrong time.